



OCTOBER 1-2, 2020  
BRIEFINGS

# MAY THE TRUST BE WITH YOU:

EMPOWERING TRUSTZONE-M  
WITH MULTIPLE TRUSTED ENVIRONMENTS

# ABOUT US



## DANIEL OLIVEIRA (🐦@danieljcoliv)

- PhD Student at University of Minho, Portugal
- Active member of ESRGv3
- Worked in a number of research projects from the automotive industry (Bosch) to embedded security (Hex Five Security, Inc)



## SANDRO PINTO (🐦@sandro2pinto)

- Research Scientist & Professor at University of Minho, Portugal
- Active and leading member of ESRGv3
- Working as an External Collaborator in Hex Five Security, Inc

# BACKGROUND

- R&D with Arm TrustZone technology since 2013
- We have published 15+ papers/articles about TrustZone
- Highlights:

### SoK: Understanding the Prevailing Security Vulnerabilities in TrustZone-assisted TEE Systems

David Cerdeira  
Centro Algoritmi  
Universidade do Minho  
david.cerdeira@dei.uminho.pt

Nuno Santos  
INESC-ID / Instituto Superior Técnico  
Universidade de Lisboa  
nuno.santos@inesc-id.pt

Pedro Fonseca  
Department of Computer Science  
Purdue University  
pfonseca@purdue.edu

Sandro Pinto  
Centro Algoritmi  
Universidade do Minho  
sandro.pinto@dei.uminho.pt

**Abstract**—Hundreds of millions of mobile devices worldwide rely on Trusted Execution Environments (TEEs) built with Arm TrustZone for the protection of security-critical applications (e.g., DRM) and operating system (OS) components (e.g., Android kernels). TEEs are often assumed to be highly secure; however, over the past years, TEEs have been successfully attacked multiple times, with highly damaging impact across various platforms. Unfortunately, these attacks have been possible by the presence of security flaws in TEE systems. In this paper, we aim to understand which types of vulnerabilities and limitations affect existing TrustZone-assisted TEE systems, what are the main challenges to build them correctly, and what contributions can be borrowed from the research community to overcome them. To this end, we present a security analysis of popular TrustZone-assisted TEE systems (targeting Cortex-A processors) developed by Qualcomm, Trustonic, Huawei, Nvidia, and Linaro. By studying publicly documented exploits and vulnerabilities as well as by reverse engineering the TEE firmware, we identified several critical vulnerabilities across existing systems which makes it legitimate to raise reasonable concerns about the security of commercial TEE implementations.

**Index Terms**—TEE, TrustZone, Security Vulnerabilities, Arm

#### I. INTRODUCTION

Trusted Execution Environments (TEEs) are a key security mechanism to protect the integrity and confidentiality of applications. By leveraging dedicated hardware, TEEs enable the execution of security-sensitive applications inside protected domains isolated from the platform's operating system (OS). Arm TrustZone [1] has become the de facto hardware technology to implement TEEs in mobile environments and has been employed in industrial control systems [2], servers [3], and low-end devices [4]. In the future, where trillions of TrustZone-enabled IoT devices are expected worldwide [5], TEEs can provide secure environments for data processing at the edge.

TrustZone-assisted TEEs are generally assumed to be more secure than modern OSes due to the hardware-based separation enforced by TrustZone technology and their smaller Trusted Computing Base (TCB), which is several orders of magnitude smaller than standard OSes<sup>1</sup>. For this reason, TEEs have become widely adopted for securing mobile devices against malware [6–10]. For instance, Android platforms incorporate TrustZone-assisted TEEs to secure application-specific operations involving, e.g., user authentication [11], online banking [12], or DRM [13]. Unfortunately, some of these systems have been exploited over the past years, which casts doubt on the real security guarantees that existing commercial TEEs can effectively provide.

In this paper, we perform a systematic study of publicly disclosed vulnerabilities in commercial TrustZone-assisted TEEs for Arm Cortex-A devices. Despite the existence of multiple security reports affecting such systems, this information tends to be scattered and, in certain cases, unverified, which makes it difficult to obtain a comprehensive understanding of the prevailing vulnerabilities and overall security properties of these systems. To fill this gap, we analyzed 207 TEE bug reports spanning a nearly 5 years, from 2013 until mid-2018, focusing on widely deployed TEE systems developed for Arm-based devices by five major vendors: Qualcomm, Trustonic, Huawei, Nvidia, and Linaro. We examined and categorized numerous vulnerabilities, in particular, some of those that have been leveraged to carry out successful attacks. From our analysis, along with the manual inspection of TEE firmware, we have gained multiple insights about the extent and causes of existing vulnerabilities, and about potential solutions to mitigate them.

One first observation is that TEE systems have a long history of critical implementation bugs. Numerous bugs have been (and continue to be) found inside TEE applications – named Trusted Applications (TAs) – and inside the trusted kernel responsible for managing the TEE runtime. Many bugs involve classic input validation errors, such as buffer overflows. As shown by multiple attacks, these bugs can be leveraged to hijack Android's Linux kernel or to entirely compromise the TEE kernel of devices featuring TEEs by Qualcomm [14, 15], Trustonic [16, 17], or Huawei [18].

Second, exploiting vulnerable TAs is facilitated by the numerous architectural deficiencies of TrustZone-assisted TEE systems. For instance, the memory protection mechanisms commonly found in modern OSes, e.g., ASLR or page guards, are almost absent or ill-implemented in most analyzed systems. TEE systems also tend to expose a large attack surface, including dangerous TEE kernel system calls that can be invoked by TAs. For example, on Qualcomm's TEE, any TA can map in memory regions of the host OS. As a result, by hijacking a vulnerable TA, e.g., leveraging a buffer overflow, an attacker can easily control Android [15].

Third, important hardware properties are overlooked in most TrustZone systems at the architectural and microarchitectural levels, which can compromise the security of the TEE. Some vulnerabilities are caused by unexpected behavior of trusted hardware components due to microarchitectural side-channels (e.g., in caches) [19–23]. Others are caused by components that can be leveraged to exfiltrate sensitive data from TEE-protected

### Demystifying Arm TrustZone: A Comprehensive Survey

SANDRO PINTO, Centro Algoritmi, Universidade do Minho  
NUNO SANTOS, INESC-ID, Instituto Superior Técnico, Universidade de Lisboa

The world is undergoing an unprecedented technological transformation, evolving into a state where ubiquitous Internet-enabled “things” will be able to generate and share large amounts of security- and privacy-sensitive data. To cope with the security threats that are thus foreseeable, system designers can find in Arm TrustZone hardware technology a most valuable resource. TrustZone is a System-on-Chip and CPU system-wide security solution, available on today's Arm application processors and present in the new generation Arm microcontrollers, which are expected to dominate the market of smart “things.” Although this technology has remained relatively underground since its inception in 2004, over the past years, numerous initiatives have significantly advanced the state of the art involving Arm TrustZone. Motivated by this revival of interest, this paper presents an in-depth study of TrustZone technology. We provide a comprehensive survey of relevant work from academia and industry, presenting existing systems into two main areas, namely, Trusted Execution Environments and hardware-assisted virtualization. Furthermore, we analyze the most relevant weaknesses of existing systems and propose new research directions within the realm of tiniest devices and the Internet of Things, which we believe to have potential to yield high-impact contributions in the future.

**CCS Concepts:** • Computer systems organization → Embedded and cyber-physical systems; • Security and privacy → Systems security; Security in hardware; Software and application security;

**Additional Key Words and Phrases:** TrustZone, security, virtualization, TEE, survey, Arm

#### ACM Reference format:

Sandro Pinto and Nuno Santos. 2019. Demystifying Arm TrustZone: A Comprehensive Survey. *ACM Comput. Surv.* 51, 6, Article 130 (January 2019), 36 pages.  
<https://doi.org/10.1145/3291047>

#### 1 INTRODUCTION

Arm TrustZone consists of hardware security extensions introduced into Arm application processors (Cortex-A) in 2004 [1, 63]. More recently, TrustZone has been adapted to cover the new generation of Arm microcontrollers (Cortex-M) [65, 113]. TrustZone follows a System-on-Chip (SoC) and CPU system-wide approach to security. This technology is centered around the concept of protection domains named *secure world* and *normal world*. The software executed by the

This work has been partially supported by COMPETE: POCI-01-0145-FEDER-007043, by COMPETE 2020/Portugal 2020/União Europeia within the project Mobile Security Ticketing (No. 11388), which is presented by Link Consulting Tecnologias de Informação SA, and by FCT-Fundação para a Ciência e Tecnologia within the Project Scope: UIDB/00119/2013, UIDX/00320/2013, SFRH/B3-B1/15236/2017, PTDC/EEI-SC/1741/2014 (Alyssa). Authors' addresses: S. Pinto, Centro Algoritmi, Universidade do Minho, Campus de Azúvmes, 4800-058, Portugal; email: sandro.pinto@dei.uminho.pt; N. Santos, Rua Alves Redel, Lisboa, 1000-029 Lisboa, Portugal; email: nuno.santos@inesc-id.pt.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, to republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from [permissions@acm.org](https://www.acm.org/permissions).

© 2019 Association for Computing Machinery.  
0360-0300/2019/01-ART130 \$15.00  
<https://doi.org/10.1145/3291047>

IEEE  
Symposium on Security and Privacy  
2020

ACM  
Computing Surveys  
2019

BLACKHAT20

ESRGv3

# AGENDA

01

## Introduction

Current state of IoT security, motivation, and goals

02

## Arm TrustZone-M “101”

TrustZone technology in a nutshell

03

## TrustZone Pitfalls

What we have learned so far: the limitations of the dual-world model

04

## uTango

Empowering TrustZone-M with multiple trusted environments

05

## The Numbers

Tests, experiments, and results (Arm Musca-B1)

06

## “Live” Demo

Demo of uTango multi-world IoT stack

07

## Conclusion

Takeway points

# Introduction

---

Current state of IoT security, motivation, and goals



03 Jan 2019

IDC Forecasts Worldwide Spending on the Internet of Things to Reach \$745 Billion in 2019, Led by the Manufacturing, Consumer, Transportation, and Utilities Sectors

EGHAM, U.K., August 29, 2019


**Gartner Says 5.8 Billion Enterprise and Automotive IoT Endpoints Will Be in Use in 2020**

arm

**A trillion devices | A trillion dollars**

One trillion new IoT devices will be produced by 2035

Annual Production of IoT devices



**THE INTERNET OF THINGS**

**Popular Internet of Things Forecast of 50 Billion Devices by 2020 Is Outdated**

Warning: All projections for the Internet of Things are subject to change

2016



2018




2018



2017

We need to address this **cat-and-mouse game** of increasing hacks and patches that **threatens** the IoT

2018



2015



2018



2017




"We are witnessing cyberattacks on **critical** infrastructure, health services systems captured and **held to ransom** and home electronics devices used as **internet gateways** by **hackers**."



# IOT SECURITY PROBLEM

## IoT-based devices requirements

- Connectivity (e.g WiFi, Bluetooth, ZigBee)
- Upgradability (e.g. firmware, features, security)
- Sensitive information (e.g. keys, certificates)



## IoT systems are vertical silos!

All code blocks are **combined** into a single memory image. Therefore, if **one block** is **compromised**, the attacker can **most certainly** exploit any other component – there's **no separation** on this traditional model (e.g., FreeRTOS TCP/IP vulnerabilities) .

## Isolate! TrustZone to the rescue

TrustZone addresses this problem by leveraging a system-wide architecture which is **split into two worlds**: a **secure** and a **non-secure** world. A trust barrier isolates the sensitive code from the rich and untrustworthy software.



But, is TrustZone enough?

# KEY MOTIVATION POINTS

01

## EXPLORE TRUSTZONE-M

TrustZone-M is a fairly **new** technology, with a **very few** number of products/works and very **little** research.

Which are the **key architectural differences** between the two **siblings**?

02

## TRUSTZONE CLASSIC MODEL IS FALLING SHORT

We observed how the **dual-world model** is becoming **impractical** and **limited** to address the continuous growth of the attack surface in small IoT devices.

Which are the **main pitfalls** of TrustZone? And how can we **mitigate** them?

03

## EMPOWER TZ WITH MULTIPLE TRUSTED WORLDS

Apply a new **zero-trust model** to all pieces of code and isolate them in completely **isolated domains** (i.e. 'containers'), by following a complete **horizontal multi-domain** security schema.

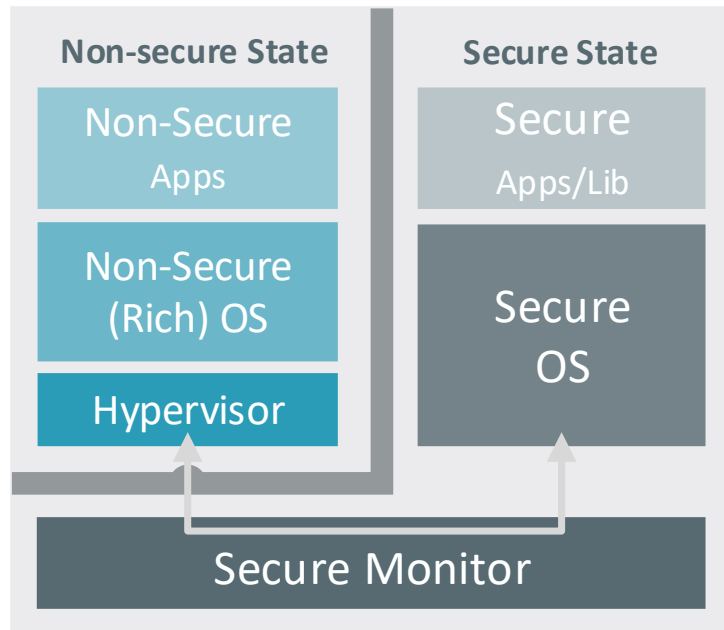
How to leverage the TrustZone-M **hardware features** to implement a **multi-world** approach?



# Arm TrustZone-M “101”


TrustZone technology in a nutshell

# TRUSTZONE VS TRUSTZONE-M



## TrustZone (for Cortex-A)

- Introduced in 2004
- Execution state is determined by the NS bit
- Monitor mode / SMC instruction
- Memory and devices: TZASC and TZPC
- Interrupts: GIC - IRQs vs FIQs



## TrustZone-M (for Cortex-M)

- Introduced in 2018
- Execution state is memory mapped based
- No Monitor mode / SG instruction
- Memory and devices: SAU and IDAU
- Interrupts: NVIC – secure and non-secure IRQs

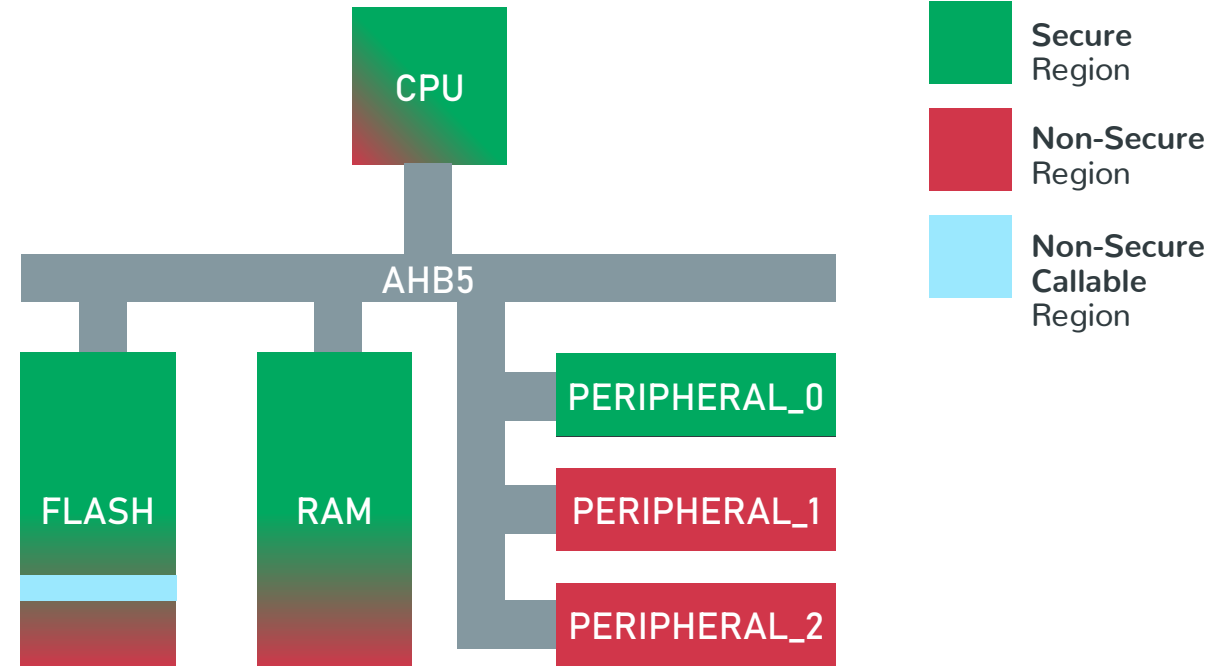
1

2

3

4

# EXECUTION STATE IS MEMORY MAP BASED



On Cortex-M MCUs, SoC resources (FLASH, RAM, peripherals) are mapped to a memory address

- The 4GB memory space is **partitioned** into:
  - **Secure** (software can access secure or non-secure resources)
  - **Non-Secure** (software can only access non-secure resources)
  - **Non-Secure Callable** (memory area only used for state transition)
- The **processor state** is **dependent** on the memory space definition: when the processor is running code in a **Secure** region it is in **secure state**, otherwise it is in **Non-Secure state**



1


2

3

4

# SECURITY STATES

## TRANSITION/SG INST



Cortex-M MCUs have two operation modes: **Handler** (for exception handling) and **Thread** (for normal application code – priv/unpriv)

- TrustZone-M **replicates** operation modes for the **secure** and **non-secure** states - the security states are **orthogonal** to the existing processor modes


1

2

3

4

# SECURITY STATES TRANSITION/SG INST



- **Secure** software can call non-secure software using the **BLXNS** <Ri> inst.
  - Return address is pushed onto the **secure stack**
  - **Link Register** is set to a special value called **FNC\_RETURN**
  - Performing a **branch** to this **LR** automatically triggers the unstacking of the **true return address** from the secure stack and a return transition to the **S** world

1


2

3

4

# SECURITY STATES

## TRANSITION/SG INST



- **Non-Secure** software cannot directly call the **secure** software
  - A **SecureFault** (or a **HardFault**) exception will be triggered
  - Prevents direct calls into protected code



1

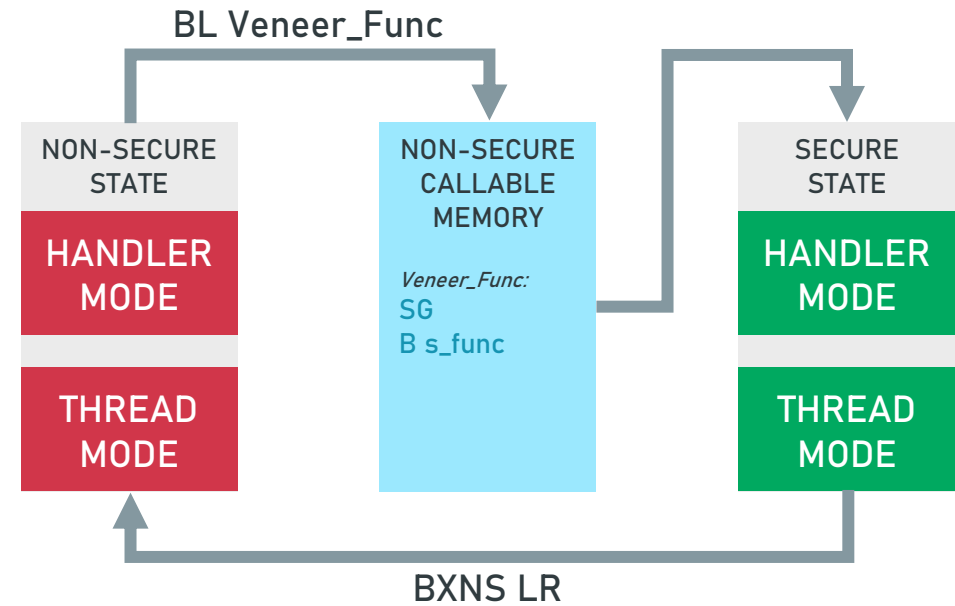
2

3

4

# SECURITY STATES

## TRANSITION/SG INST



- **Non-Secure** software can **only** call **secure** software through specific entry points located in a **Non-Secure Callable** (NSC) memory region
  - The first instruction in the API must be an **SG** (secure gateway)
  - The **SG** instruction must be in an NSC region
  - The last instruction must be a **branch** to the **secure** function
- After completion, the **secure** function returns to the **non-secure** state by using a **BXNS** instruction

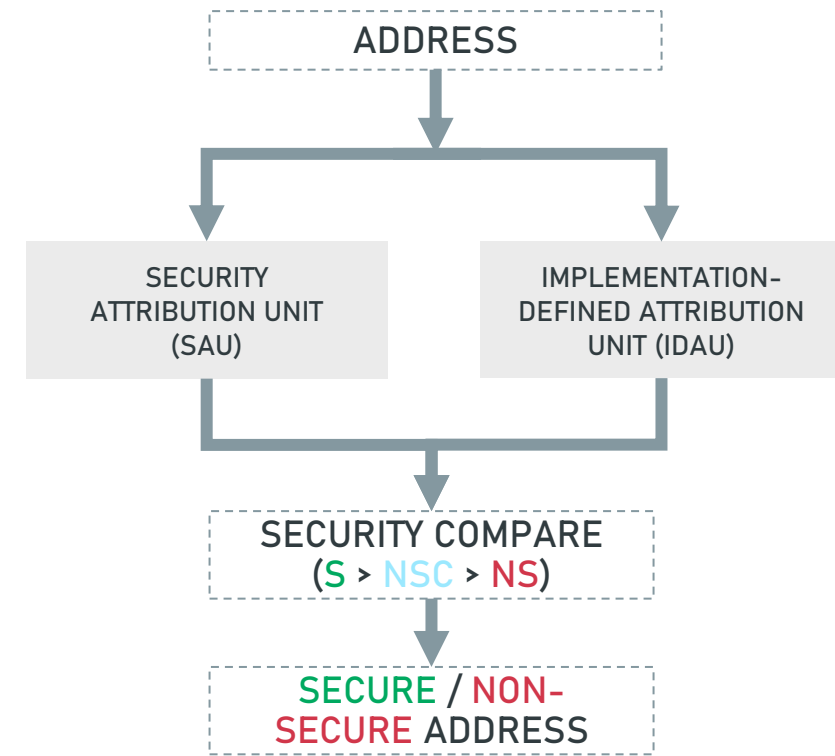
1

2

3

4

# MEMORY & DEVICES PARTITIONING



Memory and devices partitioning is defined by address

- The security state of a memory region is selected by a combination of the **IDAU** and **SAU** attribution units
  - The result is the **logical OR operation** between the **security configuration** of both memory controllers

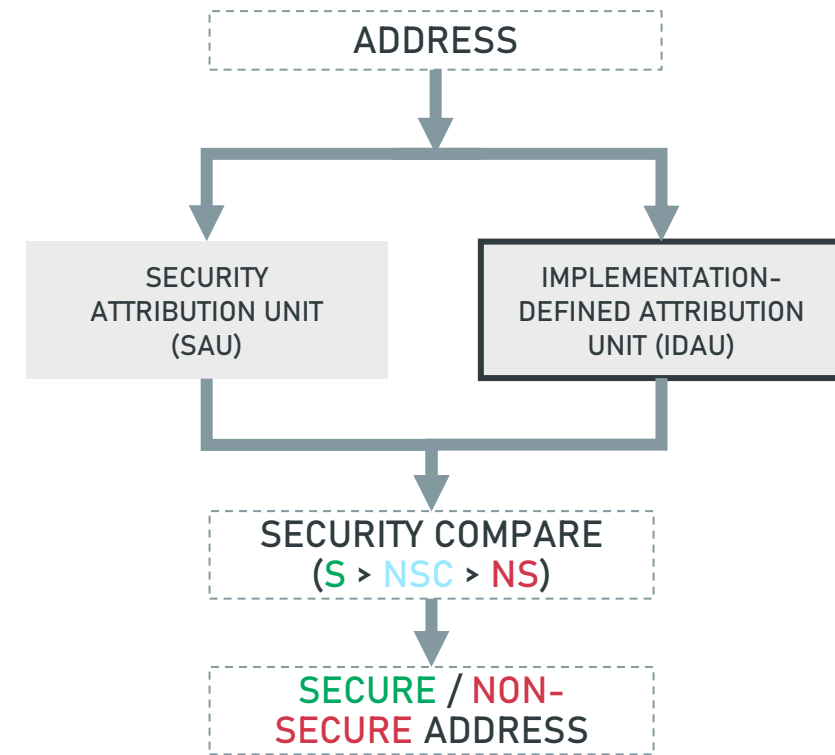
1

2

3

4

# MEMORY & DEVICES PARTITIONING



- **IDAU**
  - Provides **static** address **partitioning**
  - **Non-Programmable**
  - **Customized** by the **silicon vendor** to define a fixed memory map
  - Supports up to 256 regions



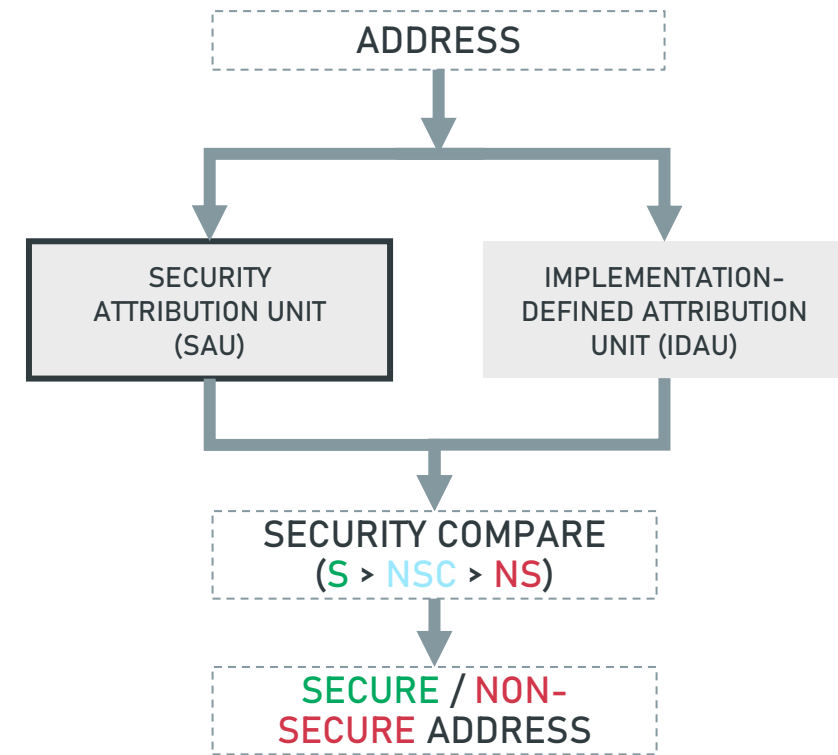
1

2

3

4

# MEMORY & DEVICES PARTITIONING



- **SAU**
  - Provides **dynamic address partitioning**
  - **Programmable** in secure state
  - Always available on all chips with TrustZone support (even if the number of regions is configured to zero)
  - The **number of regions** is defined by the chip designer (typically 8 regions)


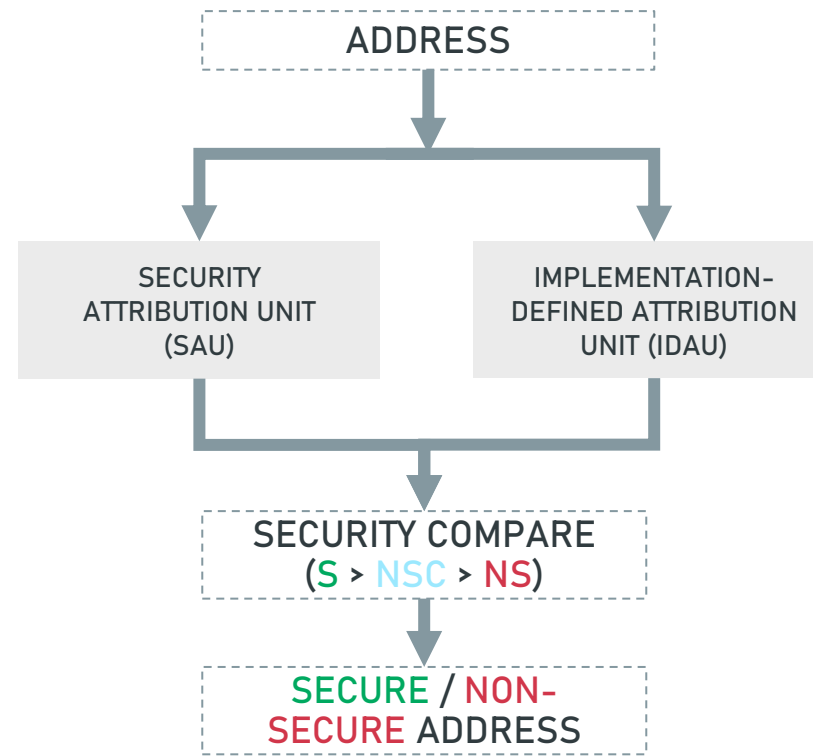
1

2

3

4

# MEMORY & DEVICES PARTITIONING



The result is the logical OR operation between the security configuration of SAU and IDAU

SAU	IDAU	Security State
Non-Secure	Non-Secure	Non-Secure
Secure	Non-Secure	Secure
Non-Secure	Secure	Secure
Secure	Secure	Secure
NSC	Secure	Secure
NSC	Non-Secure	NSC


1

2

3

4

# MEMORY & DEVICES PARTITIONING



The Security state is extended to other SoC resources through “gate-keepers” (bus filters)

- **SAU** and **IDAU** will affect the output from the **CPU to the bus**, whereas **bus filters** affect the accesses from other **bus masters to the resources**
- e.g. Arm’s SSE-200 subsystem specification provides two components that configure **system-wide access rights to memory and peripherals**:
  - MPC: Memory Protection Controller
  - PPC: Peripheral Protection Controller




1

2

3

4

# INTERRUPT HANDLING



Interrupt management is similar to ARMv7-M, where a subset of the core registers is pushed automatically into the stack

- Supports two separated exception vector tables (**Secure** and **Non-Secure**)
- Each interrupt can be assigned to the **secure** or **non-secure** state through a new NVIC register (ITNS – Interrupt Target Non-Secure)
- **S** or **NS** interrupts can share the **same priority levels**, or **secure** interrupts can have higher priority – i.e. DoS (bit PRIS on AIRCR register)

# TrustZone Pitfalls

What we have learned so far: the limitations of the dual-world model

Point: Qualcomm chips leak crypto data from secure execution

Qualcomm chips leak crypto data from secure execution

Qualcomm TEE Bug

LG, Motorola hacked: New Security At Risk

Experts discovered vulnerabilities in Qualcomm TrustZone that Check Point researchers called 'unprecedented access' because of the extremely sensitive data stored in secure elements.

Share

Facebook Twitter LinkedIn Email



# The Cyber-Security source

HOME | **NEWS & FEATURES** | BUYER'S GUIDE | OPINION | E

# Android smartphone fuzzing reveals 'gaping hole' in trusted execution environment



# WHAT HAVE WE LEARNED SO FAR

## SoK: Understanding the Prevailing Security Vulnerabilities in TrustZone-assisted TEE Systems

David Cerdeira  
Centro Algoritmi  
Universidade do Minho  
david.cerdeira@dei.uminho.pt

Nuno Santos  
INESC-ID / Instituto Superior Técnico  
Universidade de Lisboa  
nuno.santos@inesc-id.pt

Pedro Fonseca  
Department of Computer Science  
Purdue University  
pfonseca@purdue.edu

Sandro Pinto  
Centro Algoritmi  
Universidade do Minho  
sandro.pinto@dei.uminho.pt

**Abstract**—Hundreds of millions of mobile devices worldwide rely on Trusted Execution Environments (TEEs) built with Arm TrustZone for the protection of security-critical applications (e.g., DRM) and operating system (OS) components (e.g., Android keystore). TEEs are often assumed to be highly secure; however, over the past years, TEEs have been successfully attacked multiple times, with highly damaging impact across various platforms. Unfortunately, these attacks have been possible by the presence of security flaws in TEE systems. In this paper, we aim to understand which types of vulnerabilities and limitations affect existing TrustZone-assisted TEE systems, what are the main challenges to build them correctly, and what contributions can be borrowed from the research community to overcome them. To this end, we present a security analysis of popular TrustZone-assisted TEE systems (targeting Cortex-A processors) developed by Qualcomm, Trustonic, Huawei, Nvidia, and Linaro. By studying publicly documented exploits and vulnerabilities as well as by reverse engineering the TEE firmware, we identified several critical vulnerabilities across existing systems which makes it legitimate to raise reasonable concerns about the security of commercial TEE implementations.

**Index Terms**—TEE, TrustZone, Security Vulnerabilities, Arm

### I. INTRODUCTION

Trusted Execution Environments (TEE) are a key security mechanism to protect the integrity and confidentiality of applications. By leveraging dedicated hardware, TEEs enable the execution of security-sensitive applications inside protected domains isolated from the platform's operating system (OS). Arm TrustZone [1] has become the de facto hardware technology to implement TEEs in mobile environments and has been employed in industrial control systems [2], servers [3], and low-end devices [4]. In the future, where trillions of TrustZone-enabled IoT devices are expected worldwide [5], TEEs can provide secure environments for data processing at the edge.

TrustZone-assisted TEEs are generally assumed to be more secure than modern OSes due to the hardware-based separation enforced by TrustZone technology and their smaller Trusted Computing Base (TCB), which is several orders of magnitude smaller than standard OSes. For this reason, TEEs have become widely adopted for securing mobile devices against malware [6–10]. For instance, Android platforms incorporate TrustZone-assisted TEEs to secure application-specific operations involving, e.g., user authentication [11], online banking [12], or DRM [13]. Unfortunately, some of these systems have been exploited over the past years, which casts doubt on the real security guarantees that existing commercial TEEs can effectively provide.

In this paper, we perform a systematic study of publicly disclosed vulnerabilities in commercial TrustZone-assisted TEEs for Arm Cortex-A devices. Despite the existence of multiple security reports affecting such systems, this information tends to be scattered and, in certain cases, unverified, which makes it difficult to obtain a comprehensive understanding of the prevailing vulnerabilities and overall security properties of these systems. To fill this gap, we analyzed 207 TEE bug reports spanning a nearly 5 years, from 2013 until mid-2018, focusing on widely deployed TEE systems developed for Arm-based devices by five major vendors: Qualcomm, Trustonic, Huawei, Nvidia, and Linaro. We examined and categorized numerous vulnerabilities, in particular, some of those that have been leveraged to carry out successful attacks. From our analysis, along with the manual inspection of TEE firmware, we have gained multiple insights about the extent and causes of existing vulnerabilities, and about potential solutions to mitigate them.

One first observation is that TEE systems have a long history of *critical implementation bugs*. Numerous bugs have been (and continue to be) found inside TEE applications – named Trusted Applications (TAs) – and inside the trusted kernel responsible for managing the TEE runtime. Many bugs involve classic input validation errors, such as buffer overflows. As shown by multiple attacks, these bugs can be leveraged to hijack Android's Linux kernel or to entirely compromise the TEE kernel of devices featuring TEEs by Qualcomm [14, 15], Trustonic [16, 17], or Huawei [18].

Second, exploiting vulnerable TAs is facilitated by the *numerous architectural deficiencies* of TrustZone-assisted TEE systems. For instance, the memory protection mechanisms commonly found in modern OSes, e.g., ASLR or page guards, are almost absent or ill-implemented in most analyzed systems. TEE systems also tend to expose a large attack surface, including dangerous TEE kernel system calls that can be invoked by TAs. For example, on Qualcomm's TEE, any TA can map in memory regions of the host OS. As a result, by hijacking a vulnerable TA, e.g., leveraging a buffer overflow, an attacker can easily control Android [15].

Third, *important hardware properties are overlooked* in most TrustZone systems at the architectural and microarchitectural levels, which can compromise the security of the TEE. Some vulnerabilities are caused by unexpected behavior of trusted hardware components due to microarchitectural side-channels (e.g., in caches) [19–23]. Others are caused by components that can be leveraged to exfiltrate sensitive data from TEE-restricted

## Main Contributions:

### Extensive

- More than 200 TEE bug reports
- Reverse engineering of TEE firmware
- All major TEE systems

### Classification

- Architectural deficiencies
- Implementation bugs
- Hardware properties

System	Critical	Severe	Medium	Low	Total
Qualcomm TEE	52	19	12	9	92
Trustonic TEE	1	-	0	4	5
Huawei TEE	-	2	-	1	3
Nvidia TEE	-	5	1	4	10
Linaro TEE	-	-	2	1	3
Other	-	1	7	3	11
TEE Total	53	27	22	22	124
FreeRTOS	-	-	5	8	13
VxWorks	2	2	5	1	10
Linux	242	254	393	758	1647

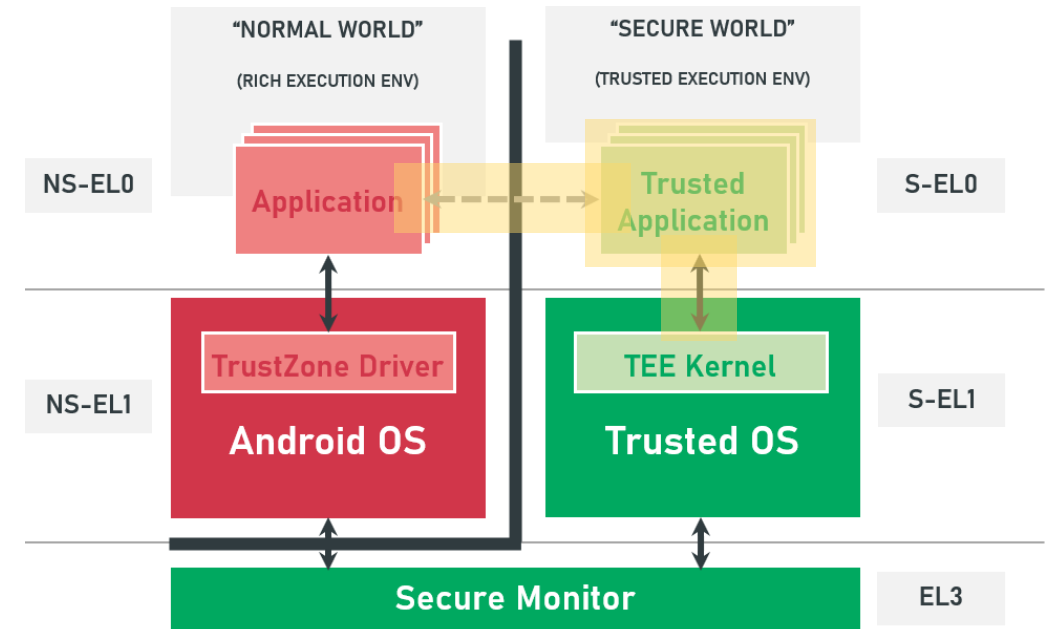
# Architctural Issues (TEE attack surface)

1

Wide interfaces in TEE  
componentes

2

Excessively large TEE  
TCBs



These wide interfaces are presented in TAs, the Trusted-OS, and even on device drivers

- The **Widevine TA** (DRM tool) implements 70 different commands (many of them manipulate complex media data structures)
- The **Qualcomm's TEE** kernel provides to TAs 69 syscalls (e.g. mapping normal world memory)
- The **Trustonic TEE** shares the fingerprint device driver with almost all TAs, which might compromise user authentication

Wide interfaces in TEE components promote designs  
more prone to vulnerabilities \*

\*Jo Van Bulck et al. "A Tale of Two Worlds: Assessing the Vulnerability of Enclave Shielding Runtimes", ACM CCS, 2019.




# Architeturational Issues (TEE attack surface)

1

Wide interfaces in TEE  
componentes

2

Excessively large TEE  
TCBs



System	Core (bin / src)	TAs
Qualcomm TEE (Google Pixel XL)	1.61MB / -	2.71MB
Trustonic TEE (Samsung S7)	350KB / -	5,02MB
Huawei TEE (Huawei P8 Lite)	744KB / -	479KB
Nvidia TEE (Nvidia Tegra)	97KB / 123Kloc	80KB
Linaro TEE (Hikey960)	365KB / 210Kloc	-
Linux (4.14.rc7)	19MB / 15Mloc	-
seL4 (kernel)	166.5KB / 19Kloc	-

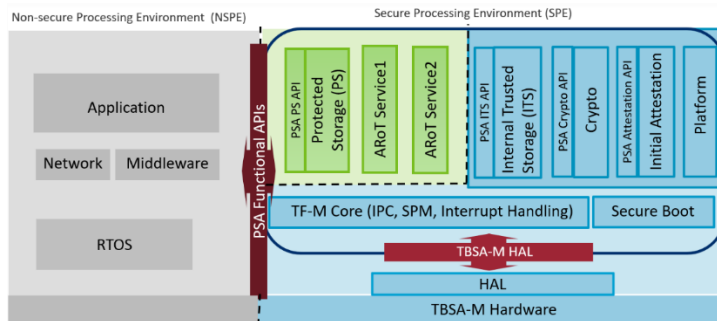
We have measured binary sizes and the number of lines of code of 5 commercial TEEs

- TEEs are fairly **bigger** than similar software systems, such as the **seL4** (security-driven microkernel)
- TAs can even **increase the TCB** with megabytes of code that need **also** to be **trusted** alongside the TEE

TEEs are too **big** and too **complex** to be secure!

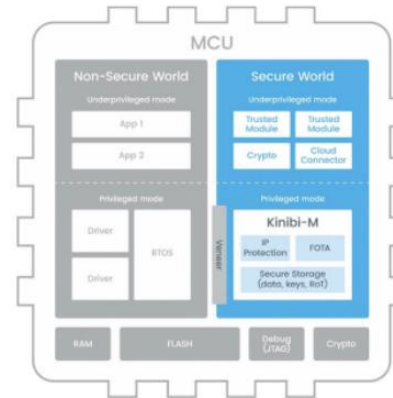
# ... AND WHAT ABOUT TRUSTZONE-M?

## Example 1 Trusted Firmware-M



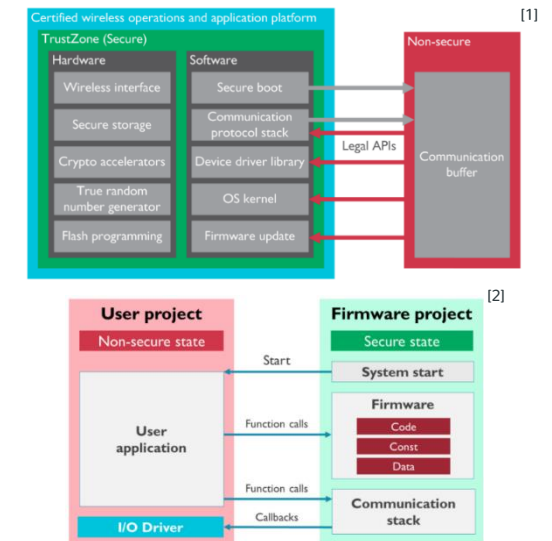
Open governance project, that “provides a reference implementation of secure world software for (...) Armv8-M.”

## Example 2 Trustonic Kinibi-M



Commercial “secure TEE operating system”, where “the platform is adapted from Trustonic’s existing Kinibi technology that already secures more than 1.7 billion smartphones and tablets.”

## Other Examples Across Arm Documentation



[1] Arm, “TrustZone technology for the ARMv8-M architecture”, 2017.

[2] Keil, “Using TrustZone on ARMv8-M”, 2016.

We by no means intend to diminish the value of these solutions...

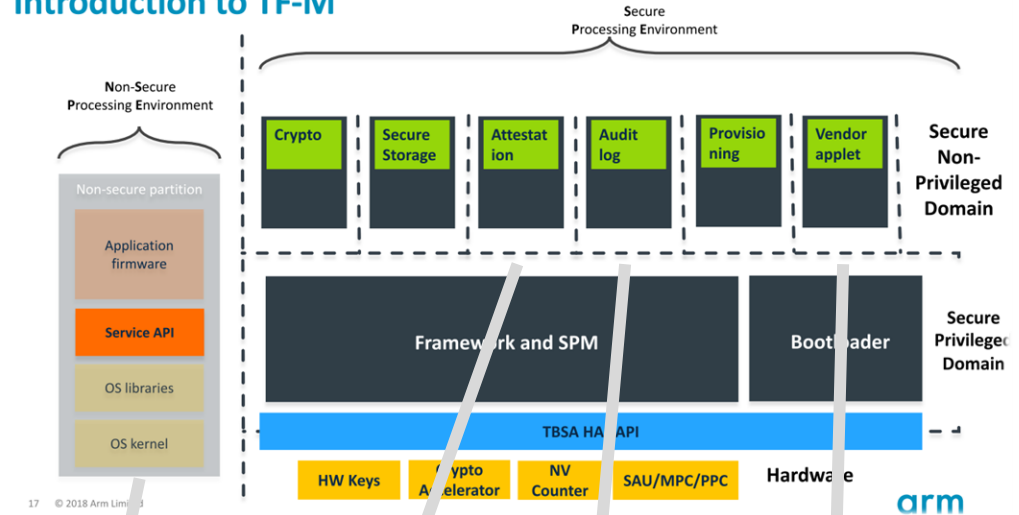
... but with what we have seen so far, it is very likely that the same problems will happen again in the near future.

# uTango

Empowering TrustZone-M with multiple trusted environments

# is there a better solution?

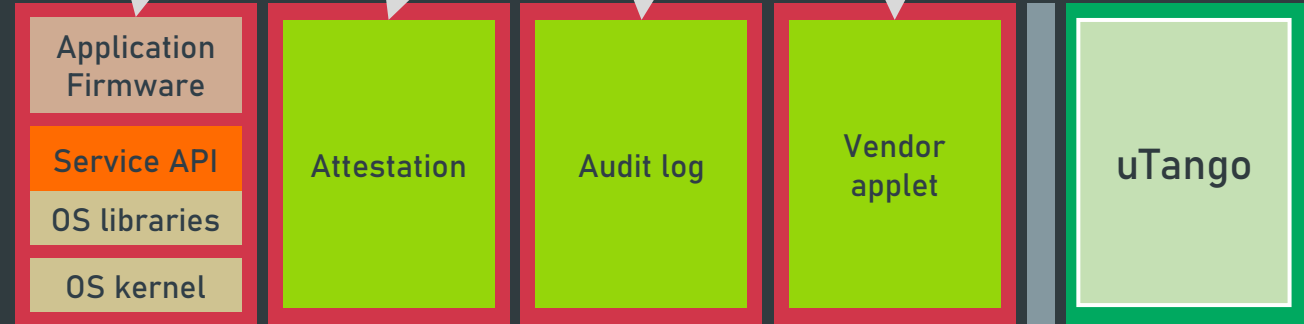
## Introduction to TF-M



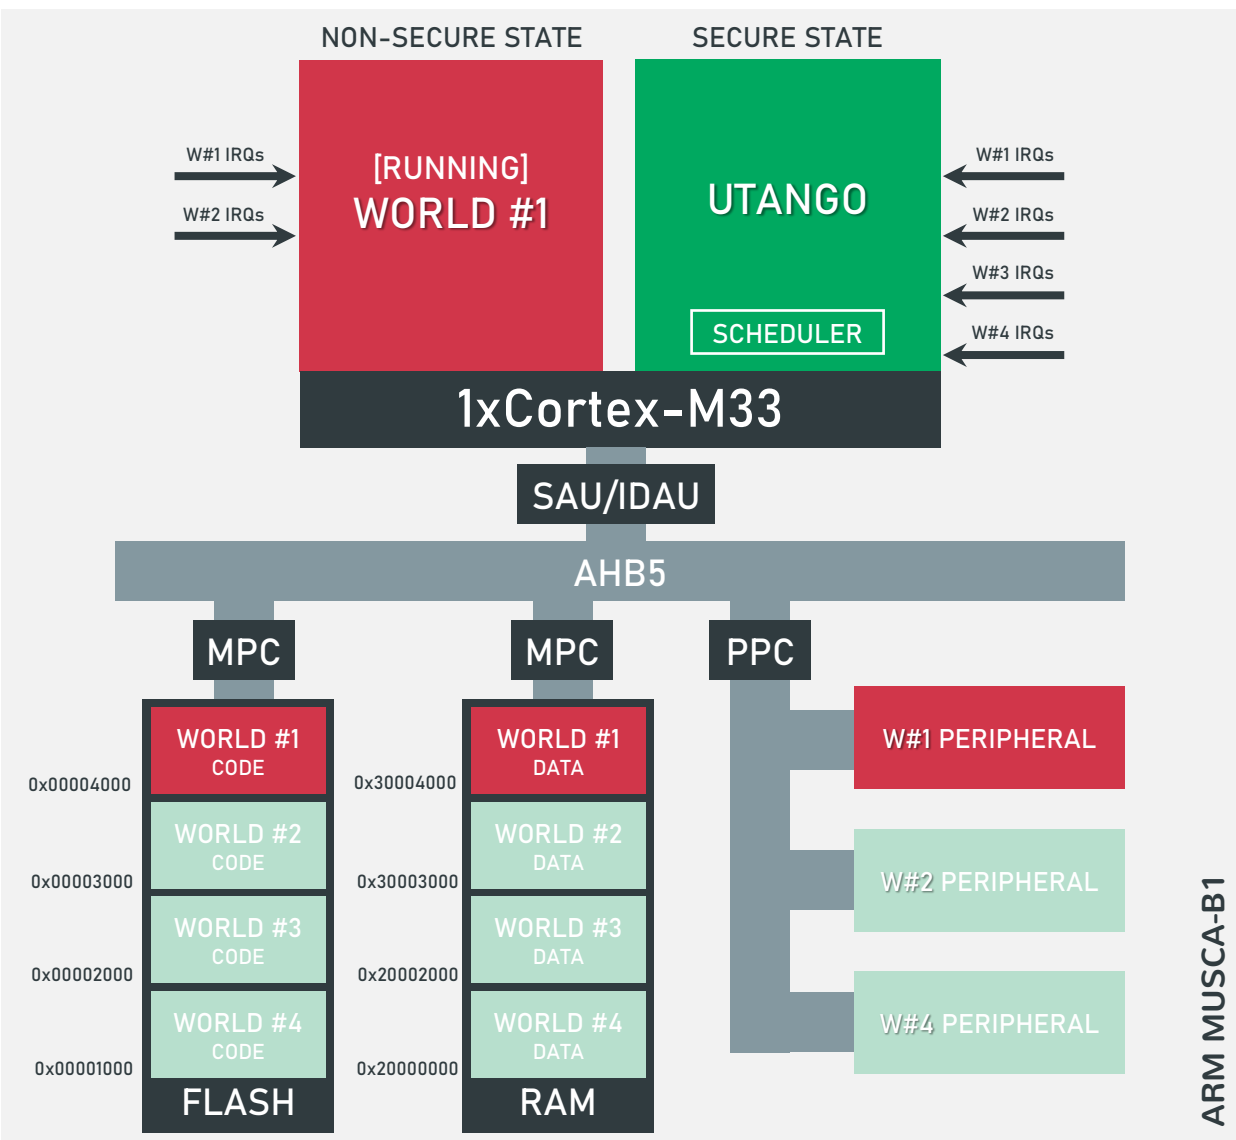
expands the **LIMITED** TrustZone-M dual-world model

applies a **ZERO-TRUST** model to all pieces of code

implements a **MULTI-WORLD** security schema with equally-secure worlds



# ARCHITECTURE & TOP-10 HIGH-LEVEL DETAILS



You can view **uTango** as a TEE  $\mu$ kernel that provides the mechanisms that allow the consolidation of **multiple, equally secure, trusted environments** on a TrustZone-M-enabled platform

1

**uTango** sits on the **Secure** side and, therefore, is the **first** software component to run, working as the secure bootloader of the system.

2

**Worlds** are guaranteed to run in a **round-robin** fashion, supported by the **uTango scheduler**

3

The extended multiple **trusted environments** (or worlds) run in the **Non-Secure** side, and are preserved in the secure side when suspended

4

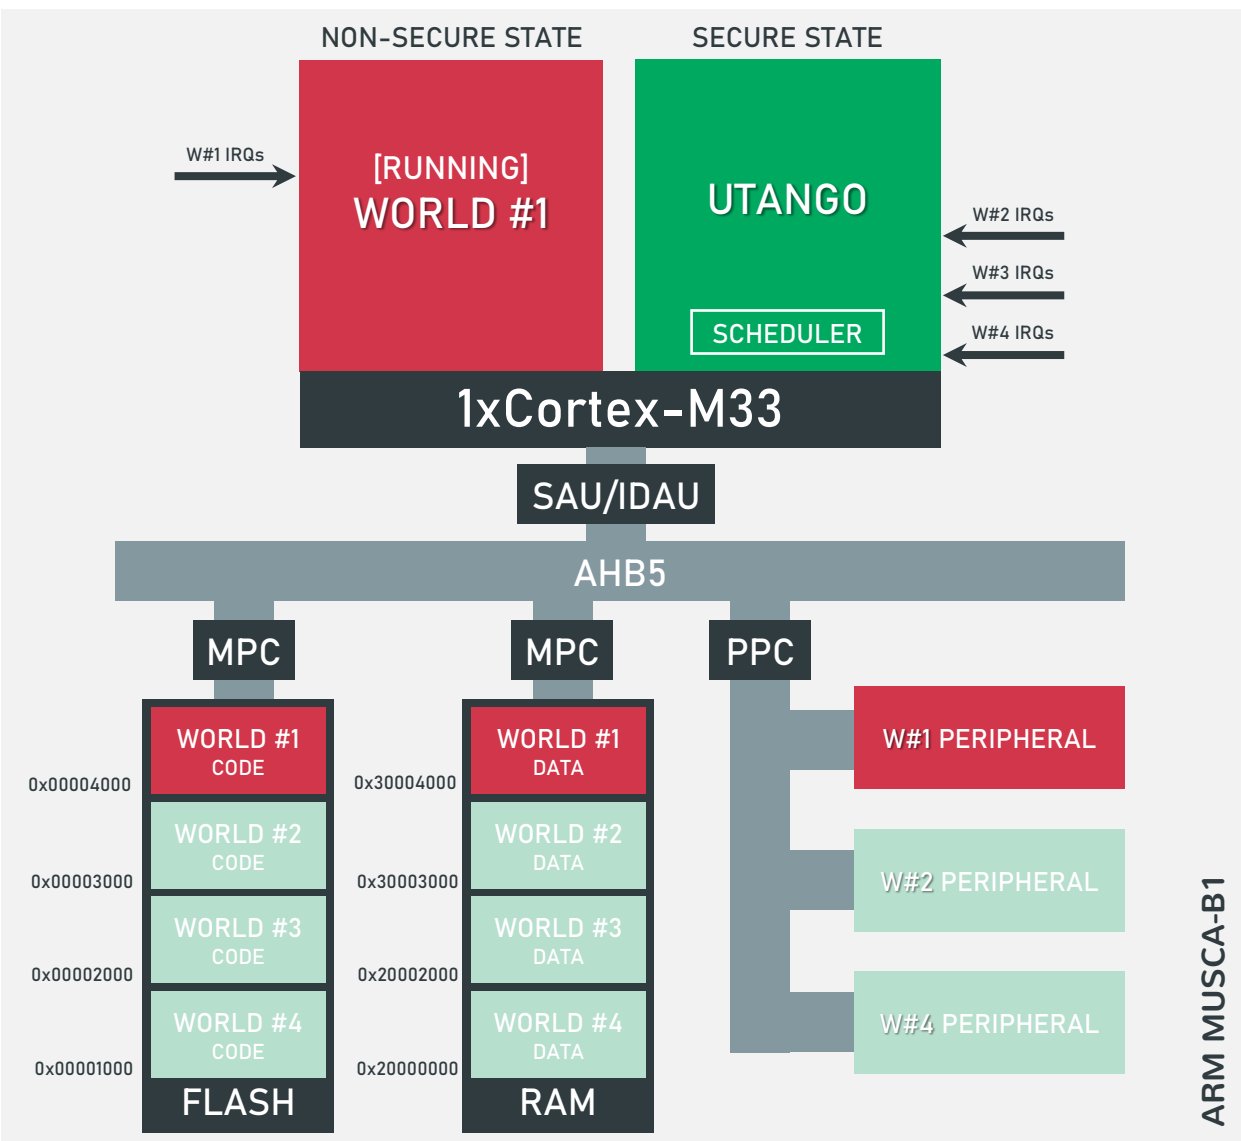
**SAU** is **reconfigured** at each scheduling point, to guarantee the **isolation** of each world and its resources (running or suspended)

5

**Interrupts** security assignment and **state** (enable/disable) is reconfigured at each scheduling point, depending on the world resumed or suspended



# ARCHITECTURE & TOP-10 HIGH-LEVEL DETAILS



You can view **uTango** as a TEE  $\mu$ kernel that provides the mechanisms that allow the consolidation of **multiple, equally secure, trusted environments** on a TrustZone-M-enabled platform

- 6 The **execution state** of each world is **saved or restore** during the context switching into/from a **control block structure**
- 7 **uTango scheduler** is supported by the **architectural system timer** (SysTick), which is **banked** between states
- 8 **MPC** and **PPC** bus filter controllers are configured in **boot time** to match the worlds memory space.
- 9 **uTango** is **statically** configured (specify which resources will be used by each world, memory area, interrupts assigned, vector table)
- 10 **uTango** is supporting **Arm Musca-B1** and other commercial platforms are in the roadmap.

# UTANGO: A “TOY” USE-CASE

## WORLD #1

Bare-metal application featuring a serial terminal

### Memory Regions:


- Code: FLASH
- Data: Internal SRAM1
- Devices: UART0

## WORLD #2

Bare-metal application blinking a LED at each timer tick interrupt (1ms)

### Memory Regions:

- Code: FLASH
- Data: Internal SRAM2
- Devices: PWM0, TIMER0



## UTANGO

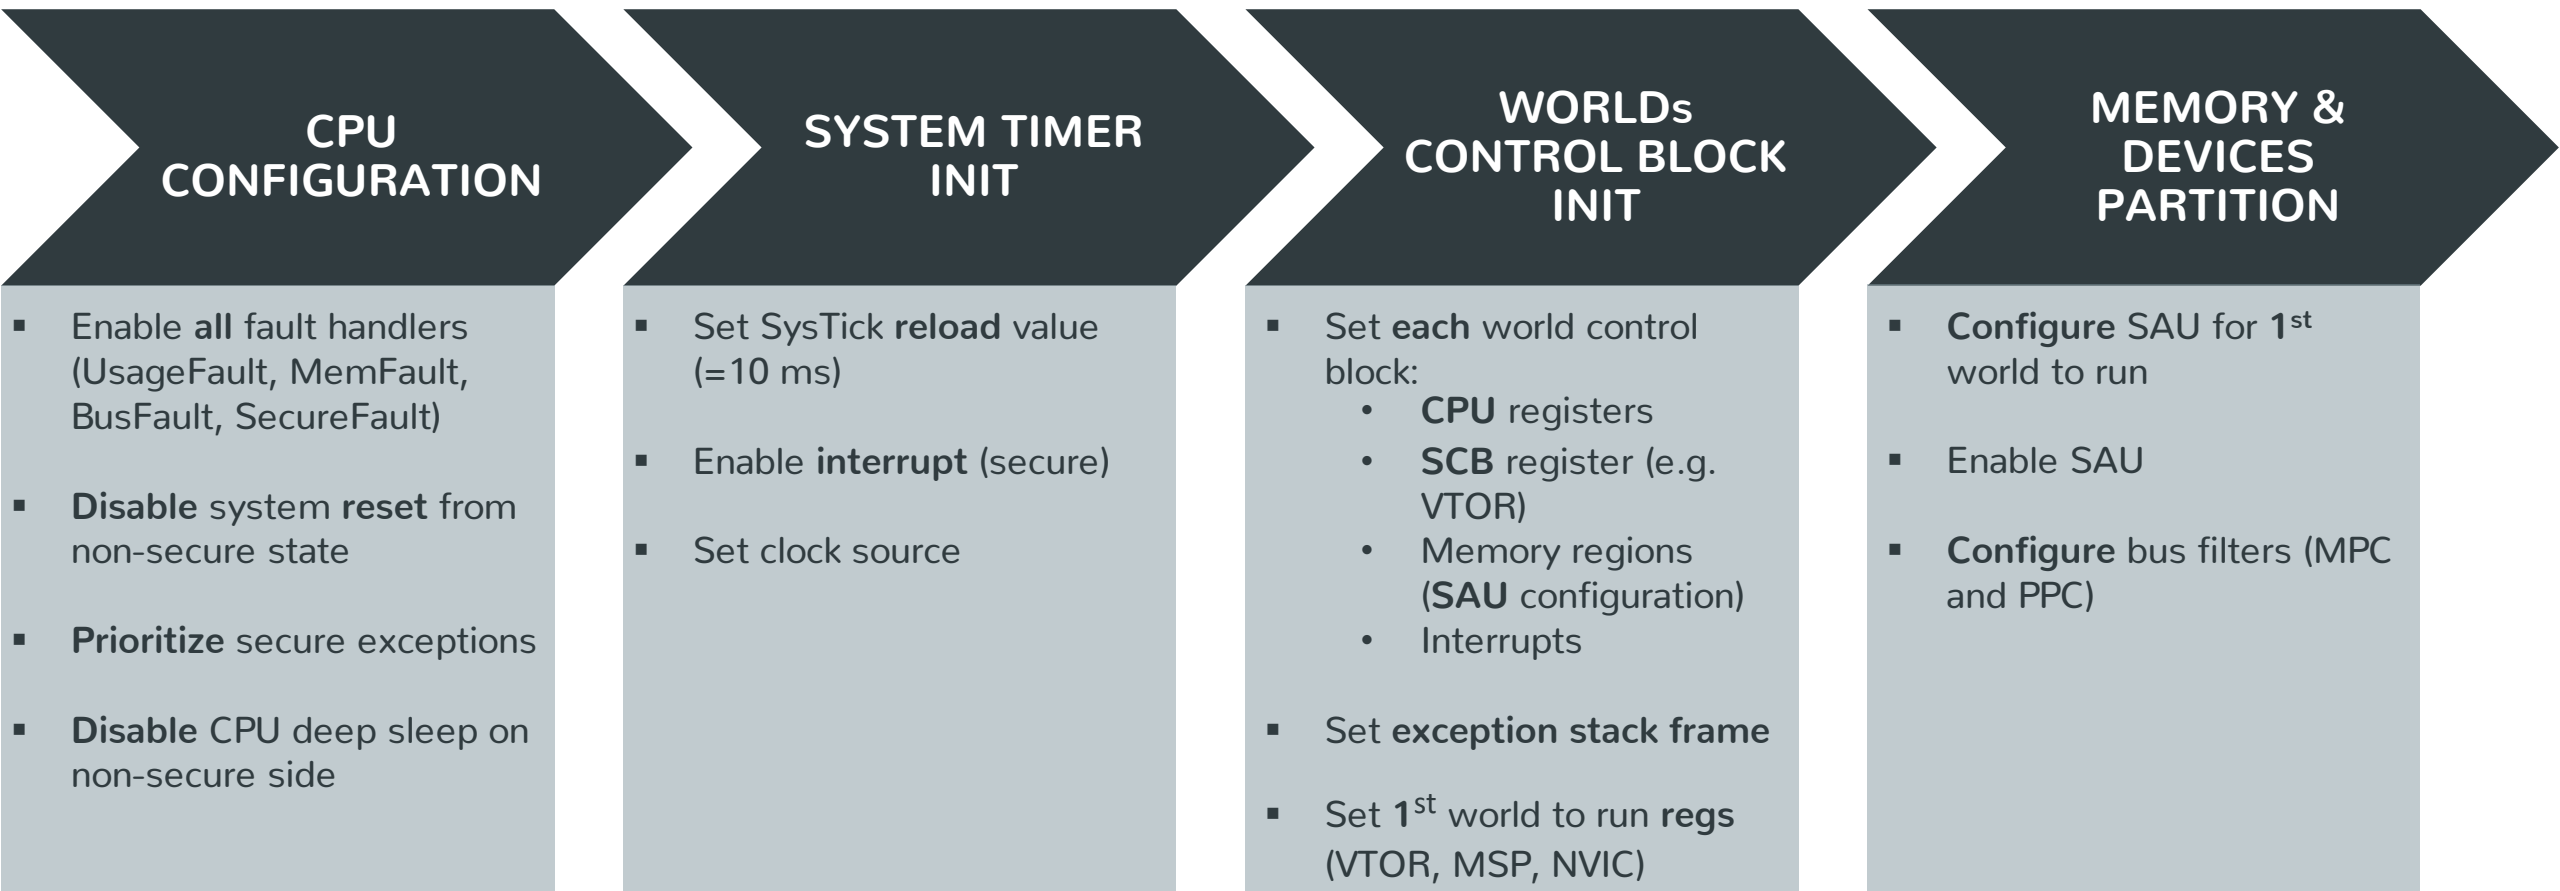
### Configurations:


- N worlds: 2
- Tick: 10 ms

### Memory Regions:

- Code: Internal SRAM0 (TCM)
- Data: Internal SRAM0 (TCM)

# BOOT SEQUENCE

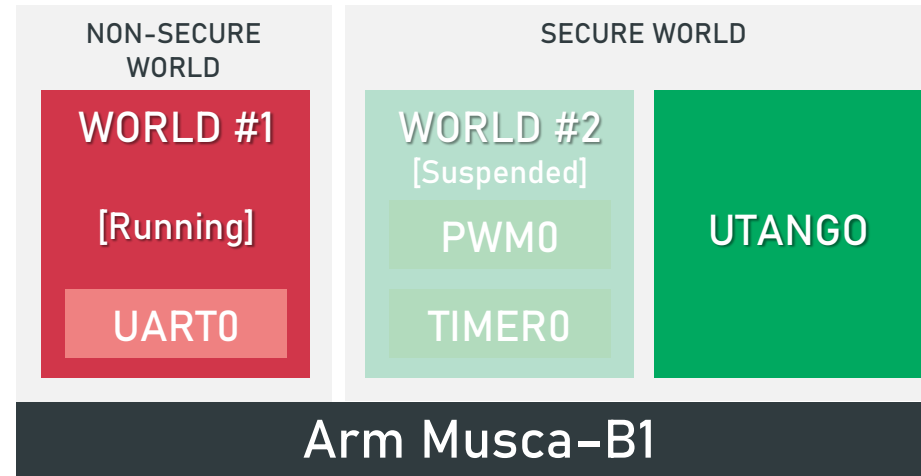





# BOOT SEQUENCE

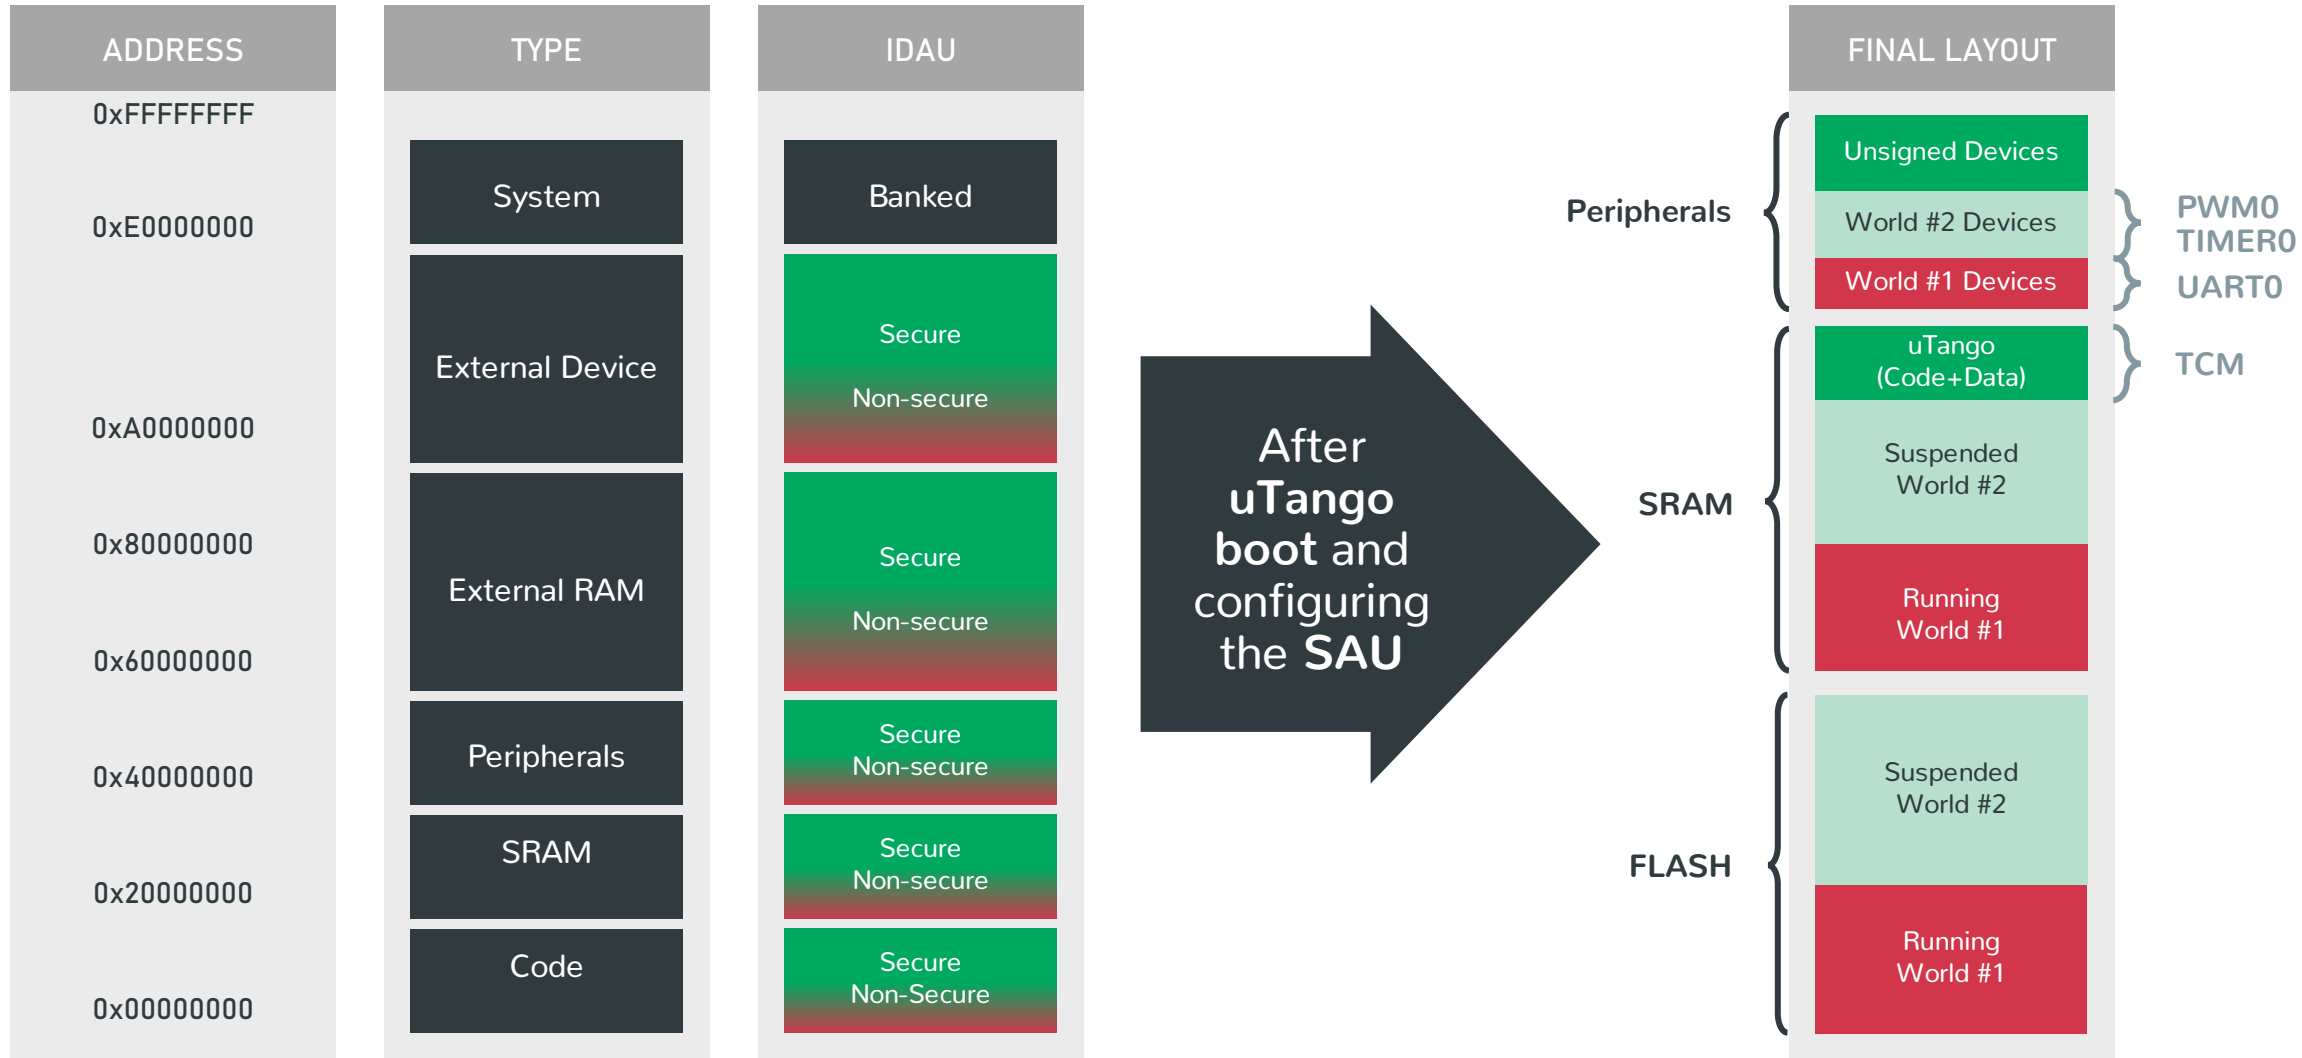
## UTANGO START


- Enable **System Timer**
- **KICKOFF!** Jump to 1<sup>st</sup> world by means of a BLXNS instruction (non-secure function call, that switches from secure to non-secure state)



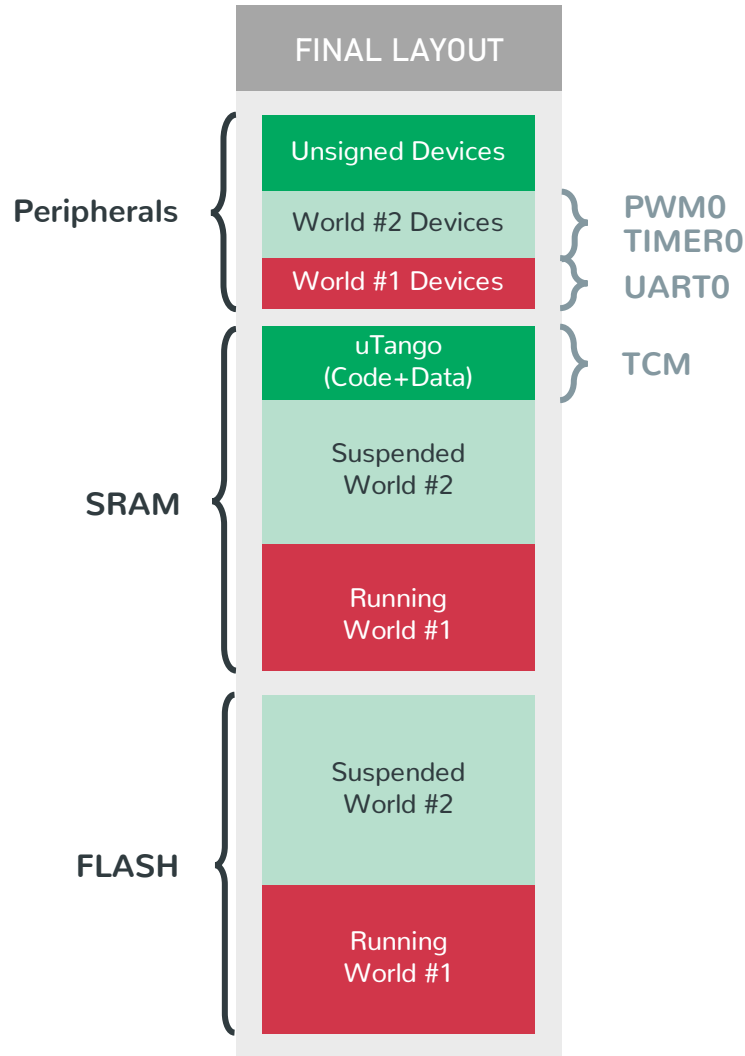


# MEMORY & DEVICES PARTITION





# MEMORY & DEVICES PARTITION



uTango sets up the **memory space** and creates **isolation boundaries** between worlds.

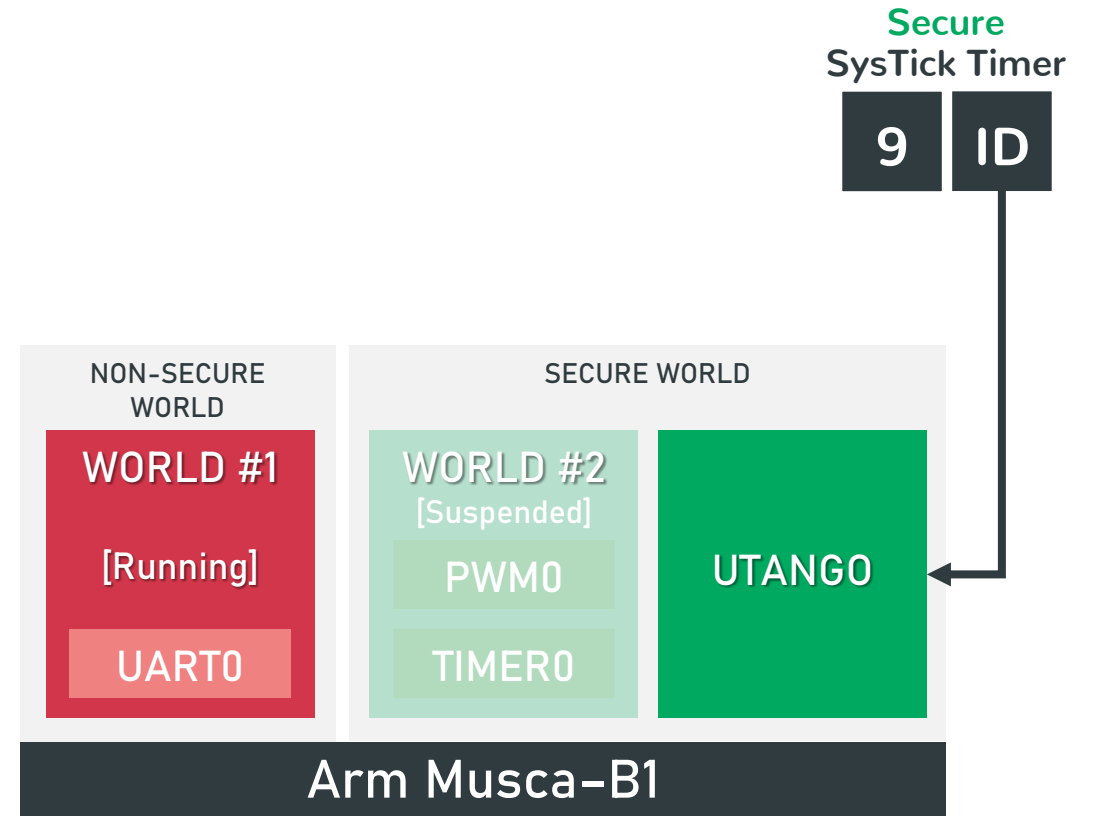
**Bus filters** are also configured to match the **memory** and **devices** partition settings


- The SAU is **enabled**
- The **non-secure** execution context for the **1<sup>st</sup>** world to run is prepared
- Remaining worlds are **suspended** in the **secure** side until a **next context switching**

- The **assumption** is that other **bus master** besides the CPU (e.g. DMA, crypto engines) are **sandboxed** into the **secure** side
- This **prevents** the **reconfiguration** of each bus filter during the context switching and the inherent performance penalties

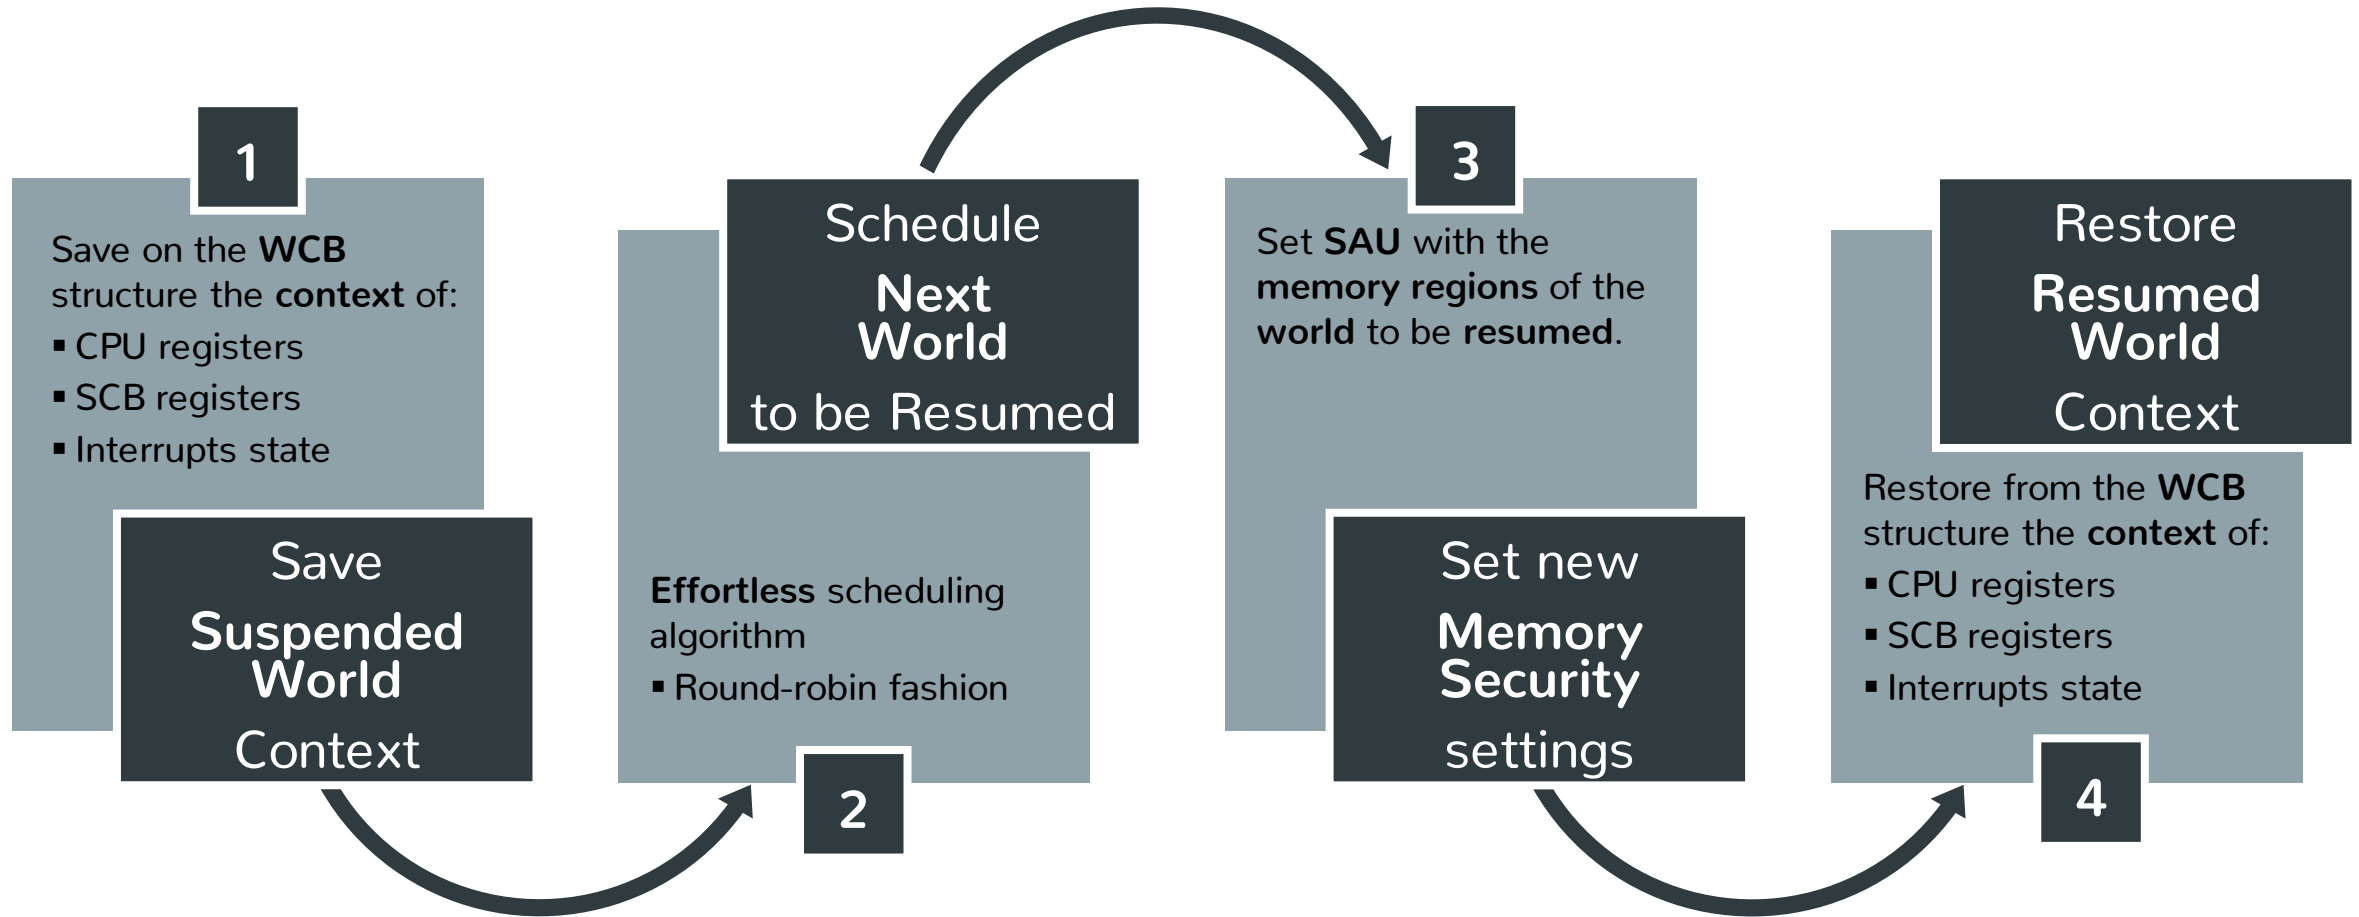


what happens  
during a  
**SCHEDULING**  
point?

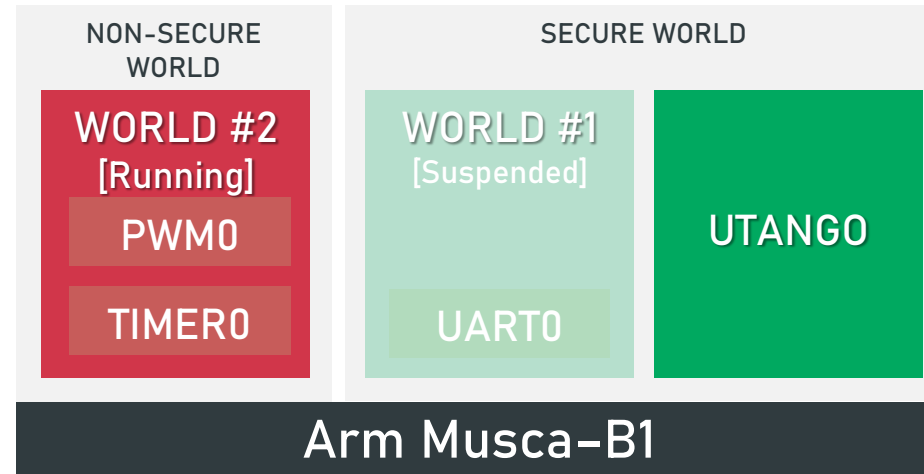





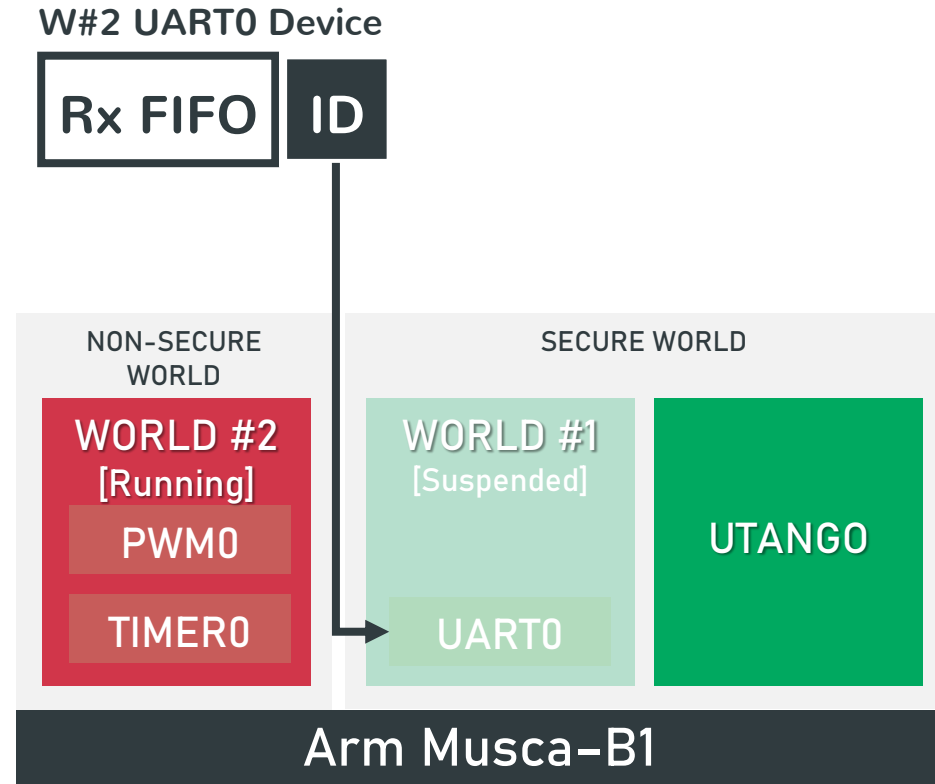
# SCHEDULER



# SCHEDULER




and what happens  
if an interrupt of a  
**SUSPENDED WORLD**  
takes place?



# INTERRUPT HANDLING


For now...

when an **interrupt** of a **suspended world** takes place, the interrupt will only be **served** when the **world is resumed** on a next scheduling point




**W#1: RTOS**  
no priority  
**W#2: TCP/IP**  
no priority  
**uTango**  
tick: 4t

- **Impracticable**, but “workable” at least to materialize uTango proof of concept
- In rough terms, the worst time in which an interruption may take to be served is given by the following equation:  
$$(worlds - 1) * tick$$




# INTERRUPT HANDLING



**W#1: RTOS**  
priority 1  
**W#2: TCP/IP**  
priority 2  
**uTango**  
tick: 4t

Currently, we are working ...

to allow an interrupt from a **suspended world** to be able to **preempt** uTango and **force** a new **context switch** to **resume** that world



- The preemption behavior will be **priority-driven**.
- A world with **higher criticality** can preempt a **lower criticality** world, but the opposite is not possible!

# The Numbers

Tests, experiments, and results (Arm Musca-B1)




# EXPERIMENTAL SETUP

## Hardware

- Arm Musca-B1 board (SSE-200)
- CPU0 @40MHz 2KB\_ICache
- CPU1 @40(-160)MHz 2KB\_ICache

## System Configs

- uTango v0.2
  - W1: LED Blinking
  - W2: Zephyr (control. a servo motor)
  - W3: Serial Terminal
  - W4: TCP/IP Stack



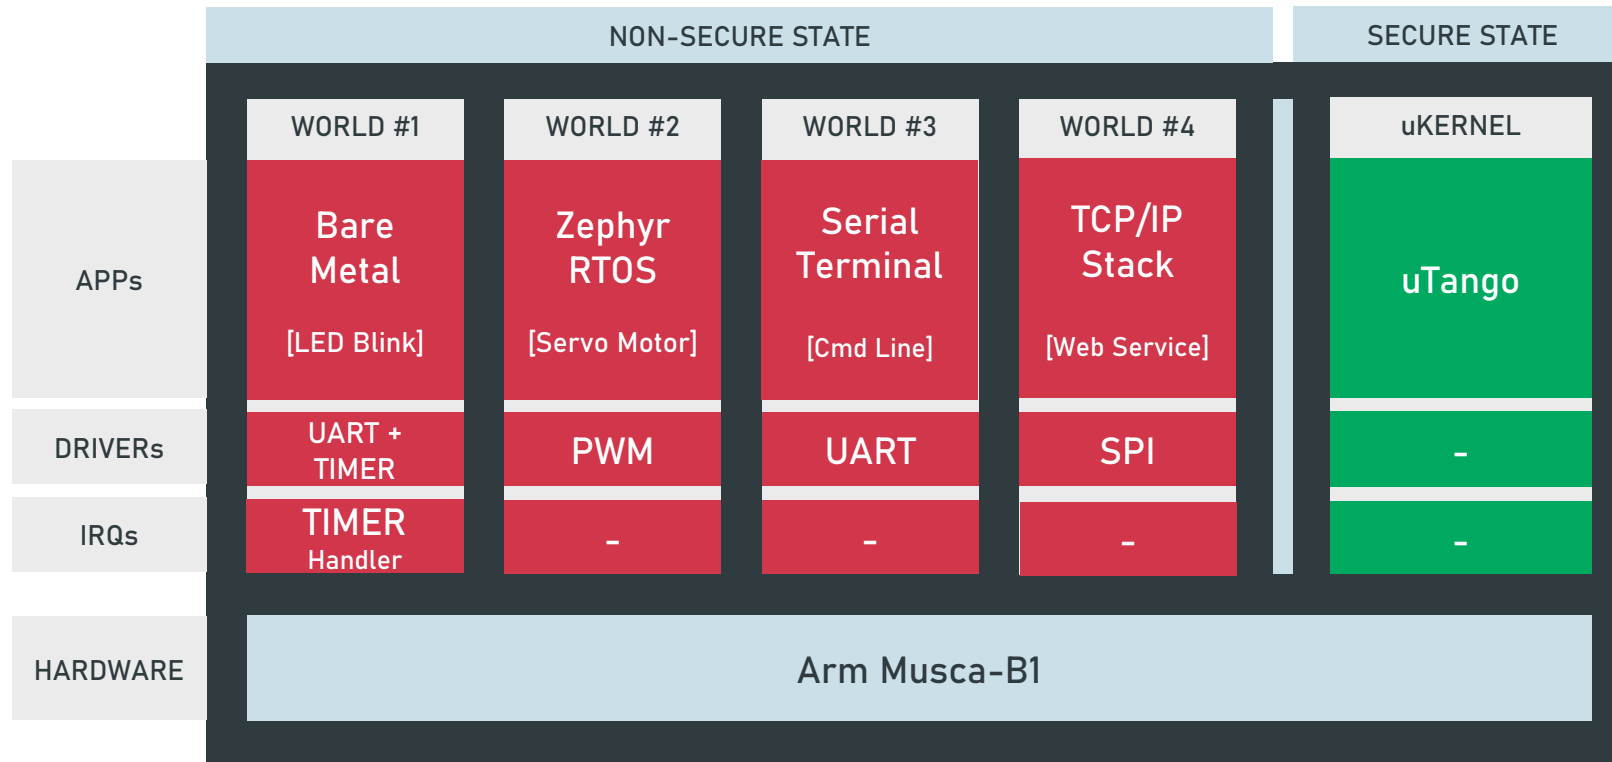
## Software

- Compiler:
  - GNU Arm Embedded Toolchain (9-2020-q2-update)
  - Optimization: -O0
- Measuring Unit:
  - Data Watchpoint and Trace (DWT) unit

## Metrics

- Performance Overhead
  - Time that uTango kernel spends while switching between worlds
- Trusted Computing Base (TCB)
  - Memory footprint (size)
  - Code Complexity (lines of code)

# UTANGO CONFIGURATION



# MEMORY FOOTPRINT

text	data	bss	total
4618	244	1056	5918 bytes

file extension	total code
.h	1248
.c	574
.S	181
.s	142


uTango relies on the **principle of minimal implementation**, striving to **reduce** the amount of **code** that executes at the **highest privilege level**, which is reflected in a TCB size of around **6 KB** with less than **900 lines of code**.

# PERFORMANCE OVERHEAD

Context Switch Time			
uTango Configuration	2 Worlds	3 Worlds	4 Worlds
Clock Cycles	110	110	110

We consider **performance overhead** as the **amount of time** that uTango kernel spends while switching between worlds, i.e. the **context/world switch time**.

The assessed results show that for all **uTango configurations**, the time taken during the context switch is precisely **110 clock cycles**.



# PERFORMANCE OVERHEAD

CPU Freq (MHz)	Context Switch Time (ns)		
40	2750	2750	2750
160	688	688	688

uTango Tick (ms)	Performance Overhead (%)		
1	0.07	0.07	0.07
5	0.01	0.01	0.01
10	0.007	0.007	0.007

160 MHz

For all system configurations (i.e. 2, 3, 4 worlds), the performance **degradation** (CPU frequency at 160 MHz and uTango running in a TCM memory) is **null** or **very small** ( $< \sim 0.1\%$ ).

# Conclusion

---

Takeaway Points

# TAKEAWAYS

1

We demonstrate why the **traditional dual-world** security model provided by TrustZone is becoming **impractical** and **inherently insecure** as the number and complexity of software integrated on embedded devices is increasing at a rapid pace.

2

We presented a **disruptive** approach that **solely** relies on the hooks of the elephant in the room: **Arm TrustZone-M**. We demonstrate how to take advantage of the same TrustZone hardware security primitives to **extend** the **dual-world model** to a **multi-world approach**.

3

We evaluate our system in a **real-world platform** and targeting a **real use-case scenario** to demonstrate that **security** doesn't always need to come at a **significant cost** (uTango achieve values of **performance overhead < 1%**).

We have presented the **1<sup>st</sup> Multi-World TEE** for modern **TrustZone-M** devices which are expected to be deployed on **billions of tomorrow's Arm MCUs**.



# ONGOING WORK

1

Wrapping up the **first release version** of uTango as soon as we finish the new interrupt handling mechanism.

[github.com/danielRep/utango](https://github.com/danielRep/utango)

2

Support for other **commercial platforms** is on the way (e.g. Nuvoton M2351, NXP LPC55S69).

3

Next major steps will focus on **multi-core support** and a **schedulability analysis** of uTango to improve **predictability** and extend the target to **mixed-criticality systems**.

# THANK YOU!

Daniel Oliveira (UMinho) | Sandro Pinto (UMinho)

**daniel.oliveira@dei.uminho.pt**

**LinkedIn** - <https://www.linkedin.com/in/danieljcoliv/>

**ResearchGate** - [https://www.researchgate.net/profile/Daniel\\_Oliveira12](https://www.researchgate.net/profile/Daniel_Oliveira12)

**Github** - <https://github.com/danielRep/>

**sandro.pinto@dei.uminho.pt**

**LinkedIn** - <https://www.linkedin.com/in/sandro-pinto-phd-40535455/>

**ResearchGate** - [https://www.researchgate.net/profile/Sandro\\_Pinto2](https://www.researchgate.net/profile/Sandro_Pinto2)

**Github** - <https://github.com/sandro2pinto/>

Q&A