



# CAPTURING ODAY EXPLOITS WITH PERFECTLY PLACED HARDWARE TRAPS

Cody Pierce . Matt Spisak . Kenneth Fitch


**ENDGAME.**

# INTRODUCTION


## EXPLOIT DETECTION IS A MOVING TARGET

- ▶ Exploitation is increasingly more sophisticated
- ▶ Creativity in exploitation is hard to plan for in the Security Development Lifecycle (SDL)
- ▶ A well financed attacker armed with 0days has the advantage


# EARLY PREVENTION TO MAINTAIN THE ADVANTAGE




# EARLY PREVENTION TO MAINTAIN THE ADVANTAGE




**INTRODUCING  
HARDWARE ASSISTANCE**






## PERFORMANCE MONITORING UNIT

- ▶ A special unit in microprocessor architectures to enable hardware level performance and system information. Often used to optimize hardware and software
- ▶ The PMU can be programmed to record dozens of different hardware “events”
- ▶ Traditionally reserved for developers and system architects



## PERFORMANCE MONITORING UNIT FOR SECURITY


- ▶ “Security Breaches as PMU Deviation: Detecting and Identifying Security Attacks Using Performance Counters”, Yuan et al., 2011
- ▶ “CFIMon: Detecting Violation of Control Flow Integrity using Performance Counters”, Xia et al., 2012
- ▶ “kBouncer: Efficient and Transparent ROP Mitigation”, Pappas, 2012
- ▶ “Transparent ROP Detection using CPU Performance Counters”, Li & Crouse, 2014



## BRANCH PREDICTION UNIT

- ▶ A unit in microprocessor architectures dedicated to improving the prediction of branch destinations to increase instruction pipeline efficiency
- ▶ Better branch prediction can have a large effect on processor performance
- ▶ Misprediction penalties can be many clock cycles due to flushing and filling the correct branch into the instruction pipeline
- ▶ Indirect branches can be common in C++ applications and predicting them is crucial to performance






**INTRODUCING  
CONTROL FLOW INTEGRITY**

# CONTROL FLOW INTEGRITY

- ▶ Enforcement of legitimate control flow in a program
- ▶ Traditionally done with compiler generated instrumentation
- ▶ Many different implementation of policy enforcement exist but the basic idea is to validate each indirect control flow transfer against a static list of trusted functions



CONTROL BRANCH



CFI POLICY ENFORCEMENT



DESTINATION



TERMINATE

## ALTERNATIVE CFI IMPLEMENTATIONS

- ▶ Control Flow Guard (CFG), Microsoft, 2014
- ▶ Control-flow Enforcement Technology (CET) , Intel, TBD?
- ▶ Return Address Protection/Indirect Control Transfer Protection (RAP/ICTP), PaX Team, 2015

## ALTERNATIVE CFI IMPLEMENTATIONS

- ▶ Control Flow Guard (CFG), Microsoft, 2014
- ▶ Control-flow Enforcement Technology (CET) , Intel, TBD?
- ▶ Return Address Protection/Indirect Control Transfer Protection (RAP/ICTP), PaX Team, 2015
- ▶ While these are **very strong** implementations they require recompilation, updated software/kernel/OS, or aren't cross platform

## SCOPING OUR RESEARCH TO FILL THE GAP

- ▶ No source code access
- ▶ Cross-Platform OS support
- ▶ 32 and 64 bit support
- ▶ No pre-processing of binaries or CFG reconstruction
- ▶ Not specific to a single bug-class or exploit technique such as Use-After-Free (UAF) or Return-oriented Programming (ROP)
- ▶ Overhead must be acceptable in benchmarks and subjective user experience

## REAL-WORLD VERIFICATION


- ▶ Approach must be verified using “real” exploits and “real” software
- ▶ Cyber Grand Challenge samples
- ▶ Research community PoCs
- ▶ Metasploit modules
- ▶ Exploit Kit samples including previous 0days
- ▶ Internally developed exploits

# CFI APPROACH

## HARDWARE-ASSISTED CONTROL FLOW INTEGRITY (HA-CFI)

- ▶ Hijacked indirect branches almost always mispredicted by BPU
- ▶ HA-CFI Approach:
  - ▶ Use Intel PMU to trap all mispredicted indirect branches
  - ▶ Requires setting counter to -1
  - ▶ Use ISR for CFI policy: validate **indirect branch** destinations in real-time
  - ▶ Initial prototype in Linux

# OUR INITIAL APPROACH




**BUT WONT ALL THOSE INTERRUPTS BE EXPENSIVE?**



# INDIRECT BRANCH CFI COMPARISON

	Source Code Required	Patching Required	Overhead	CFI Logic Frequency
<b>Binary Rewriting</b>	NO	YES	LOW	100% FOR PROTECTED CALLS
<b>Compiler Transformation</b>	YES	NO	LOW	100% FOR PROTECTED CALLS
<b>PMU-Assisted</b>	NO	NO	MEDIUM	ONLY WHEN MISPREDICTED 1%-20%

# PROGRAMMING THE PMU



## PROGRAMMING THE PMU


- ▶ Controlled by several Model Specific Registers (MSRs)
  - ▶ IA32\_PERF\_GLOBAL\_CTRL : global enable/disable of counters
  - ▶ IA32\_PERFEVTSELx : event to count, mode inclusion bits, interrupt bit
  - ▶ IA32\_PMCx : counter value
  - ▶ IA32\_PERF\_GLOBAL\_STATUS / IA32\_PERF\_GLOBAL\_OVF\_CTRL
    - ▶ counter overflow status and clear registers
- ▶ Additional references: Threads 2014 [Li et al], BH USA 2015 [Herath, Fogh]

## INDIRECT BRANCH - INTEL PMU EVENTS

EVENT NAME	UMASK	CODE	DESCRIPTION
BR_MISP_RETIRED.NEAR_CALL	0x02	0xC5	Direct and indirect mispredicted near call instructions retired
BR_MISP_EXEC.TAKEN_INDIRECT_NEAR_CALL	0xA0	0x89	Taken speculative and retired mispredicted indirect calls

- ▶ BR\_MISP\_RETIRED (PEBS) counts retired only, includes direct and indirect
- ▶ BR\_MISP\_EXEC includes speculative events == branches falsely labeled as mispredicted
- ▶ Opted to use BR\_MISP\_RETIRED.NEAR\_CALL since more precise and fewer Interrupts

# THE INTERRUPT SERVICE ROUTINE



# PMU TRAPS “ON PAPER”

IA32\_PMC0: 0xFFFFFFFF (-1) Event: 0x5102C5

PMC	INSTRUCTION
-1	0x1000: MOV rax, [rsi]
-1	0x1003: MOV rdi, [rax+0x78]
-1	0x1007: CALL rdi
foo:	
0	0xB890: MOV rax, rsp
0	0xB893: MOV [rax+0x20], r9d

RIP: 0xB890

PMU ISR

```
#ifdef WINDOWS
    ip = KTRAP_FRAME.RIP;
#else
    ip = pt_regs.rip;
//Apply CFI policy to RIP
do_cfi(ip);
```

PMI

# PMU TRAPS IN THE REAL WORLD

IA32\_PMC0: 0xFFFFFFFF (-1) Event: 0x5102C5

PMC	INSTRUCTION		
-1	0x1000:	MOV	rax, [rsi]
-1	0x1003:	MOV	rdi, [rax+0x78]
-1	0x1007:	CALL	rdi
foo:			
0	0xB890:	MOV	rax, rsp
0	0xB893:	MOV	[rax+0x20], r9d

RIP: 0xB893

Skid = 1 Instruction

PMI



## PMU TRAPS IN THE REAL WORLD

- ▶ Due to instruction skid after overflow, no guarantee saved IP is address of branch destination
- ▶ AMD docs state skid could be up to 72 instructions
- ▶ We found 1 instruction skid (or none) to be most common on Intel
- ▶ Need a more precise way to get branch target address on PMU overflow

## LBR TO THE RESCUE

- ▶ Intel Last Branch Record (LBR) can provide us precise branch addresses
- ▶ Configured and accessed via MSRs:
  - ▶ IA32\_DEBUGCTL : Enable/Disable bit, Freeze on PMI bit
  - ▶ LBR\_SELECT : filter types of branches
  - ▶ LASTBRANCH\_x\_FROM\_IP / LASTBRANCH\_x\_TO\_IP : LBR stack entries
  - ▶ LBR\_TOS : Offset that points to current top of LBR stack

# LBR TO THE RESCUE

Table 17-13. MSR\_LBR\_SELECT for Intel® microarchitecture code name Haswell

Bit Field	Bit Offset	Access	Description
CPL_EQ_0	0	R/W	When set, do not capture branches occurring in ring 0
CPL_NEQ_0	1	R/W	When set, do not capture branches occurring in ring >0
JCC	2	R/W	When set, do not capture conditional branches
NEAR_REL_CALL	3	R/W	When set, do not capture near relative calls
NEAR_IND_CALL	4	R/W	When set, do not capture near indirect calls
NEAR_RET	5	R/W	When set, do not capture near returns
NEAR_IND JMP	6	R/W	When set, do not capture near indirect jumps except near indirect calls and near returns
NEAR_REL JMP	7	R/W	When set, do not capture near relative jumps except near relative calls.
FAR_BRANCH	8	R/W	When set, do not capture far branches
EN_CALLSTACK <sup>1</sup>	9		Enable LBR stack to use LIFO filtering to capture Call stack profile
Reserved	63:10		Must be zero

LBR\_SELECT = 0x1ED ( Indirect Calls in ring > 0 )

easy first check in ISR



Table 17-8. MSR\_LASTBRANCH\_x\_FROM\_IP

Bit Field	Bit Offset	Access	Description
Data	47:0	R/O	This is the "branch from" address. See Section 17.4.8.1 for address format.
SIGN_EXt	62:48	R/O	Signed extension of bit 47 of this register.
MISPRED	63	R/O	When set, indicates either the target of the branch was mispredicted and/or the direction (taken/non-taken) was mispredicted; otherwise, the target branch was predicted.

# PMU TRAPS WITH LBR PRECISION

IA32\_PMC0: 0xFFFFFFFF (-1) Event: 0x5102C5

PMC	INSTRUCTION		
-1	0x1000: MOV	rax, [rsi]	
-1	0x1003: MOV	rdi, [rax+0x78]	
-1	0x1007: CALL	rdi	

```
foo:
0 0xB890: MOV rax, rsp
0 0xB893: MOV [rax+0x20], r9d
```

RIP: 0xB893

LBR FROM	LBR TO
0x59BC5CE5	0x75DFC3FB
0x1234	0x1000
0x8000000000001007	0xB890
0x59BC61C3	0x75DFC452

TOS →


PMI

PMU ISR

```
//Get LBR T0
tos = rdmsr(LBR_T0S);
lbr_to = rdmsr(LBR_T0 + tos)

//Apply CFI policy to LBR to
do_cfi(lbr_to);
```

# COLLECTING ALL THE DATA



# VALIDATING APPROACH W/ CYBER GRAND CHALLENGE SAMPLES

```
vRP@ubuntu:~$ miniperf -p 8491 -i 1 -e 0x51a089
Monitoring process: CROMU_00044 (8491)
```

```
80007F51FCF6DDBD 7F51FCF9D5F0
80007F51FCF9F62B 7F51FCF9E570
80007F51FCF9E734 7F51FCF9E7C0
80007F51FCF9E4D9 7F51FCF9CFC0
80007F51FCF9E69D 7F51FCF9D5D0
80007F51FCF6DDBD 7F51FCF9D5F0
80007F51FCF9F62B 7F51FCF9E570
80007F51FCF9E734 7F51FCF9E7C0
```

```
...
80007F51FCF9D69E 7F51FCF9E7C0
80007F51FCF9E4D9 7F51FCF9CFC0
8000000004032B2 41414141
```

```
403294: callq  <_ZN10CUserEntry20GetLastUnreadMessageEv>
403299: mov    %rax,-0x30(%rbp)
40329d: mov    -0x30(%rbp),%rax
4032a1: mov    (%rax),%rax
4032a4: add    $0x10,%rax
4032a8: mov    (%rax),%rax
4032ab: mov    -0x30(%rbp),%rdx
4032af: mov    %rdx,%rdi
4032b2: callq  *%rax
```

```
// Display last unread message
pCur = pUser->GetLastUnreadMessage();
printf( "From: @s\n", pCur->GetFrom().c_str() );
```

# INDIRECT BRANCH ANALYSIS - CVE-2014-0556


- ▶ Ubuntu 14.04.3 LTS x64
- Pepper Flash 14.0.0.177
- ▶ Moved to real-world Linux x64 POC, but missed hijack due to JMP
- ▶ Tweaked ActionScript POC from Chris Evans to generate additional data:
- ▶ 16 unique hijack points
- ▶ Call / JMP Analysis

ACTIONSCRIPT TRIGGER	HIJACKED CALL SITE
ByteArray.readBytes()	0x33D438: jmp rax
ByteArray.readBytes()	0x33D3BC: call qword ptr [rax]
ByteArray.readMultiByte()	0x33D1D6: call qword ptr [rax]
ByteArray.readMultiByte()	0x33D343: call qword ptr [rax+0x10]
ByteArray.readMultiByte()	0x33D1A7: call qword ptr [rax+0x10]
ByteArray.readMultiByte()	0x405358: call qword ptr [rax+0x8]
ByteArray.writeBytes()	0x33D4A8: jmp rax
ByteArray.writeBytes()	0x33D0E7: call qword ptr [rax+0x10]
ByteArray.writeMultiByte()	0x33CFFB: call qword ptr [rax+0x10]
ByteArray.writeMultiByte()	0x40805A: call qword ptr [rcx]
ByteArray.writeUTF()	0x33CE48: call qword ptr [rax]
ByteArray.writeUTFBytes()	0x33D0B8: call qword ptr [rax]
ByteArray.writeObject()	0x33D05E: call qword ptr [rax+0x10]
ByteArray.writeObject()	0x40477A: jmp rax
ByteArray.readObject()	0x33CDCE: call qword ptr [rax+0x10]
ByteArray.readObject()	0x40482B: jmp rax


## INDIRECT BRANCH - CALL VS JMP

● CALL


● JMP




`jscript9.dll`



`Flash.ocx`



`libpepflashplayer.so`




`libxul.so`

- ▶ Hijackable indirect JMP slightly more common in Linux binaries
- ▶ Indirect JMPs often used for switch statements
- ▶ For this talk we will focus exclusively on indirect CALLs


# WHAT IS A VALID INDIRECT BRANCH?

Firefox  
**167,755,264**  
Branches collected

Dromaeo JavaScript Benchmark  
[dromaeo.com](http://dromaeo.com)



# OUR FINAL DESIGN



## WHITELIST GENERATION

- ▶ Only after we were able to validate all 160M branches from ELF did we even explore real-time whitelist query
- ▶ Generate list on each image load in protected processes
- ▶ Overall approach is same on ELF and PE:
  - ▶ Find all code pointer addresses present in loaded image
  - ▶ Code pointer considered if relative or absolute address points to `.text`
  - ▶ Primarily focus on Exports, Relocations, and “Callbacks”

# IMPLEMENTATION CHALLENGES

## KEY CHALLENGES

- ▶ Receiving PMU Interrupts
- ▶ Clearing PMU Interrupts
- ▶ Thread Tracking

## RECEIVING PMU INTERRUPTS ON WINDOWS

- ▶ Modifying the Interrupt Descriptor Table (IDT) for the PMU interrupt will not work for x64 due to PatchGuard
- ▶ While investigating how Windows handles PMI, we discovered a non-exported kernel routine in hal.dll
  - ▶ **HalpSetSystemInformation()**
    - ▶ **InformationClass** of **HalProfileSourceInterruptHandler**
    - ▶ Reachable through **HalDispatchTable** export

# RECEIVING PMU INTERRUPTS ON WINDOWS

```
NT_STATUS _HalpSetSystemInformation(HAL_SET_INFORMATION_CLASS InformationClass, ULONG BufferSize,
PVOID *Buffer) {
// ...

if(InformationClass == HalProfileSourceInterruptHandler) {
    if(BufferSize != 4)
        return STATUS_INFO_LENGTH_MISMATCH;

    if(HalpFeatureBits & 1 == 0)
        return STATUS_INVALID_DEVICE_REQUEST;

    if(ProfilingProcessId == 0) {
        _HalpPerfInterruptHandler = Buffer[0];
        if(Buffer[0] != NULL)
            ProfilingProcessId = PsGetCurrentProcessId();

    } else {
        if(PsGetCurrentProcessId() != ProfilingProcessId)
            return STATUS_INVALID_DEVICE_REQUEST;
        _HalpPerfInterruptHandler = Buffer[0];
        ProfilingProcessId = (Buffer[0] ? ProfilingProcessId : 0);
    }

    return STATUS_SUCCESS;
}

// ...
}
```

# RECEIVING PMU INTERRUPTS ON WINDOWS

- ▶ Pass in the interrupt handler function and it will be called when a PMI occurs

```
NTSTATUS status;
PVOID buffer[1];

buffer[0] = profileSourceInterruptHandler;
status = HalpSetSystemInformation(HalProfileSourceInterruptHandler,
                                   sizeof(PVOID),
                                   buffer);
```

- ▶ Calling (from the same process) with a NULL pointer deregisters the handler

```
NTSTATUS status;
PVOID buffer[1];

buffer[0] = NULL;
status = HalpSetSystemInformation(HalProfileSourceInterruptHandler,
                                   sizeof(PVOID),
                                   buffer);
```

## CLEARING PMU INTERRUPTS ON WINDOWS

- ▶ Another issue encountered involved unmasking PMU interrupts from the handler
- ▶ PMU interrupts are delivered by the APIC
- ▶ In order to acknowledge an interrupt has been handled and to receive future interrupts, a register in the APIC needs to be written
- ▶ How this is accomplished depends on the APIC interface used, which differs between Windows versions

# CLEARING PMU INTERRUPTS ON WINDOWS

## xAPIC

- ▶ Existed since Pentium 4
- ▶ Windows 7
- ▶ APIC Registers are accessed through mapped physical memory
- ▶ Register access accomplished using physical memory mapped into kernel virtual memory via `MmMapIoSpace`

## x2APIC

- ▶ Introduced in Nehalem microarch
- ▶ Windows 8/8.1
- ▶ APIC Registers are accessed via MSRs
- ▶ Interface can be accessed with a single `_writemsr` intrinsic
  - `_writemsr(LVT_x2APIC_PMI, 0xFE)`

## RECEIVING PMU INTERRUPTS ON LINUX

- ▶ Setting this up on Linux is even simpler
- ▶ Register for a Non-Maskable Interrupt (NMI) handler

```
register_nmi_handler(NMI_LOCAL,  
                     our_nmi_handler,  
                     NMI_FLAG_FIRST,  
                     "hacfi_pmi");  
  
unregister_nmi_handler(NMI_LOCAL, "hacfi_pmi");
```

## THREAD TRACKING

- ▶ We don't want to monitor the entire system
- ▶ Monitoring can be restricted to a few "high threat" executables
- ▶ The PMU doesn't know anything about thread or process context

## THREAD TRACKING ON WINDOWS

- ▶ Not so straightforward
- ▶ Windows has no (explicit) mechanism for executing arbitrary code at thread context switches
- ▶ Without some sort of callback when a thread quantum starts execution, we don't know when to turn on the PMU counters
- ▶ This is a problem


## ASYNCHRONOUS PROCEDURE CALLS TO THE RESCUE

- ▶ “When an APC is queued to a thread, the system issues a software interrupt. The next time the thread is scheduled, it will run the APC function.” - Microsoft
- ▶ Perfect! We could just use APCs to get callbacks, and re-queue a new one whenever we finish the previous
- ▶ Not quite that simple, since we don’t track all threads and don’t know when a monitored quantum has ended
- ▶ Also, scheduling an APC for the current thread, from an APC handler, leads to an endless APC loop due to the software interrupt

## OUR APC SOLUTION

1. Schedule a kernel APC for every thread we want to track
2. Configure PMU to trap all mispredicted branches
3. When we see an interrupt for the wrong thread, schedule a new APC for the previous thread on the processor (or all tracked threads that don't have one currently queued)
4. Repeat

## OUR APC SOLUTION



# THREAD TRACKING ON LINUX

- ▶ Very straightforward
- ▶ `preempt_notifier_init` gives us a simple callback registration for when a thread is preempted

```
static struct preempt_notifier notifier;
static struct preempt_ops hacfi_preempt_ops = {
    .sched_in = hacfi_notifier_sched_in,
    .sched_out = hacfi_notifier_sched_out
};

static void hacfi_notifier_sched_in(struct preempt_notifier *notifier, int cpu);

static void hacfi_notifier_sched_out(struct preempt_notifier *notifier,
                                     struct task_struct *next);

preempt_notifier_init(&notifier, &hacfi_preempt_ops);
```



zakaazis@gmail.com

# RESULTS

## ANALYSIS OF RESULTS



- ▶ Performance Overhead
- ▶ Exploit Detection efficacy testing

## PERFORMANCE

- ▶ We expect there to be a lot of mispredicted branches leading to excessive interrupts
- ▶ There is also a minor fixed overhead for each quantum from the APC
- ▶ We need to test and see if this is feasible...

# HOW MANY PMU INTERRUPTS ARE WE TALKING ABOUT?

MISPREDICTED INDIRECT CALLS DURING OCTANE



# PERFORMANCE OVERHEAD

Benchmark	Baseline	HA-CFI	EMET
<b>PassMark PerformanceTest</b>	score: 940	9% score: 855	3% score: 910
<b>Dromaeo JavaScript w/ Internet Explorer</b>	325 runs/s	22% 253 runs/s	32% 220 runs/s

\*TESTING PERFORMED ON AN INTEL HASWELL CPU

## EXPLOIT DETECTION TESTING

- ▶ We needed exploits to test....
- ▶ We wanted exploits of recent CVEs for Adobe Flash, Internet Explorer, and Microsoft Office
- ▶ To Metasploit!



# EXPLOIT DETECTION TESTING - METASPLOIT

VULNERABILITY	TARGET	DETECTION RATE
CVE-2014-0497	Flash Player 11.7.700.202	100%
CVE-2014-0515	Flash Player 11.7.700.275	100%
CVE-2014-0556	Flash Player 14.0.0.145	100%
CVE-2014-0569	Flash Player 15.0.0.167	100%
CVE-2014-8440	Flash Player 15.0.0.189	100%
CVE-2015-0311	Flash Player 16.0.0.235	100%
CVE-2015-0313	Flash Player 16.0.0.296	100%
CVE-2015-0359	Flash Player 17.0.0.134	100%
CVE-2015-3090	Flash Player 17.0.0.169	90%
CVE-2015-3105	Flash Player 17.0.0.188	100%
CVE-2015-3113	Flash Player 18.0.0.160	100%
CVE-2015-5119	Flash Player 15.0.0.189	100%
CVE-2015-5122	Flash Player 18.0.0.194	100%
CVE-2014-1761	Microsoft Word 2010	100%

## EXPLOIT DETECTION TESTING

- ▶ Metasploit results were great, but what about the bad guys?
- ▶ The techniques used in an exploit matter as much or more than the actual vulnerability itself
- ▶ We don't think Metasploit is a great testbed for HA-CFI, due to lack of diversity in exploitation approach
- ▶ So we turned to VirusTotal and Exploit Kit samples collected in the wild



## EXPLOIT DETECTION TESTING - VIRUSTOTAL

- ▶ VirusTotal enabled us to test on real-world malware including previously 0day exploits
- ▶ Decided that samples from some of the more popular exploit kits would be a good basis for testing
- ▶ Using actual exploits from 'the wild' should provide a good sample of exploitation techniques
- ▶ We chose 48 unique samples for our testbed

7  
Exploit Kits

48  
Samples

20  
Unique CVEs

## EXPLOIT DETECTION TESTING - VIRUSTOTAL

- ▶ We analyzed each sample and bucketed them into three separate categories according to exploitation technique
- ▶ **ROP Technique** - Uses standard Return Oriented Programming techniques
- ▶ **ROPless Technique A** - Flash exploitation technique invoking a wrapper routine of VirtualProtect to make shellcode executable
- ▶ **ROPless Technique B** - Similar to A, but via hijacking Method.apply() in ActionScript to find and invoke VirtualProtect directly (Vitaly Toropov)

## EXPLOIT KIT DETECTION - HA-CFI VS EMET

CODE EXECUTION TECHNIQUE	# SAMPLES	HA-CFI DETECTION RATE	EMET DETECTION RATE
ROP	37	95%	100%
ROPless Technique A	1	100%	0%
ROPless Technique B	10	100%	0%

## EXPLOIT KIT DETECTION - BY BUG CLASS

BUG CLASS	# CVE'S	# SAMPLES	HA-CFI DETECTION RATE
Out-of-bounds Write	3	6	83.3%
Buffer Overflow	3	6	83.3%
Integer Overflow	2	6	100%
Use-After-Free	4	14	100%
Double Free	2	4	100%
Type Confusion	3	6	100%
Race Condition	1	4	100%
Uninitialized Memory	1	1	100%

# CASE STUDIES

## CLASSIC ROP TECHNIQUE

- ▶ CVE-2015-2419 : Double-free in jscript9 (MS15-065)


- ▶ Magnitude EK Sample

HA-CFI blocks the initial hijack.

RIP in ISR = jscript9 + 3BE32



EMET



```
jscript9 + A7541 ; JavascriptOperators::OP_SetElementI  
call edi  
  
jscript9 + 3BE32  
xchg eax,esp ; stack pivot gadget  
retn  
  
jscript9 + 4B0B5  
mov [ecx+0xC],ax ; CoE help  
retn  
  
jscript9 + 3BE33  
retn  
  
kernel32 + 42C15  
VirtualProtectStub ; mark shellcode +X
```




HA-CFI

## ROPLESS TECHNIQUE #1

- ▶ **CVE-2014-0515** : Heap overflow in Adobe Flash (patch in 13.0.0.206)
- ▶ Found in many Exploit kits and watering hole attacks
- ▶ ROPless technique re-uses VirtualProtect wrapper function in Flash image
- ▶ 2 control flow hijacks: one to VP wrapper, second one to shellcode
- ▶ Bypasses anti-ROP checks since VP invoked somewhat legitimately

## ROPLESS TECHNIQUE #1

- ▶ CVE-2014-0515 : Heap overflow in Adobe Flash (patch in 13.0.0.206)



HA-CFI



## FUTURE WORK

- ▶ Hypervisor support to enable hardware features in virtual machines
  - ▶ Last Branch Record (LBR) is not fully supported in popular hypervisors
  - ▶ Performance Monitoring Interrupts (PMI) on overflow is supported in many hypervisors
  - ▶ We wrote a patch for Xen to enable HA-CFI but it crashes randomly, anyone want to help?
  - ▶ Just-In-Time code pages are hard to validate with our current whitelist approach

# CONCLUSION

## EXPLOIT DEFENSE

- ▶ Exploit defense needs to detect and prevent exploitation at the earliest phase
- ▶ Compile-time solutions are powerful, but there is room for run-time defense too
- ▶ Defenses focused exclusively on techniques such as Return-oriented Programming can be easily circumvented as new methods get adopted
- ▶ Exploits will continue to “look normal” to bypass prevention checks

## HARDWARE ASSISTED CONTROL FLOW INTEGRITY

- ▶ CFI is a powerful first step in ensuring only trusted code paths can be executed
- ▶ Many vulnerabilities must hijack control-flow to achieve code execution
- ▶ Hardware can be leveraged for strong CFI policy enforcement of applications at run-time
- ▶ Many new hardware features are emerging that can be used for exploit defense

## CFI ENFORCEMENT

- ▶ CFI policies can be more complex
- ▶ Powerful features of the PMU interrupt on branches for prevention are the high IRQL and complete access to context information
- ▶ We have more ideas in the works to detect additional events and apply policies to detect abnormal read, writes, and cases where attackers stay within our whitelist

## SPECIAL THANKS

- ▶ Aaron Lamb, Endgame
- ▶ Gabriel Landau, Endgame
- ▶ Andrea Limbago, Endgame
- ▶ Kafeine, [malware.dontneedcoffee.com](http://malware.dontneedcoffee.com)
- ▶ Fellow researchers and vendors working on exploit defense

# QUESTIONS?