



USA 2018

AUGUST 4-9, 2018
MANDALAY BAY / LAS VEGAS

Back To The Future: A Radical Insecure Design of KVM on ARM

Baibhav Singh
Rahul Kashyap



#BHUSA / BLACKHAT EVENTS


Introduction

- Samsung Research America
 - Actively working on ARM system security
- *The opinion present are my own and not necessarily represent the opinion of my employer*

Samsung **Kn&X**

Background

- Found this issue while doing some hypervisors research
- On 25 Jan 2018, reported to Red Hat Product Security team
Secalert@redhat.com provided detail report with exploit code
- Multiple mail exchange to help them understand the problem
- **Still is not fixed**
- Decided to submit a BH paper (Thanks to Rahul, Michael Grace)
- Thanks to BH for providing the platform




Exception Vector Table

B. Sing - SRA

5

#BHUSA



Address	Exception type	Description
VBAR_ELn + 0x000	Synchronous	Current EL with SP0
+ 0x080	IRQ/vIRQ	
+ 0x100	FIQ/vFIQ	
+ 0x180	SError/vSError	
+ 0x200	Synchronous	Current EL with SPx
+ 0x280	IRQ/vIRQ	
+ 0x300	FIQ/vFIQ	
+ 0x380	SError/vSError	
+ 0x400	Synchronous	Lower EL using AArch64
+ 0x480	IRQ/vIRQ	
+ 0x500	FIQ/vFIQ	
+ 0x580	SError/vSError	
+ 0x600	Synchronous	Lower EL using AArch32
+ 0x680	IRQ/vIRQ	
+ 0x700	FIQ/vFIQ	
+ 0x780	SError/vSError	


ARM Hypervisor executes in EL2

TYPE 1

- Bare Metal Hypervisor
- Host is considered as VM
- XEN is type 1 and runs in EL2


TYPE 2

- Hypervisor is extension of Host Kernel
- Host is not considered as VM
- KVM is type 2 and runs partial in EL2



Linux and KVM

1. Updates
VBAR_EL2 to
Install STUB
Vector Table as
EL2 Exception
Vector Table
2. Switches back to
EL1



```
*****  
* Gets SPSR for BL33 entry  
*****  
  
uint32_t rpi3_get_spsr_for_bl33_entry(void)  
{  
#.if RPI3_BL33 IN AARCH32  
    INFO("BL33 will boot in Non-secure AArch32 Hypervisor mode\n");  
    return SPSR_MODE32(MODE32_hyp, SPSR_T_ARM, SPSR_E_LITTLE,  
                      DISABLE_ALL_EXCEPTIONS);  
#else  
    return SPSR_64(MODE_EL2, MODE_SP_ELX, DISABLE_ALL_EXCEPTIONS);  
#endif  
}
```

Stub Vector Table

```
ENTRY( hyp stub vectors)
    ventry el2_sync_invalid      // Synchronous EL2t
    ventry el2_irq_invalid       // IRQ EL2t
    ventry el2_fiq_invalid       // FIQ EL2t
    ventry el2_error_invalid     // Error EL2t

    ventry el2_sync_invalid      // Synchronous EL2h
    ventry el2_irq_invalid       // IRQ EL2h
    ventry el2_fiq_invalid       // FIQ EL2h
    ventry el2_error_invalid     // Error EL2h

    ventry el1_sync              // Red circle highlights this entry
    ventry el1_irq_invalid       // IRQ 64-bit EL1
    ventry el1_fiq_invalid       // FIQ 64-bit EL1
    ventry el1_error_invalid     // Error 64-bit EL1

    ventry el1_sync_invalid      // Synchronous 32-bit EL1
    ventry el1_irq_invalid       // IRQ 32-bit EL1
    ventry el1_fiq_invalid       // FIQ 32-bit EL1
    ventry el1_error_invalid     // Error 32-bit EL1

ENDPROC( hyp stub vectors)
```

VBAR_EL2

Handler in case HVC originating
from 64 bit Kernel

Filename : linux\arch\arm64\kernel\hyp-stub.s

Stub HVC Handler


```
el1_sync:  
    cmp x0, #HVC_SET_VECTORS  
    b.ne 2f  
    msr vbar_el2, x1  
    b 9f  
  
2:   cmp x0, #HVC_SOFT_RESTART  
    b.ne 3f  
    mov x0, x2  
    mov x2, x4  
    mov x4, x1  
    mov x1, x3  
    br x4          // no return  
  
3:   cmp x0, #HVC_RESET_VECTORS  
    beq 9f          // Nothing to reset!  
  
    /* Someone called kvm_call_hyp() against the hyp-stub... */  
    ldr x0, =HVC_STUB_ERR  
    eret  
  
9:   mov x0, xzr  
    eret  
ENDPROC(el1_sync)
```

Checks if register X0 == 0

Updates VBAR_EL2

Filename : linux\arch\arm64\kernel\hyp-stub.s

State Diagram 1



KVM HVC Handler

```

ell_sync:                      // Guest trapped into EL2

    mrs x0, esr_el2
    lsr x0, x0, #ESR_ELX_EC_SHIFT
    cmp x0, #ESR_ELx_EC_HVC64
    ccmp x0, #ESR_ELx_EC_HVC32, #4, ne
    b.ne    ell_trap

    mrs x1, vttbr_el2      // If vttbr is valid, the guest
    cbnz x1, ell_hvc_guest // called HVC

    /* Here, we're pretty sure the host called HVC. */
    ldp x0, x1, [sp], #16

    /* Check for a stub HVC call */
    cmp x0, #HVC_STUB_HCALL_NR
    b.hs    if

    /*
     * Compute the idmap address of __kvm_handle_stub_hvc and
     * jump there. Since we use kimage_voffset, do not use the
     * HYP VA for __kvm_handle_stub_hvc, but the kernel VA instead
     * (by loading it from the constant pool).
     *
     * Preserve x0-x4, which may contain stub parameters.
     */
    ldr x5, =__kvm_handle_stub_hvc
    ldr_l x6, kimage_voffset

    /* x5 = pa(x5) */
    sub x5, x5, x6
    br x5

```

}

}

}

Checks if HVC is from Host

Checks if register x0 < 3

Calls __kvm_handle_stub_hvc

Filenname : linux\arch\arm64\kernel\hyp-entry.s

KVM HVC Handler

B. Singh - SRA

12

#BHUSA



Filename :
linux\arch\arm64\kvm\hyp\hyp-entry.s

```
ENTRY(kvm_handle_stub_hvc)
    cmp x0, #HVC_SOFT_RESTART
    b.ne    if

    /* This is where we're about to jump, staying at EL2 */
    msr elr_el2, x1
    mov x0, #(PSR_F_BIT | PSR_I_BIT | PSR_A_BIT | PSR_D_BIT | PSR_MODE_EL2h)
    msr spsr_el2, x0

    /* Shuffle the arguments, and don't come back */
    mov x0, x2
    mov x1, x3
    mov x2, x4
    b    reset

1:   cmp x0, #HVC_RESET_VECTORS
    b.ne    if
reset:
    /*
     * Reset kvm back to the hyp stub. Do not clobber x0-x4 in
     * case we coming via HVC_SOFT_RESTART.
     */
    mrs x5, sctlr_el2
    ldr x6, =SCTRLR_ELx_FLAGS
    bic x5, x5, x6      // Clear SCTL_M and etc
    pre_disable_mmuvirtualisation
    msr sctlr_el2, x5
    isb


    /* Install stub vectors */
    adr_l  x5, __hyp_stub_vectors
    msr vbar_el2, x5
    mov x0, xzr
    eret
```

Checks register X0 == 2



Disables EL2 MMU

Back to Stub Vector

State Diagram 2



Exploitation



"You're on the host, and you can break KVM by inserting a rogue kernel module. Big deal. You can also blast the page tables, corrupt file systems, and make sure the box is on fire"

KVM Threat Model Assumption : HOST.EL1 == EL2

Security hole in privilege isolation boundary

Host kernel compromise is End of the Game !

Real World : EL1 != EL2

For Attacker Beginning of a New Game.... ☺

Attacker can exploit this issue to gain more privilege and will migrate to EL2

- Launch attack from isolated and unreachable memory.
- Can configure EL2 to get code execution from various different places
- A generic way to bypass security implemented in the kernel (LKRG),by escaping to EL2
- Attack the secure monitoring running in hypervisor mode
- Gives attacker opportunity for Blue Pill for KVM on ARM

Juicy target for attacker to perform highly sophisticated and stealthy attack

Impact

Potentially bigger impact for mobile and IoT

- Most of them are ARM based
- Chances are high that it will boot in EL2
- Single Kernel Device (More Attack Surface)

Affected Architecture: ARM v7-A and ARM v8-A with hardware virtualization

Mitigation

My two cents... You're likely vulnerable to this attack.
Patch the system by making sure Linux starts in EL1

Reference

- https://static.docs.arm.com/ddi0487/ca/DDI0487C_a_armv8_arm.pdf
- <https://developer.arm.com/products/architecture/arm-profile/docs/100942/latest/hypervisor-software>
- <https://dl.acm.org/citation.cfm?id=2541946>
- <http://www.cs.columbia.edu/~cdall/pubs/atc17-dall.pdf>
- <http://www.cs.columbia.edu/~cdall/pubs/sosp2017-neve.pdf>
- <https://lwn.net/Articles/557132/>
- http://lia.disi.unibo.it/Courses/som1516/materiale/VOSYS_BolognaKVMARM_2_12_2015.pdf