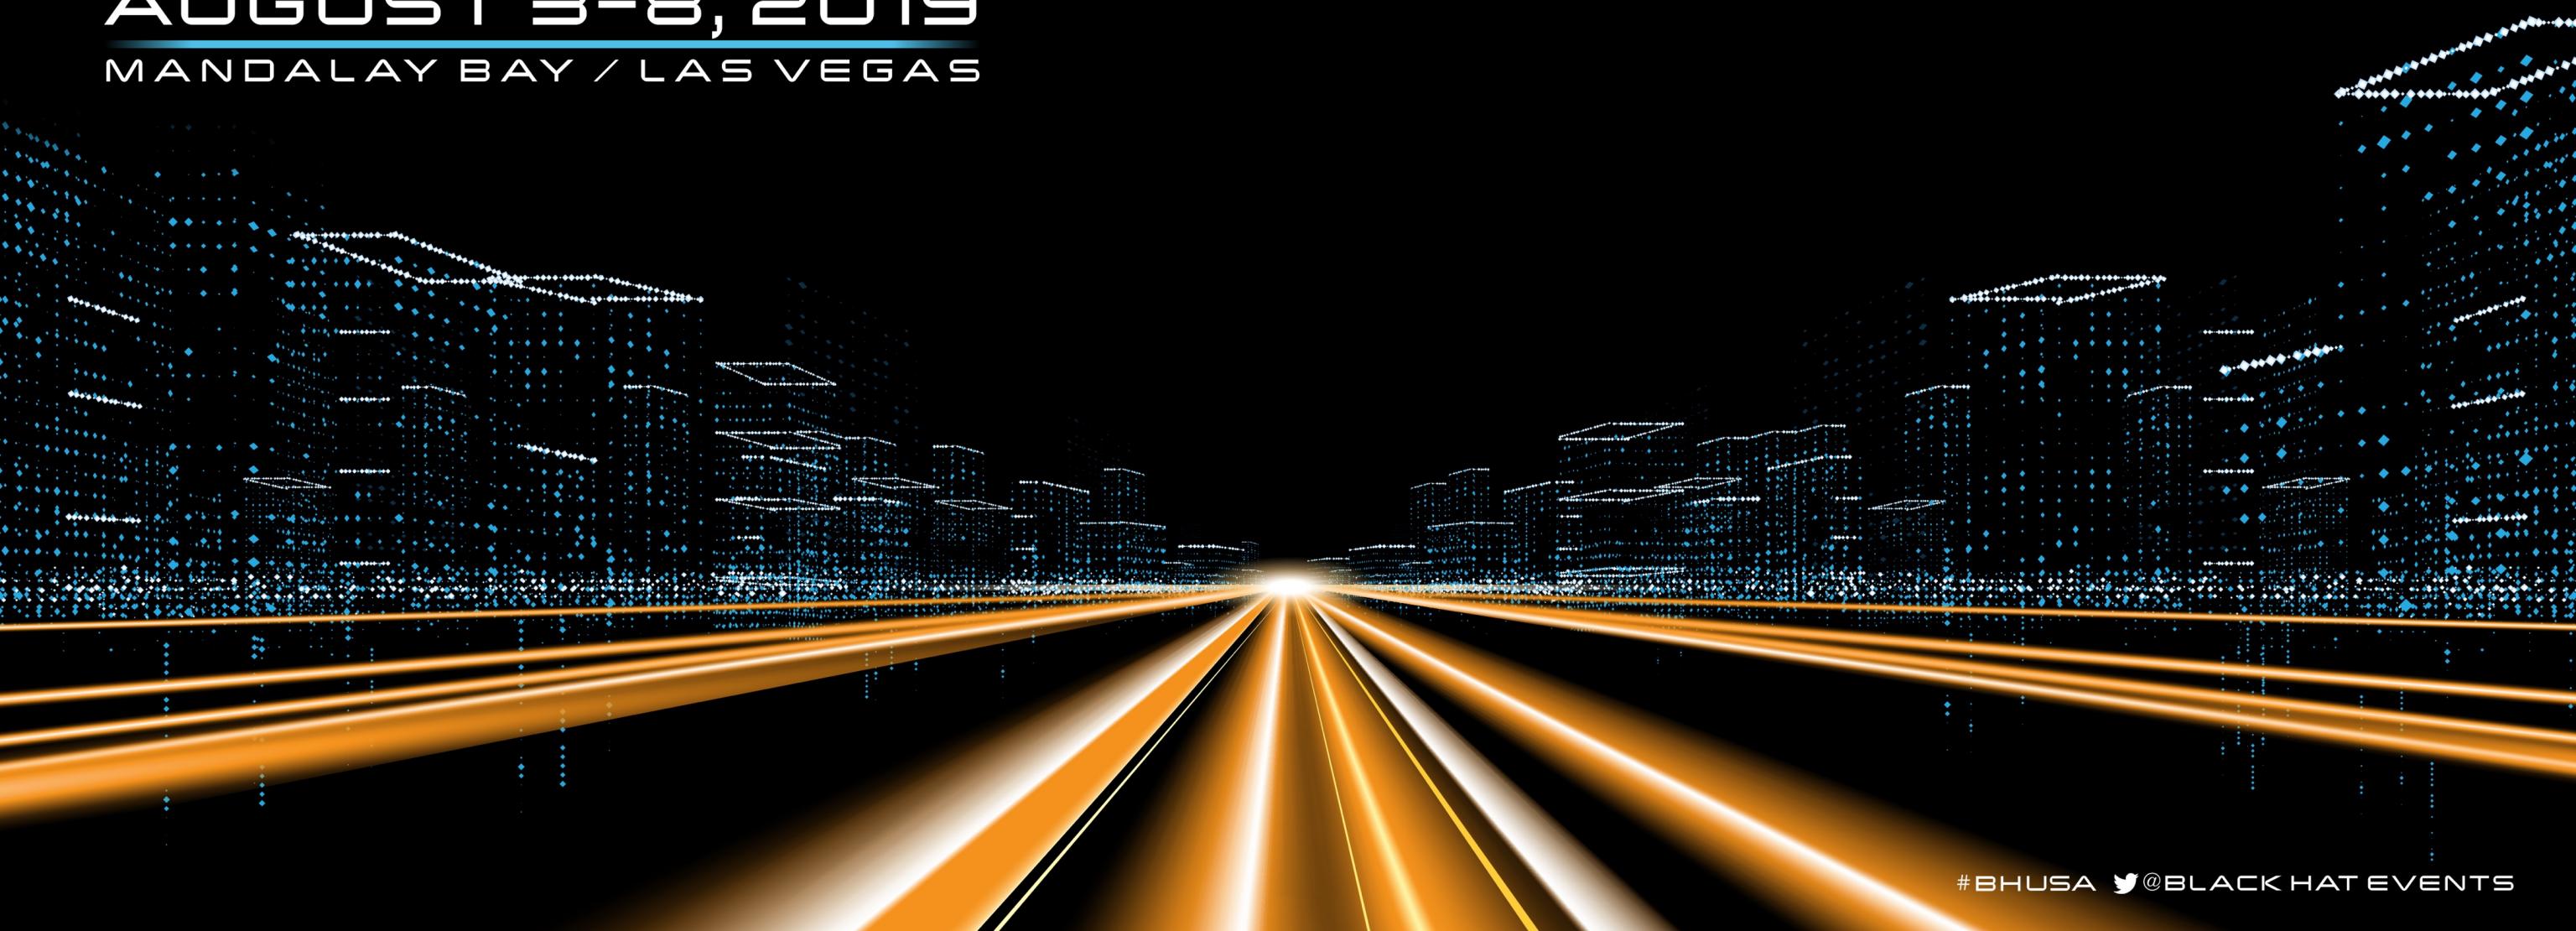




AUGUST 3-8, 2019

MANDALAY BAY / LAS VEGAS



PeriScope: An Effective Probing and Fuzzing Framework for the Hardware-OS Boundary

Dokyung Song, Felicitas Hetzelt, Dipanjan Das, Chad Spensky, Yeoul Na, Stijn Volckaert,
Giovanni Vigna, Christopher Kruegel, Jean-Pierre Seifert, Michael Franz



Remote Compromise of Peripheral Chips



BIZ & IT —

Broadcom chip bug opened 1 billion phones to a Wi-Fi-hopping worm attack

Wi-Fi chips used in iPhones and Android may revive worm attacks of old.

DAN GOODIN - 7/28/2017, 12:35 PM

Worm Attack via Arbitrary Code Execution on Broadcom Wi-Fi chip (aka BroadPwn)
– Nitay Artenstein at Black Hat USA 2017

Remote Compromise of Peripheral Chips



ZERO DAY
INITIATIVE

(Pwn2own) Samsung Galaxy S8 Shannon GPRS Stack-based Buffer Overflow Remote Code Execution Vulnerability

VULNERABILITY DETAILS

This vulnerability allows remote attackers to execute arbitrary code on vulnerable installations of Samsung Galaxy S8. User interaction is required to exploit this vulnerability in that the target must have their cellular radios enabled.

The specific flaw exists within the handling of IPCP headers. The issue results from the lack of proper validation of the length of user-supplied data prior to copying it to a fixed-length, stack-based buffer. An attacker can leverage this vulnerability to execute code under the context of the baseband processor.

Vulnerability used to achieve Arbitrary Code Execution on Shannon Baseband Chip
– Amat Cama at Mobile Pwn2Own

Remote Compromise of Peripheral Chips



ZDNet SEARCH CLOUD AI INNOVATION SECURITY MORE ▾ NEWSLETTERS ALL WRITERS LOG IN

iPhone, Android hit by Broadcom Wi-Fi chip bugs: Now Apple, Google plug flaws

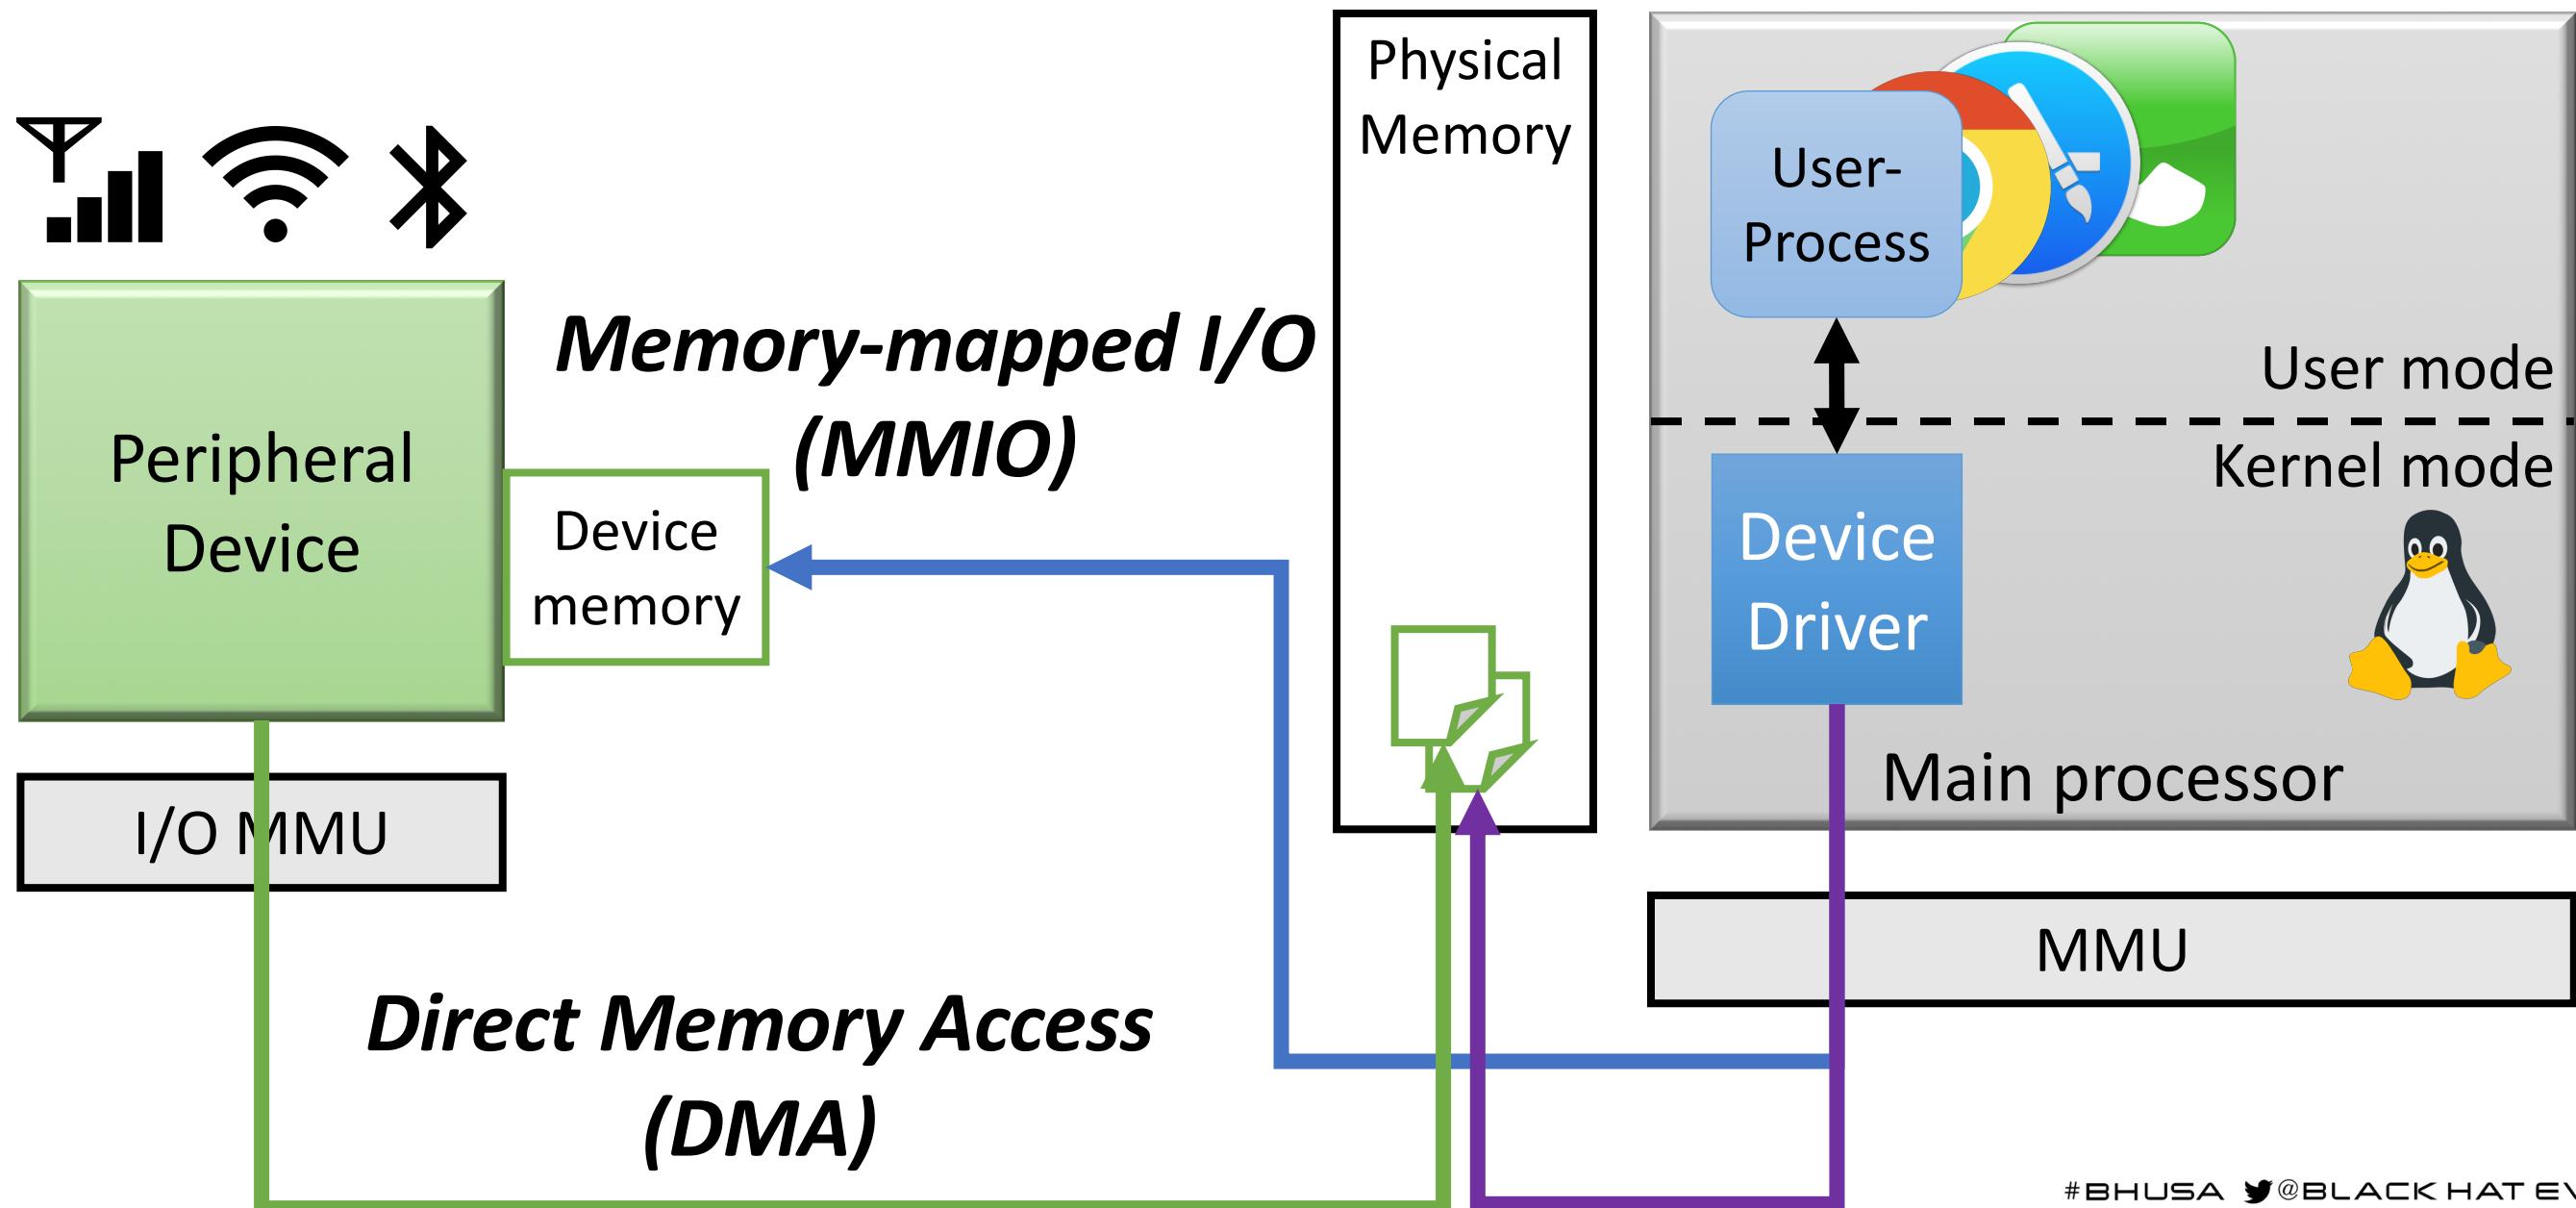
Google's Project Zero shows how attackers could target increasingly powerful Wi-Fi chips on phones as low-hanging fruit.



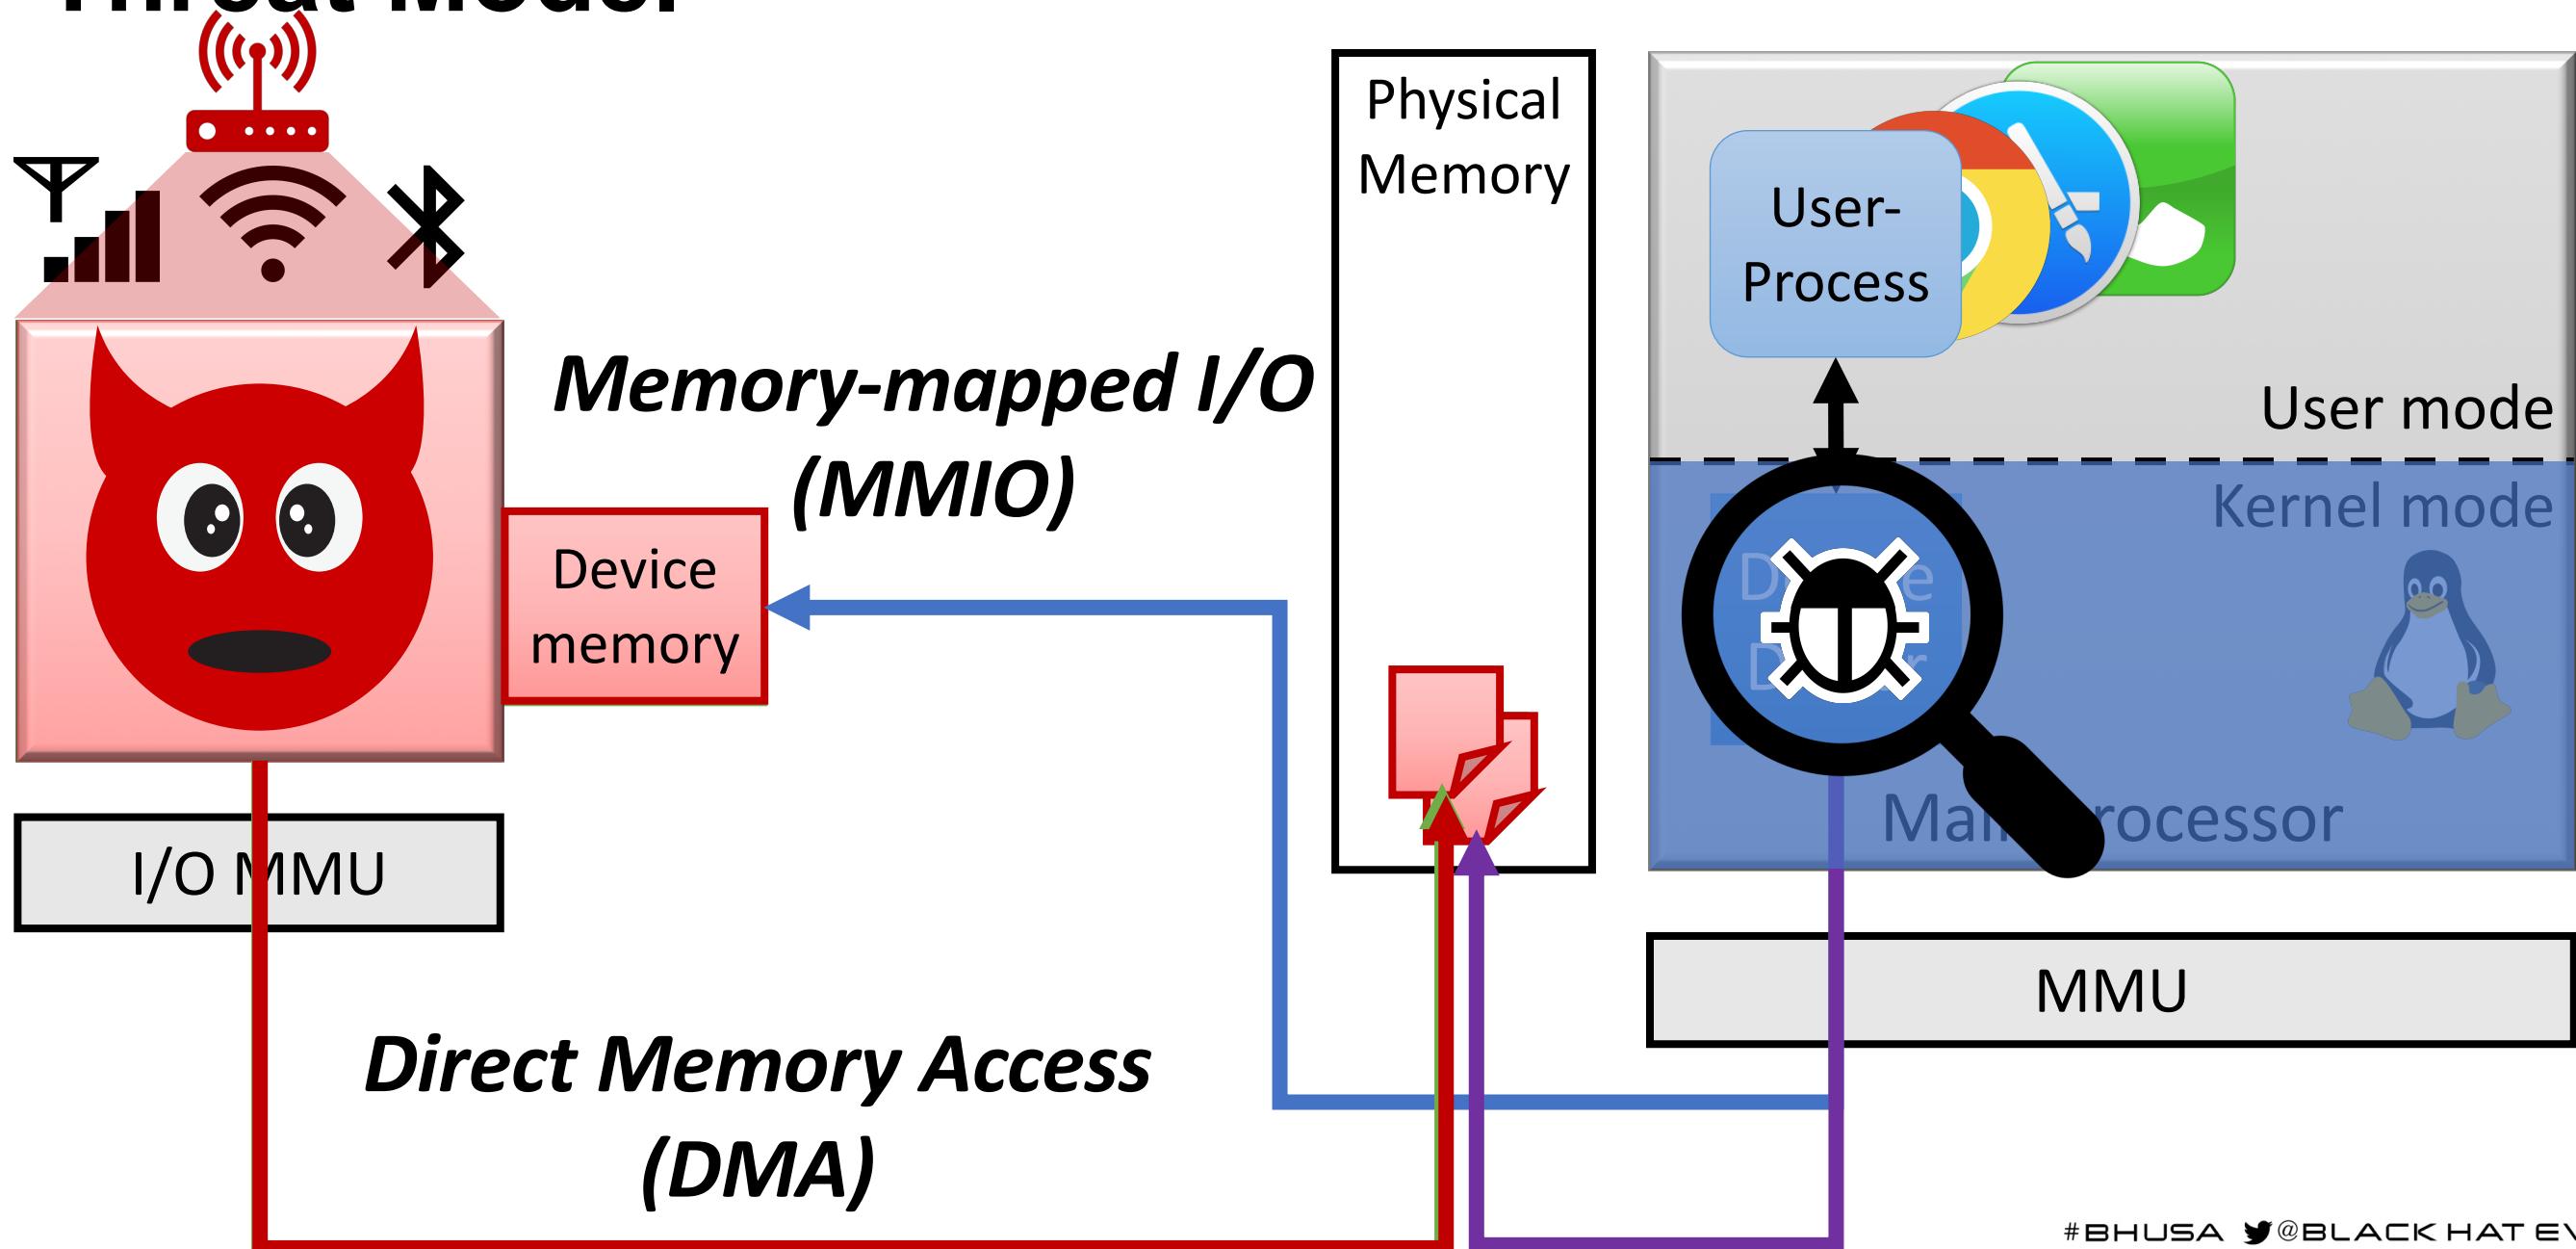
By [Liam Tung](#) | April 5, 2017 -- 10:48 GMT (03:48 PDT) | Topic: [Security](#)

Arbitrary Code Execution on Broadcom Wi-Fi chips *and* Main Processor
– Gal Beniamini at Google Project Zero

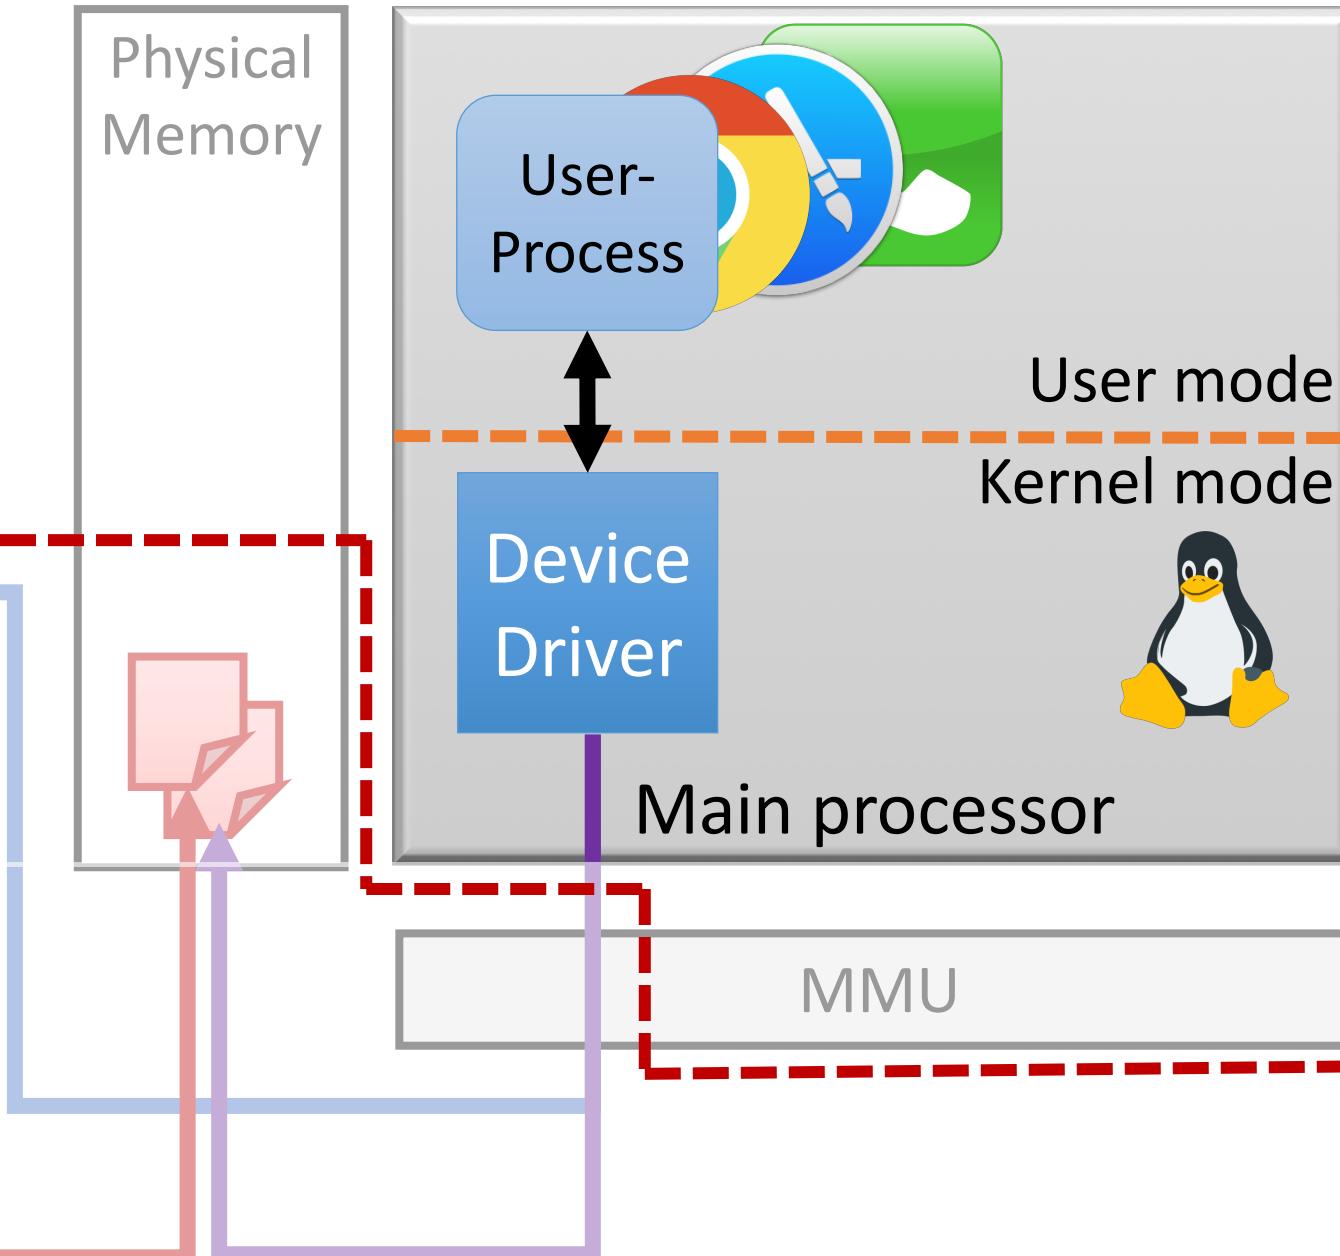
Hardware-OS Interface: MMIO and DMA



Threat Model



Hardware-OS Boundary vs. System Call Boundary



System Call Boundary

(A diverse set of) Sandboxes

- App sandbox, e.g., per-app UID
- Mandatory access control, i.e., SELinux
- Browser sandbox, e.g., seccomp
- I/O daemons



Hardware-OS Boundary

Sandbox: Page-granularity IOMMU

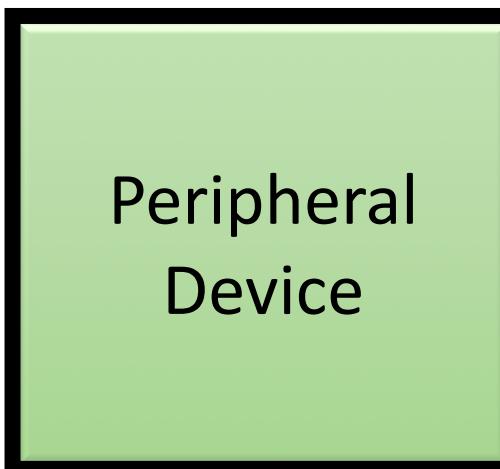
Memory Exploit Mitigations: Peripheral Firmware vs. User-mode Process

ARMv7, Xtensa, MIPS, ...
(aka microcontrollers)

↔ ARMv8

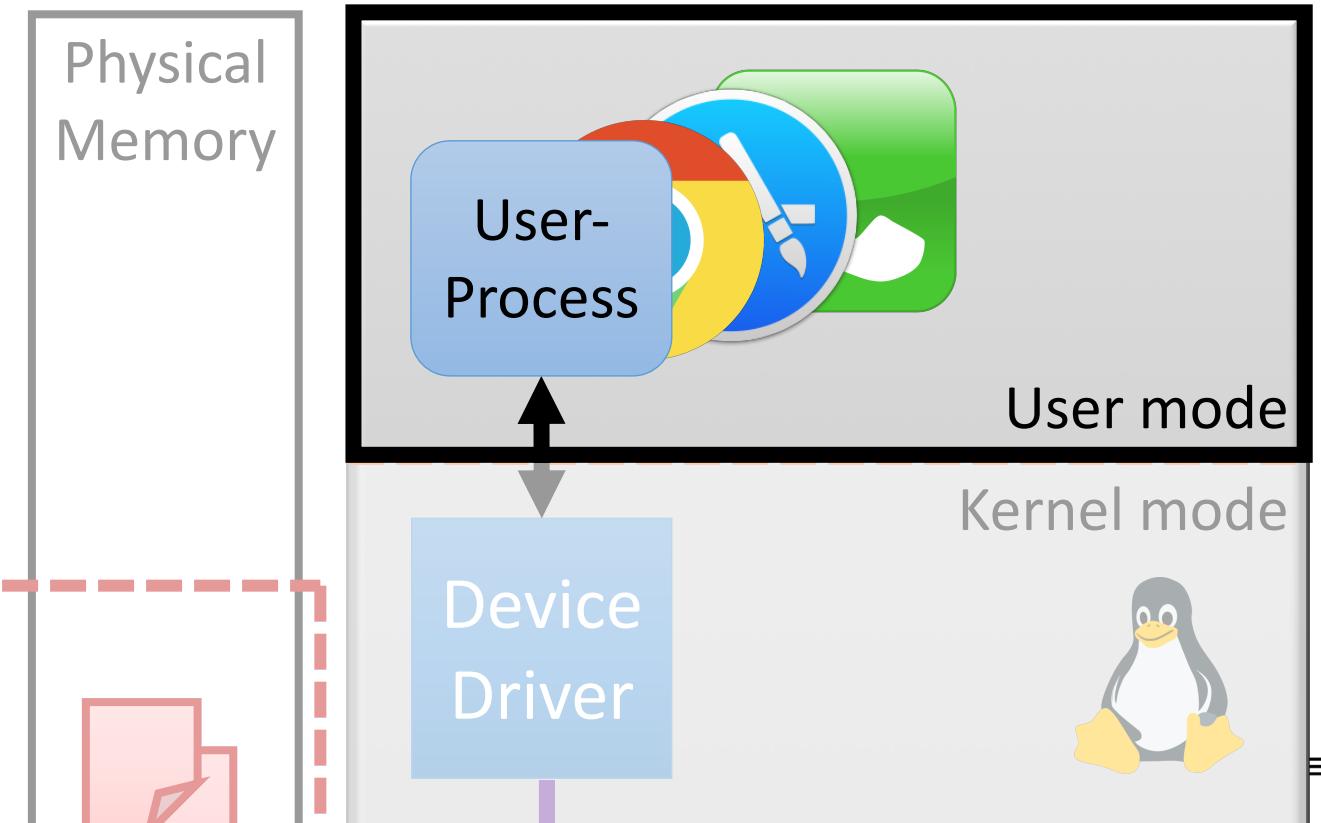
Lack of Mitigations

- Typically resource-constrained (CPU and memory)
- Firmware often runs bare-metal and lacks MMU

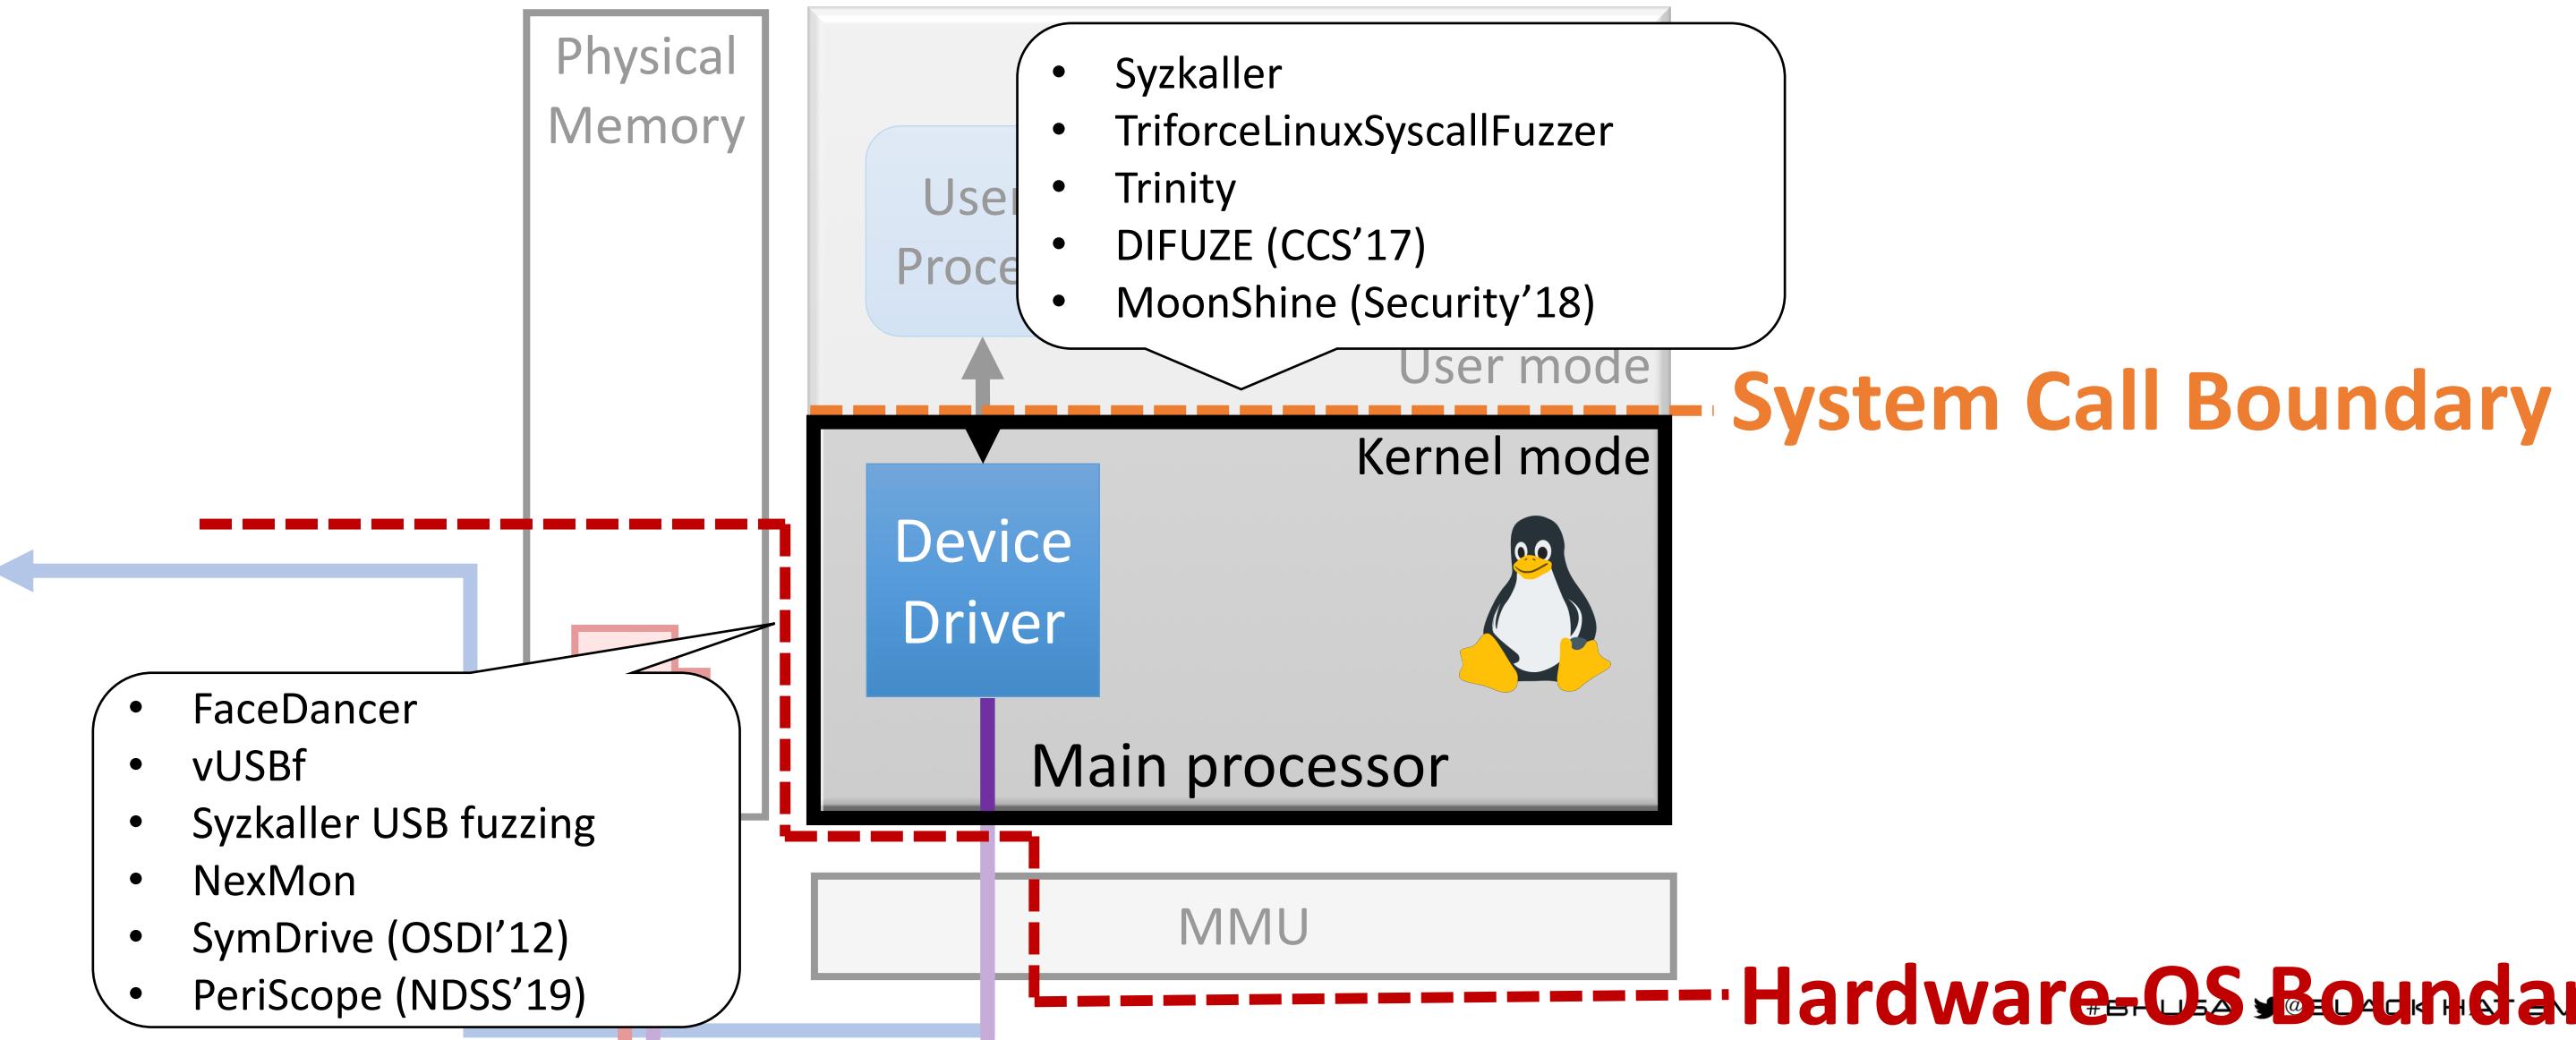


(A diverse set of) Mitigations

- ASLR and DEP
- Control-Flow Integrity
- Secure Heap Allocator



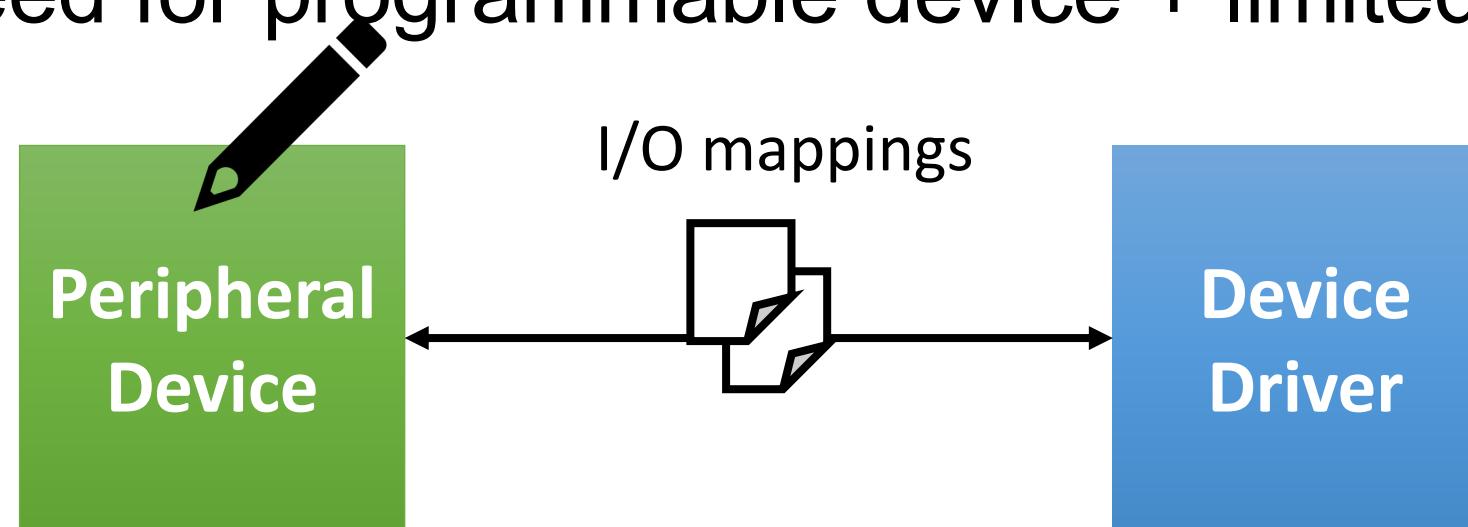
Analysis Tools: Hardware-OS Boundary vs. System Call Boundary



State-of-the-art: Analyzing HW-OS Interface (1/3)

- **Device Adaptation**

- **Pros:** Non-intrusive (OS-independent)
- **Cons:** Need for programmable device + limited visibility into driver

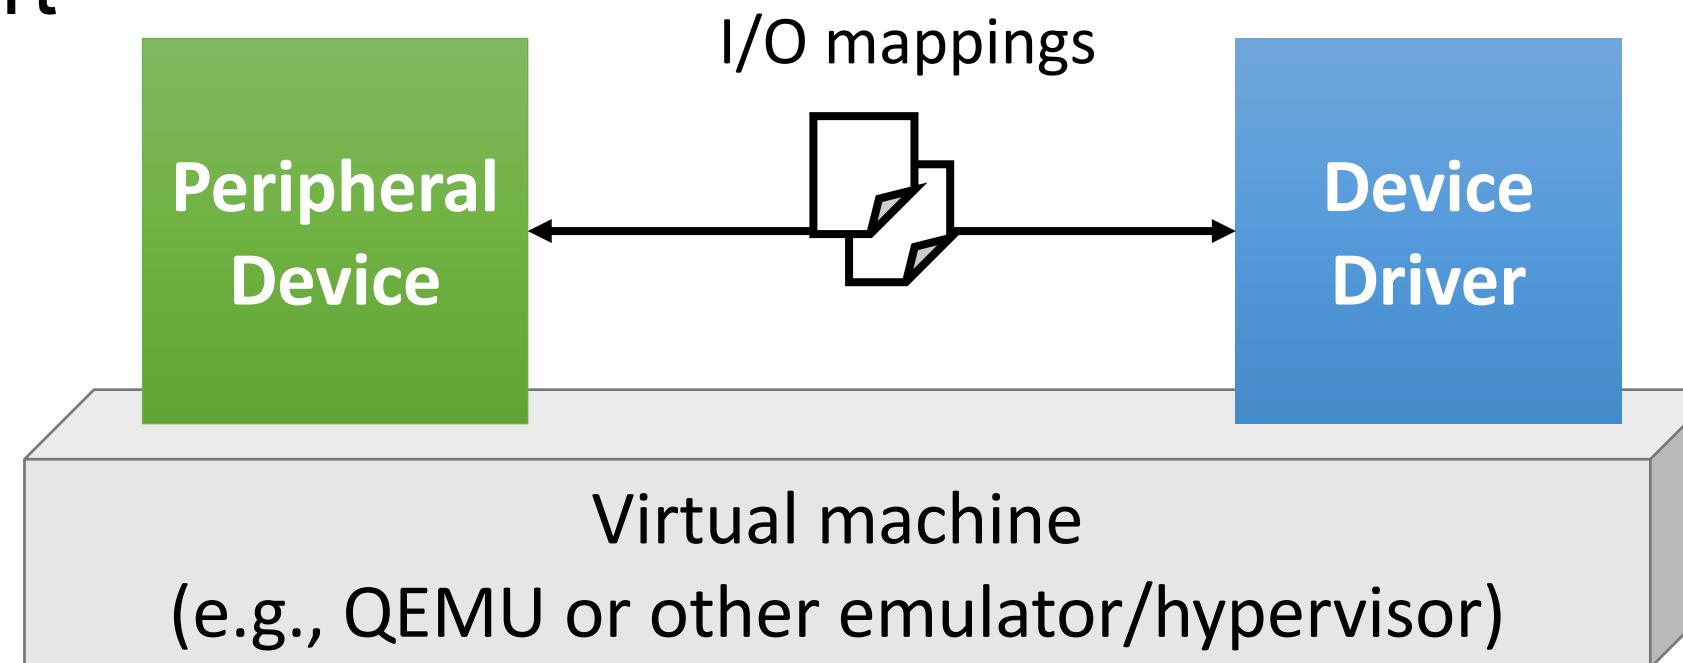


Reprogram the device
(e.g., FaceDancer21 custom USB)

State-of-the-art: Analyzing HW-OS Interface (2/3)

- ***Virtual Machine Introspection***

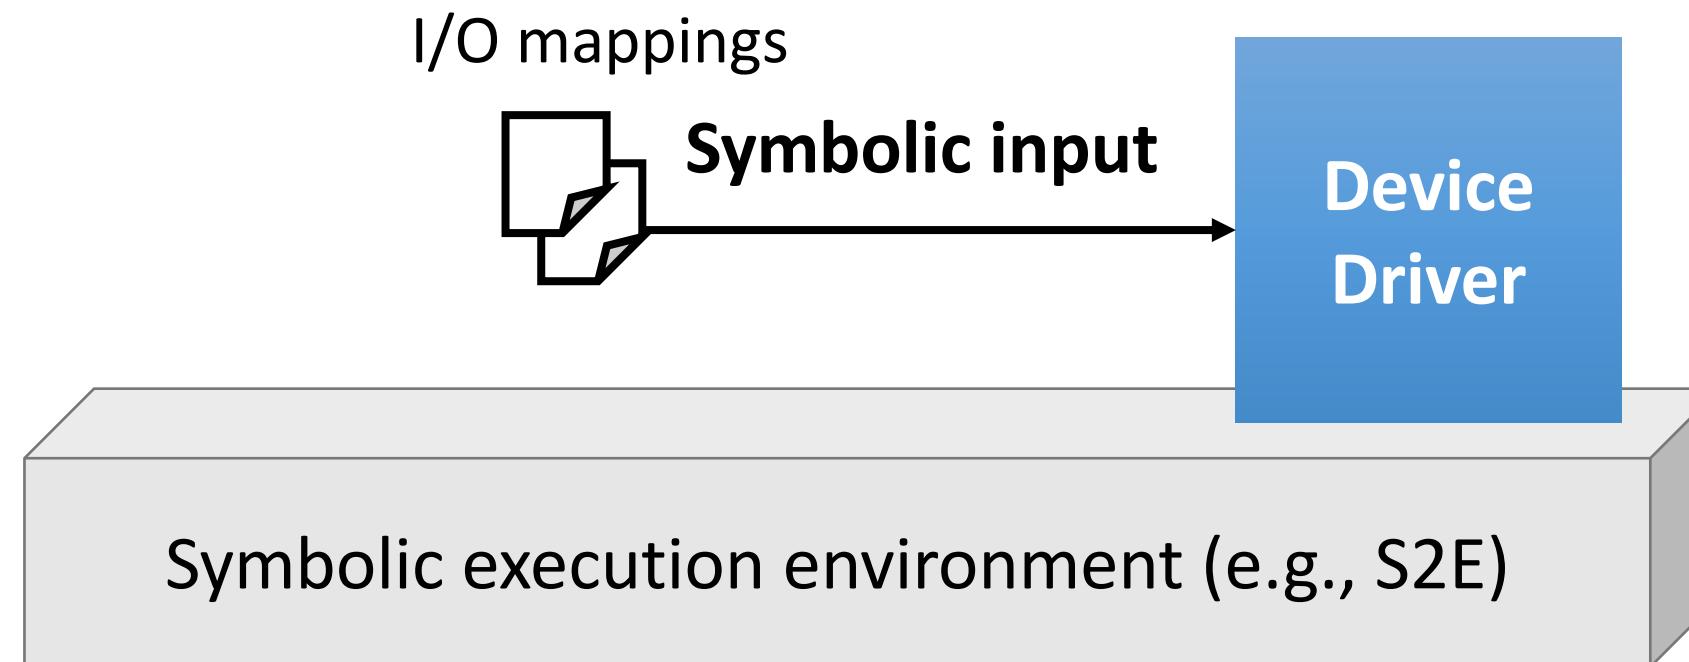
- **Pros:** High visibility yet non-intrusive
- **Cons:** Need for virtual device and/or virtualization HW support



State-of-the-art: Analyzing HW-OS Interface (3/3)

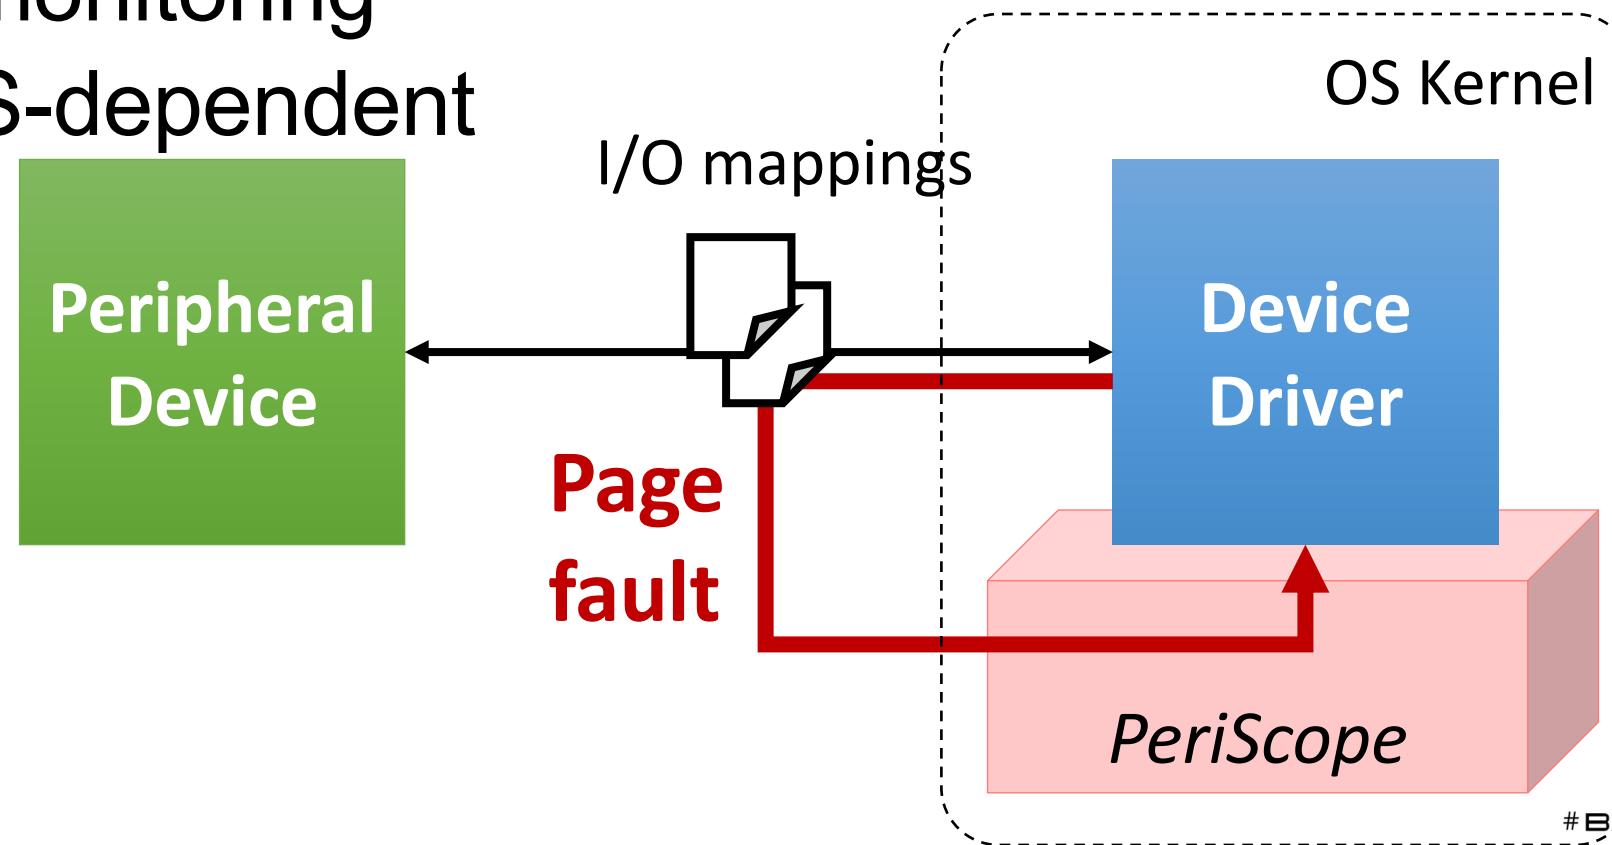
- ***Symbolic Devices***

- **Pros:** No need for physical/virtual device
- **Cons:** Inherits cons of symbolic execution

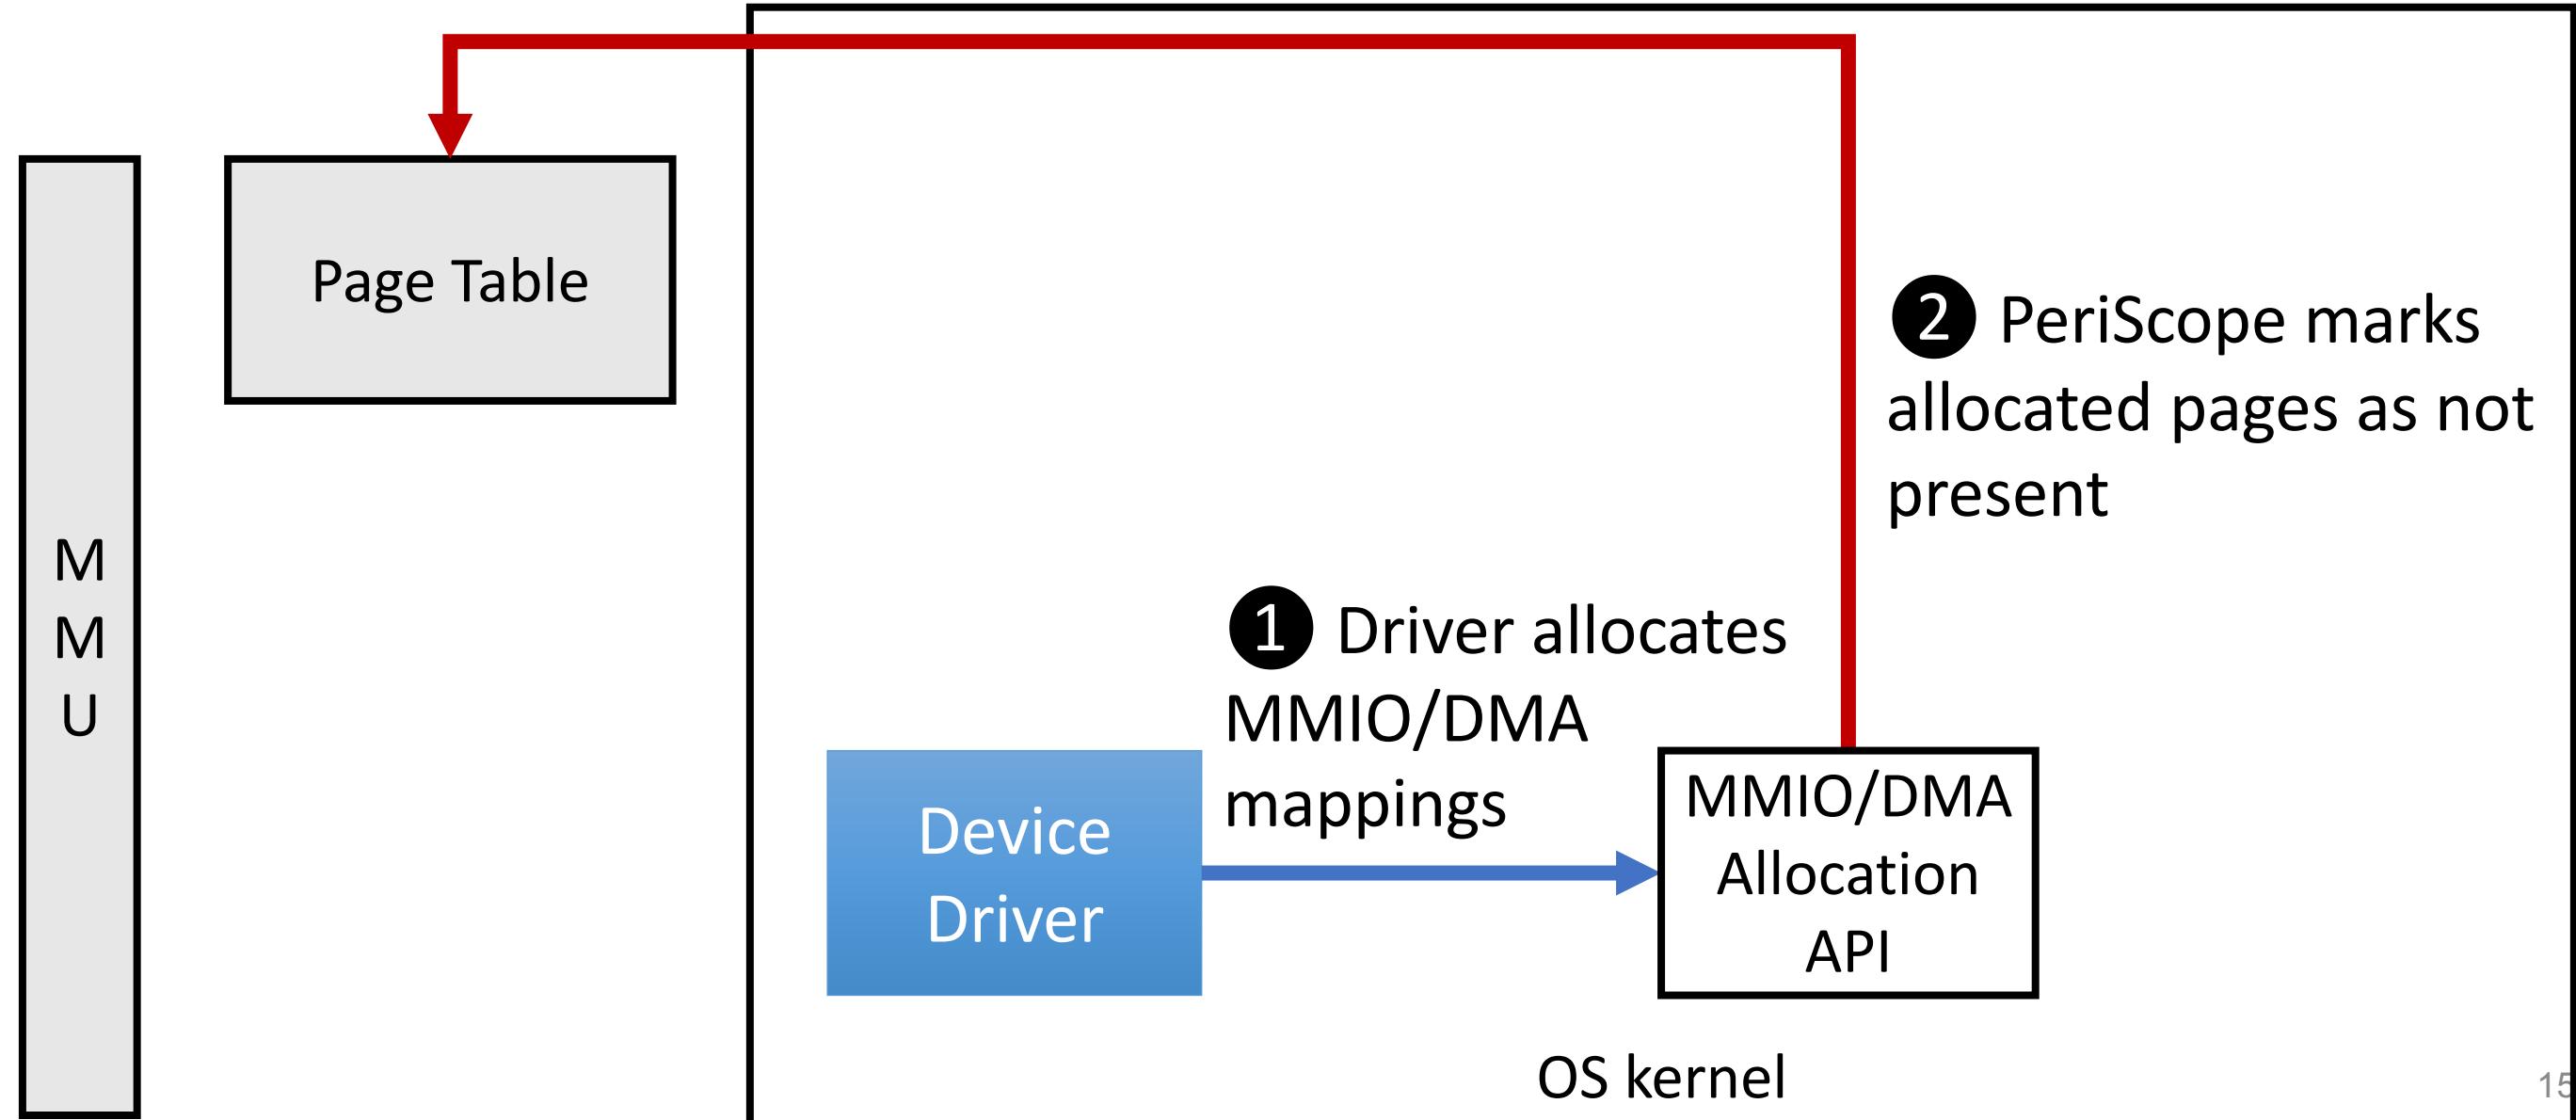


PeriScope – Our Approach

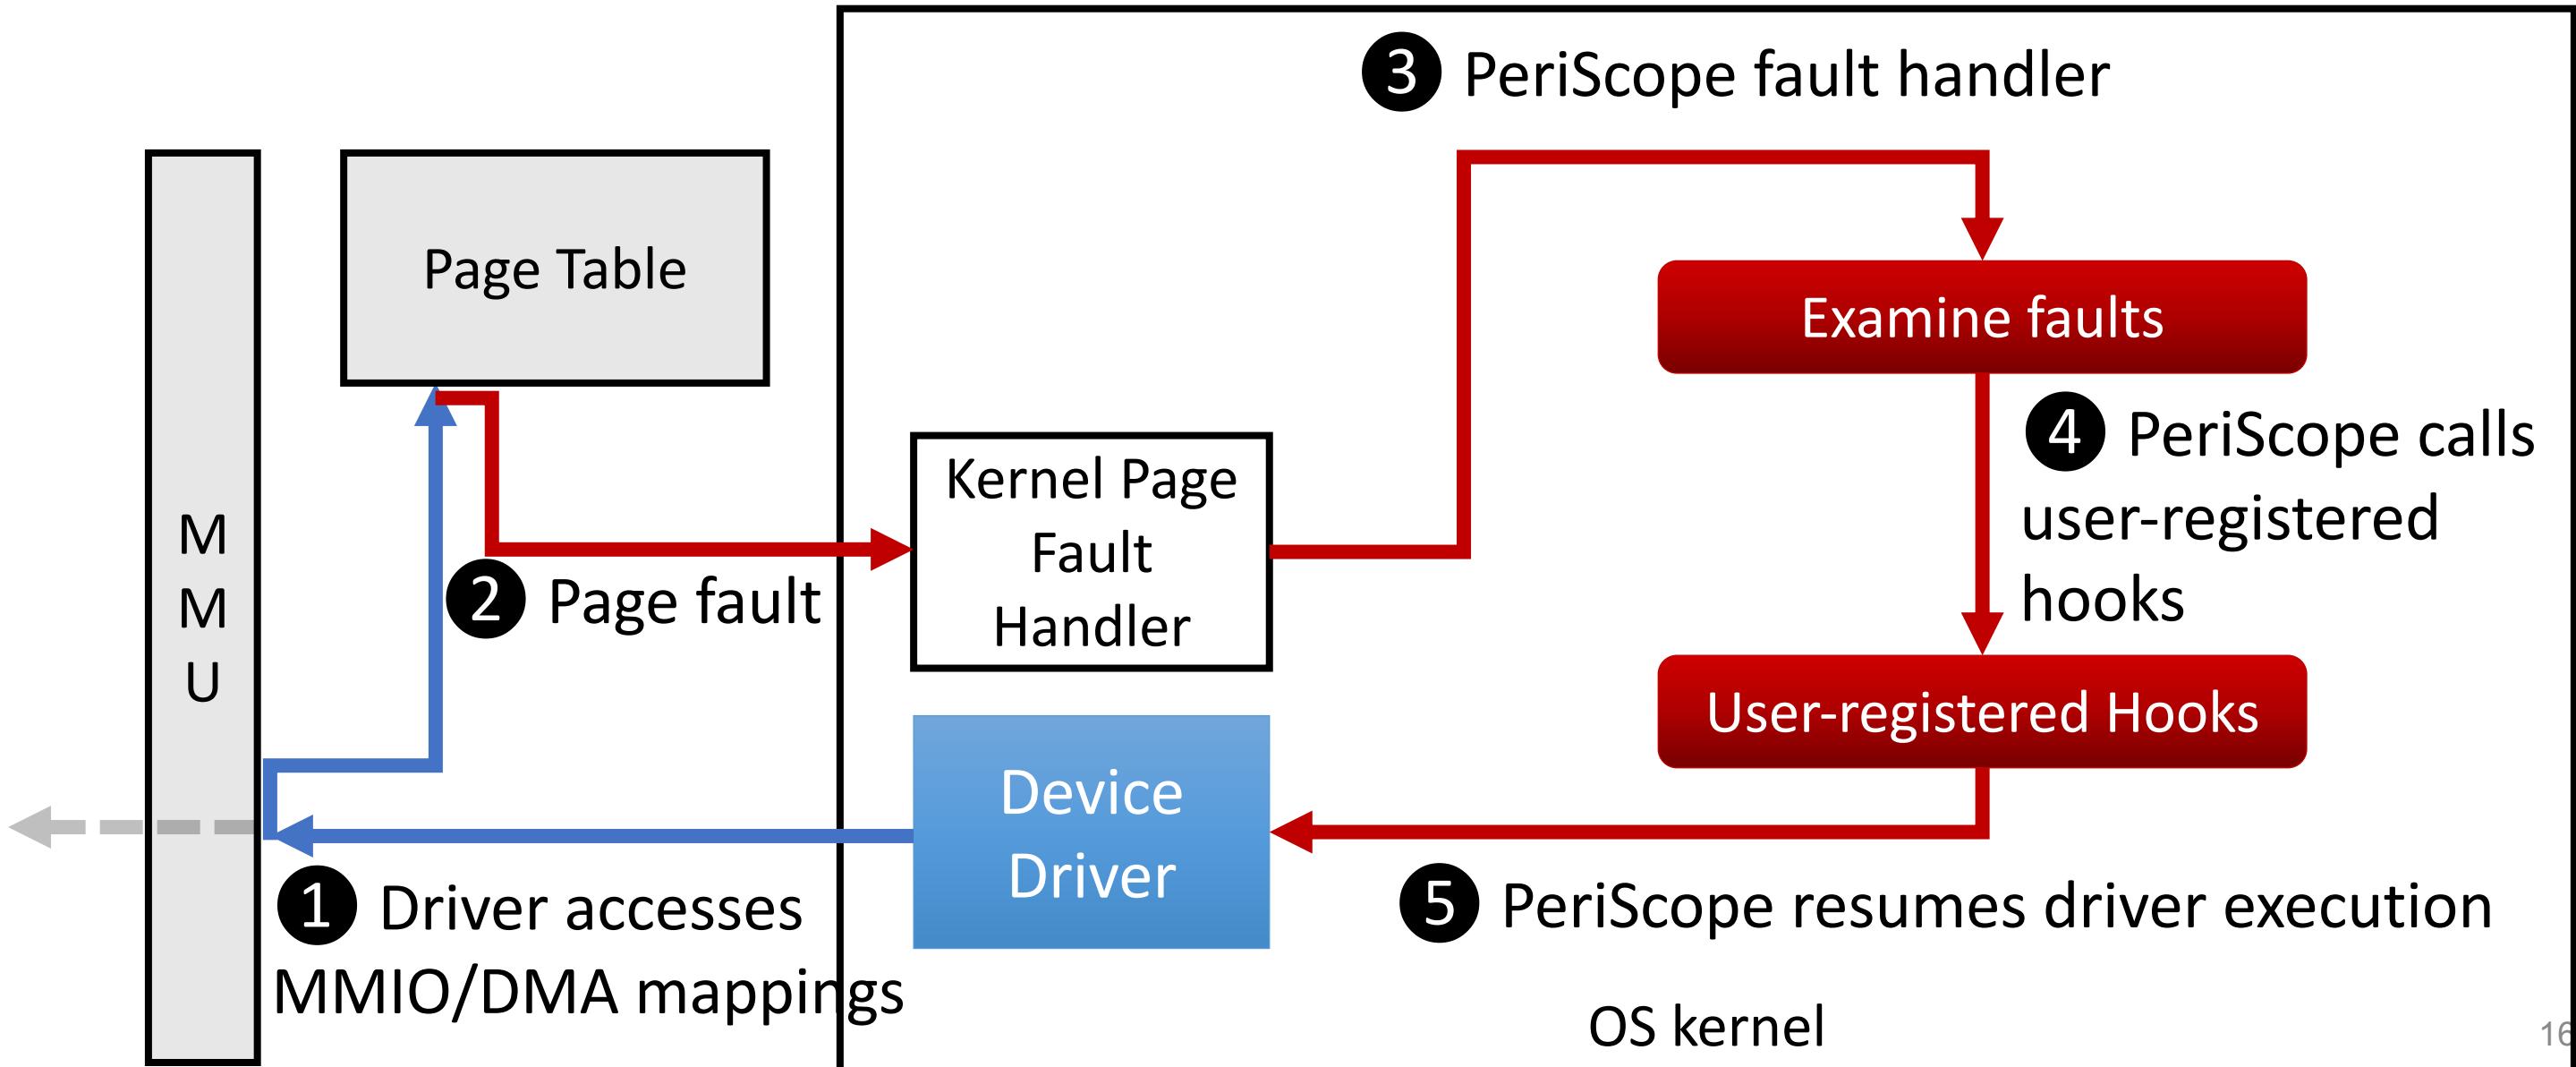
- ***In-kernel, page-fault-based monitoring***
 - Pros: No device-specific/virtual device requirement yet fine-grained monitoring
 - Cons: OS-dependent



PeriScope Overview



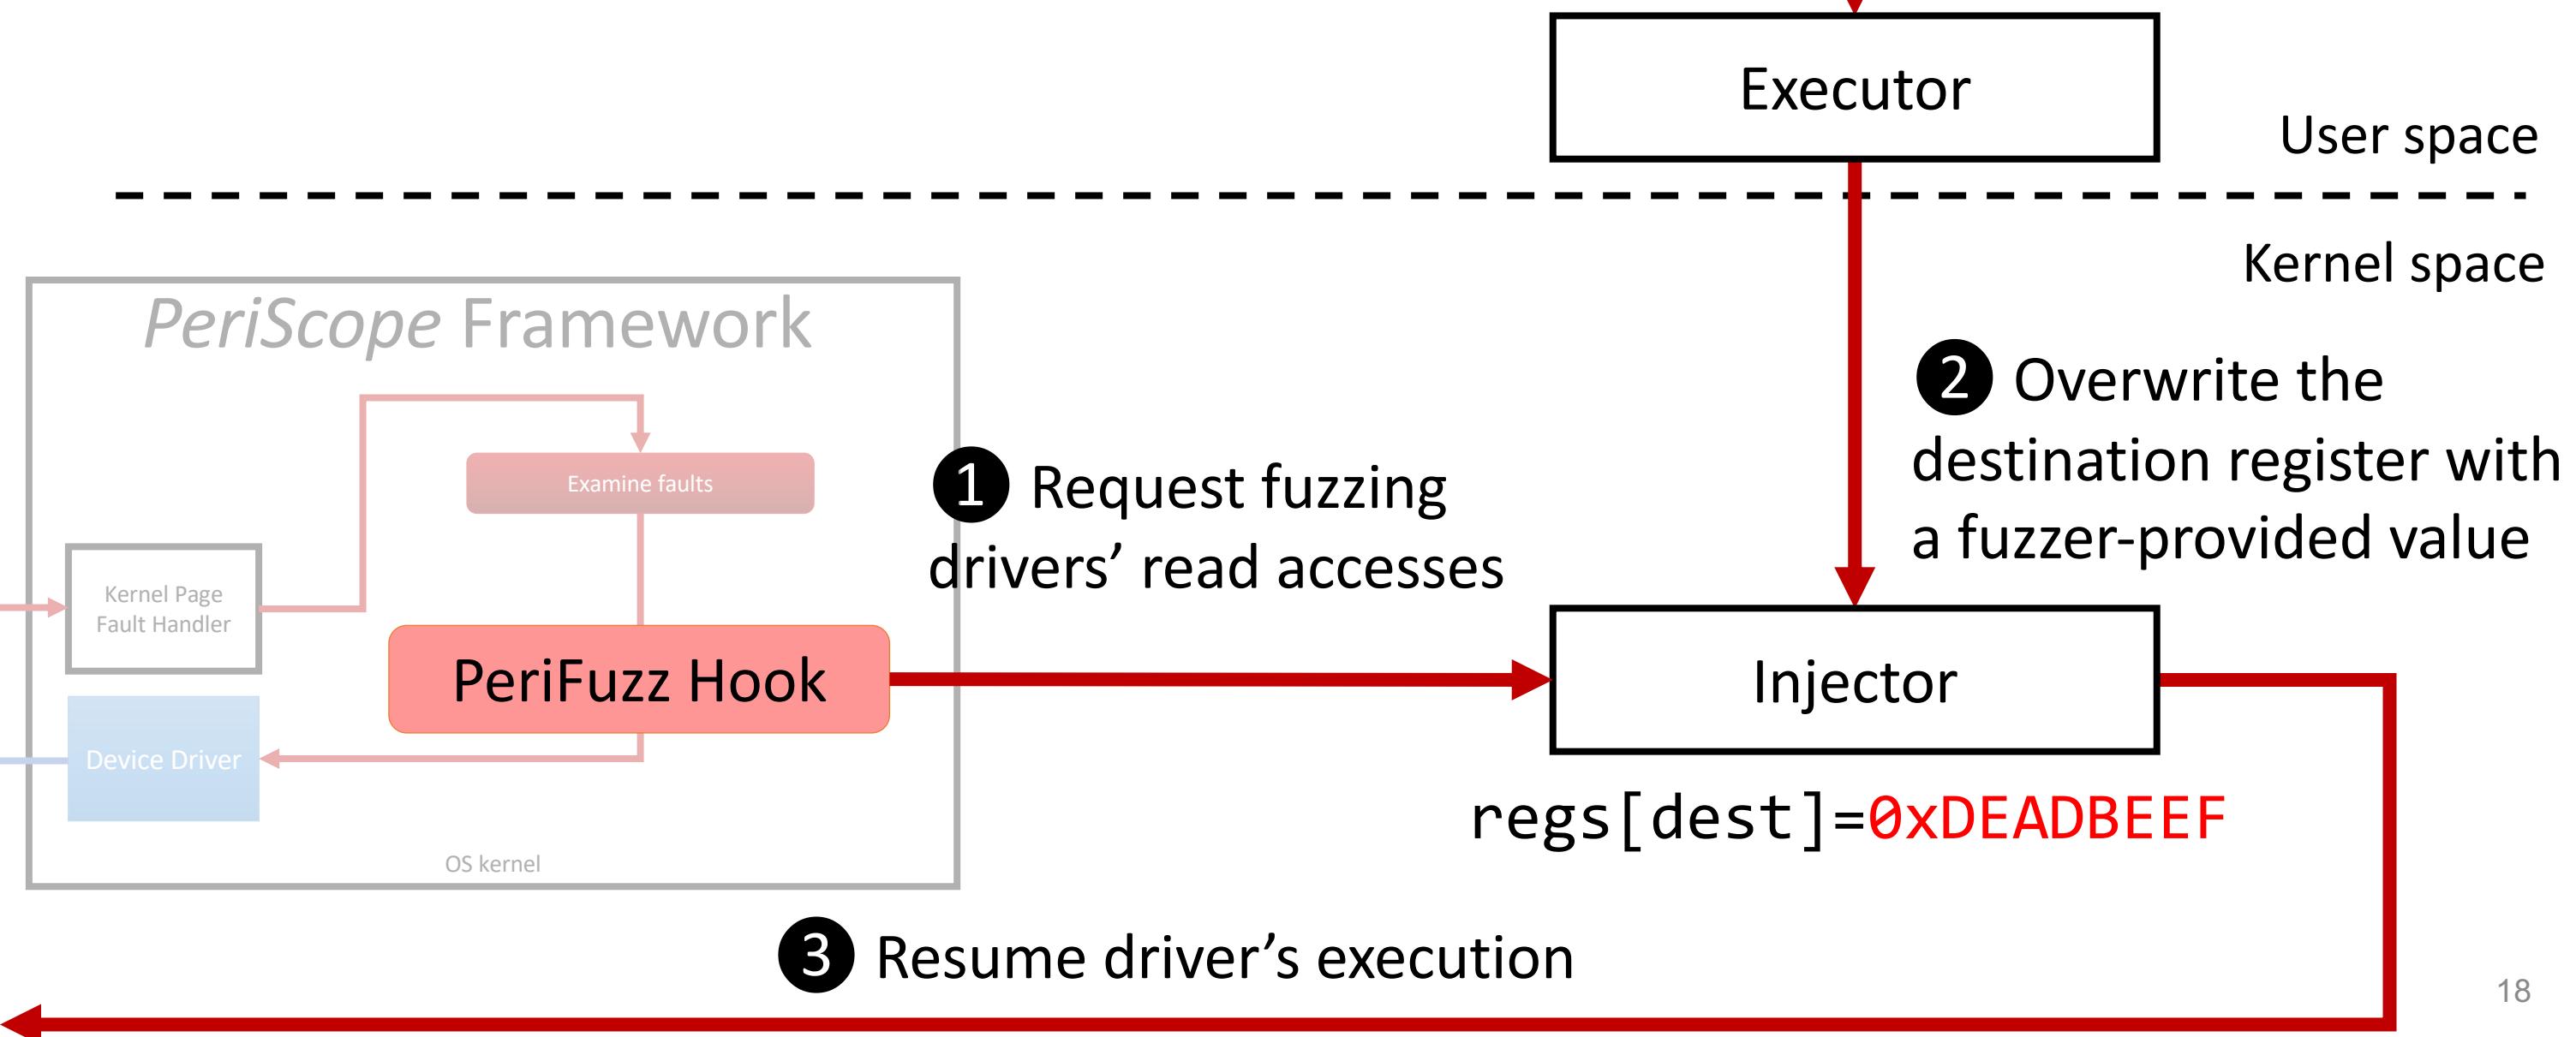
PeriScope Overview



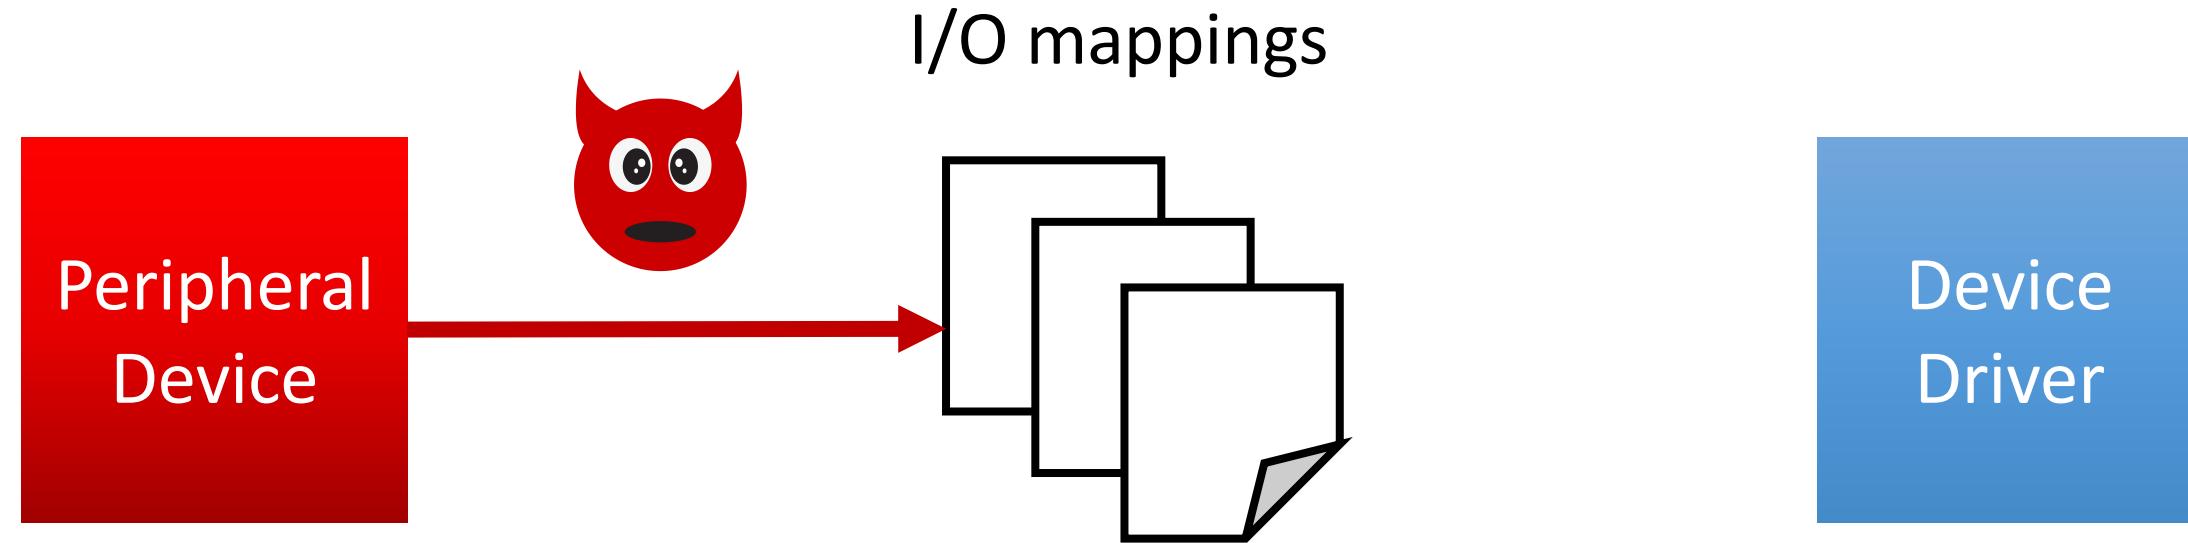
PeriFuzz – Fuzzer for the HW-OS boundary

- **Goal:** To find vulnerabilities in kernel drivers reachable from a compromised device
- Therefore, *PeriFuzz* fuzzes **Driver's Read Accesses** to MMIO and DMA mappings

PeriFuzz Overview



Threat Model Review

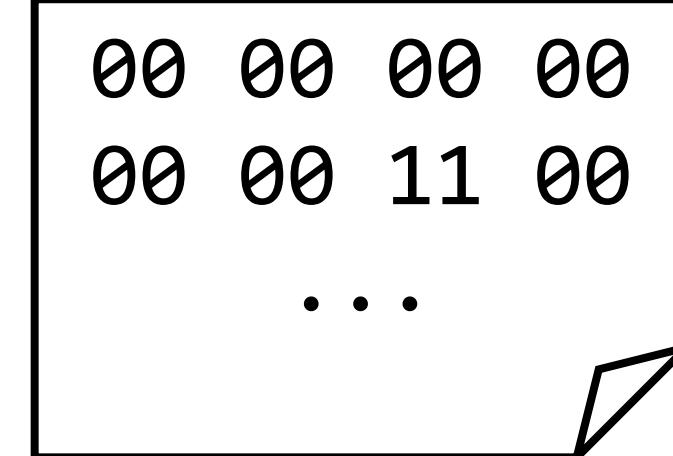


Attacker can write **any value** to the I/O mappings
even multiple times at **any time**

Potential Double-fetch Bugs in I/O Mappings

An I/O mapping

Peripheral
Device



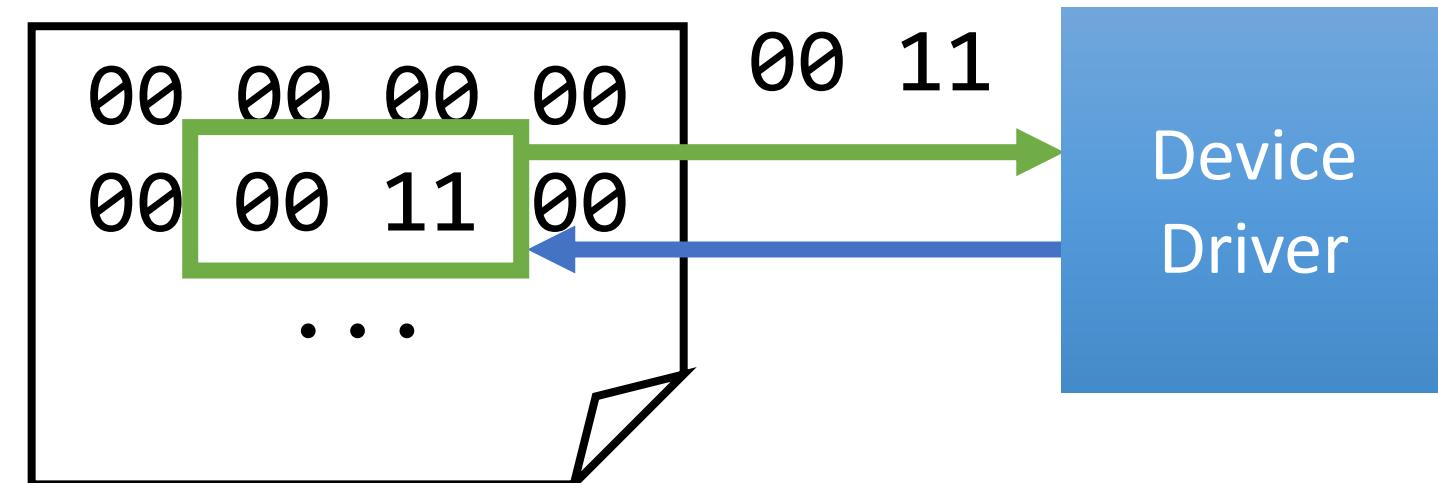
Device
Driver

```
if (*map_ptr <= 0x00FF) {  
    ...  
  
    array[*map_ptr] = ...;
```

Potential Double-fetch Bugs in I/O Mappings

An I/O mapping

Peripheral Device



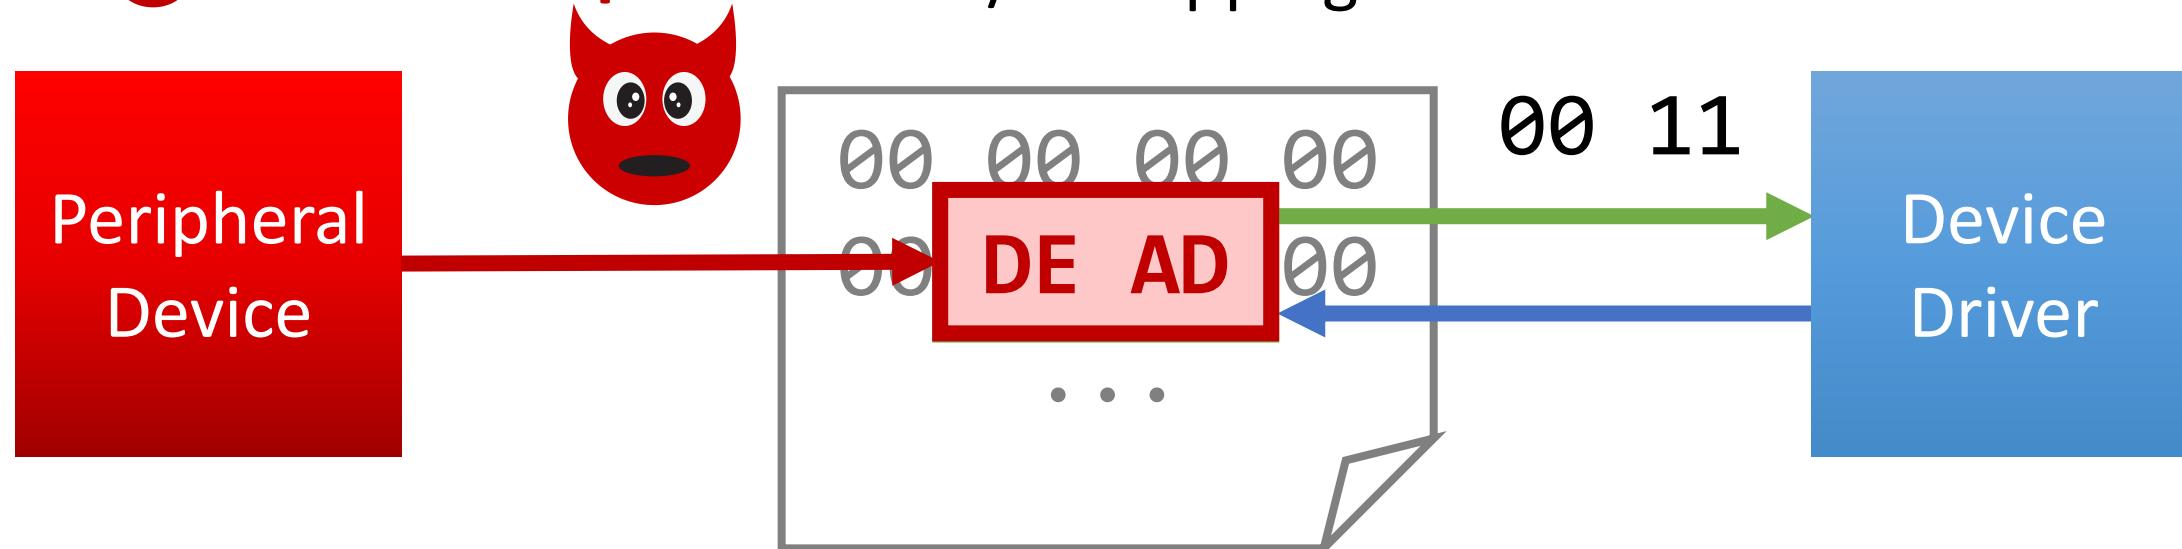
- ① First fetch
& check passes

```
if (*map_ptr <= 0x00FF) {  
    ...  
    array[*map_ptr] = ...;
```

Potential Double-fetch Bugs in I/O Mappings

② Malicious Update

An I/O mapping



① First fetch
& check passes

```
if (*map_ptr <= 0x00FF) {  
    ...  
    array[*map_ptr] = ...;
```

Potential Double-fetch Bugs in I/O Mappings

② Malicious Update

An I/O mapping



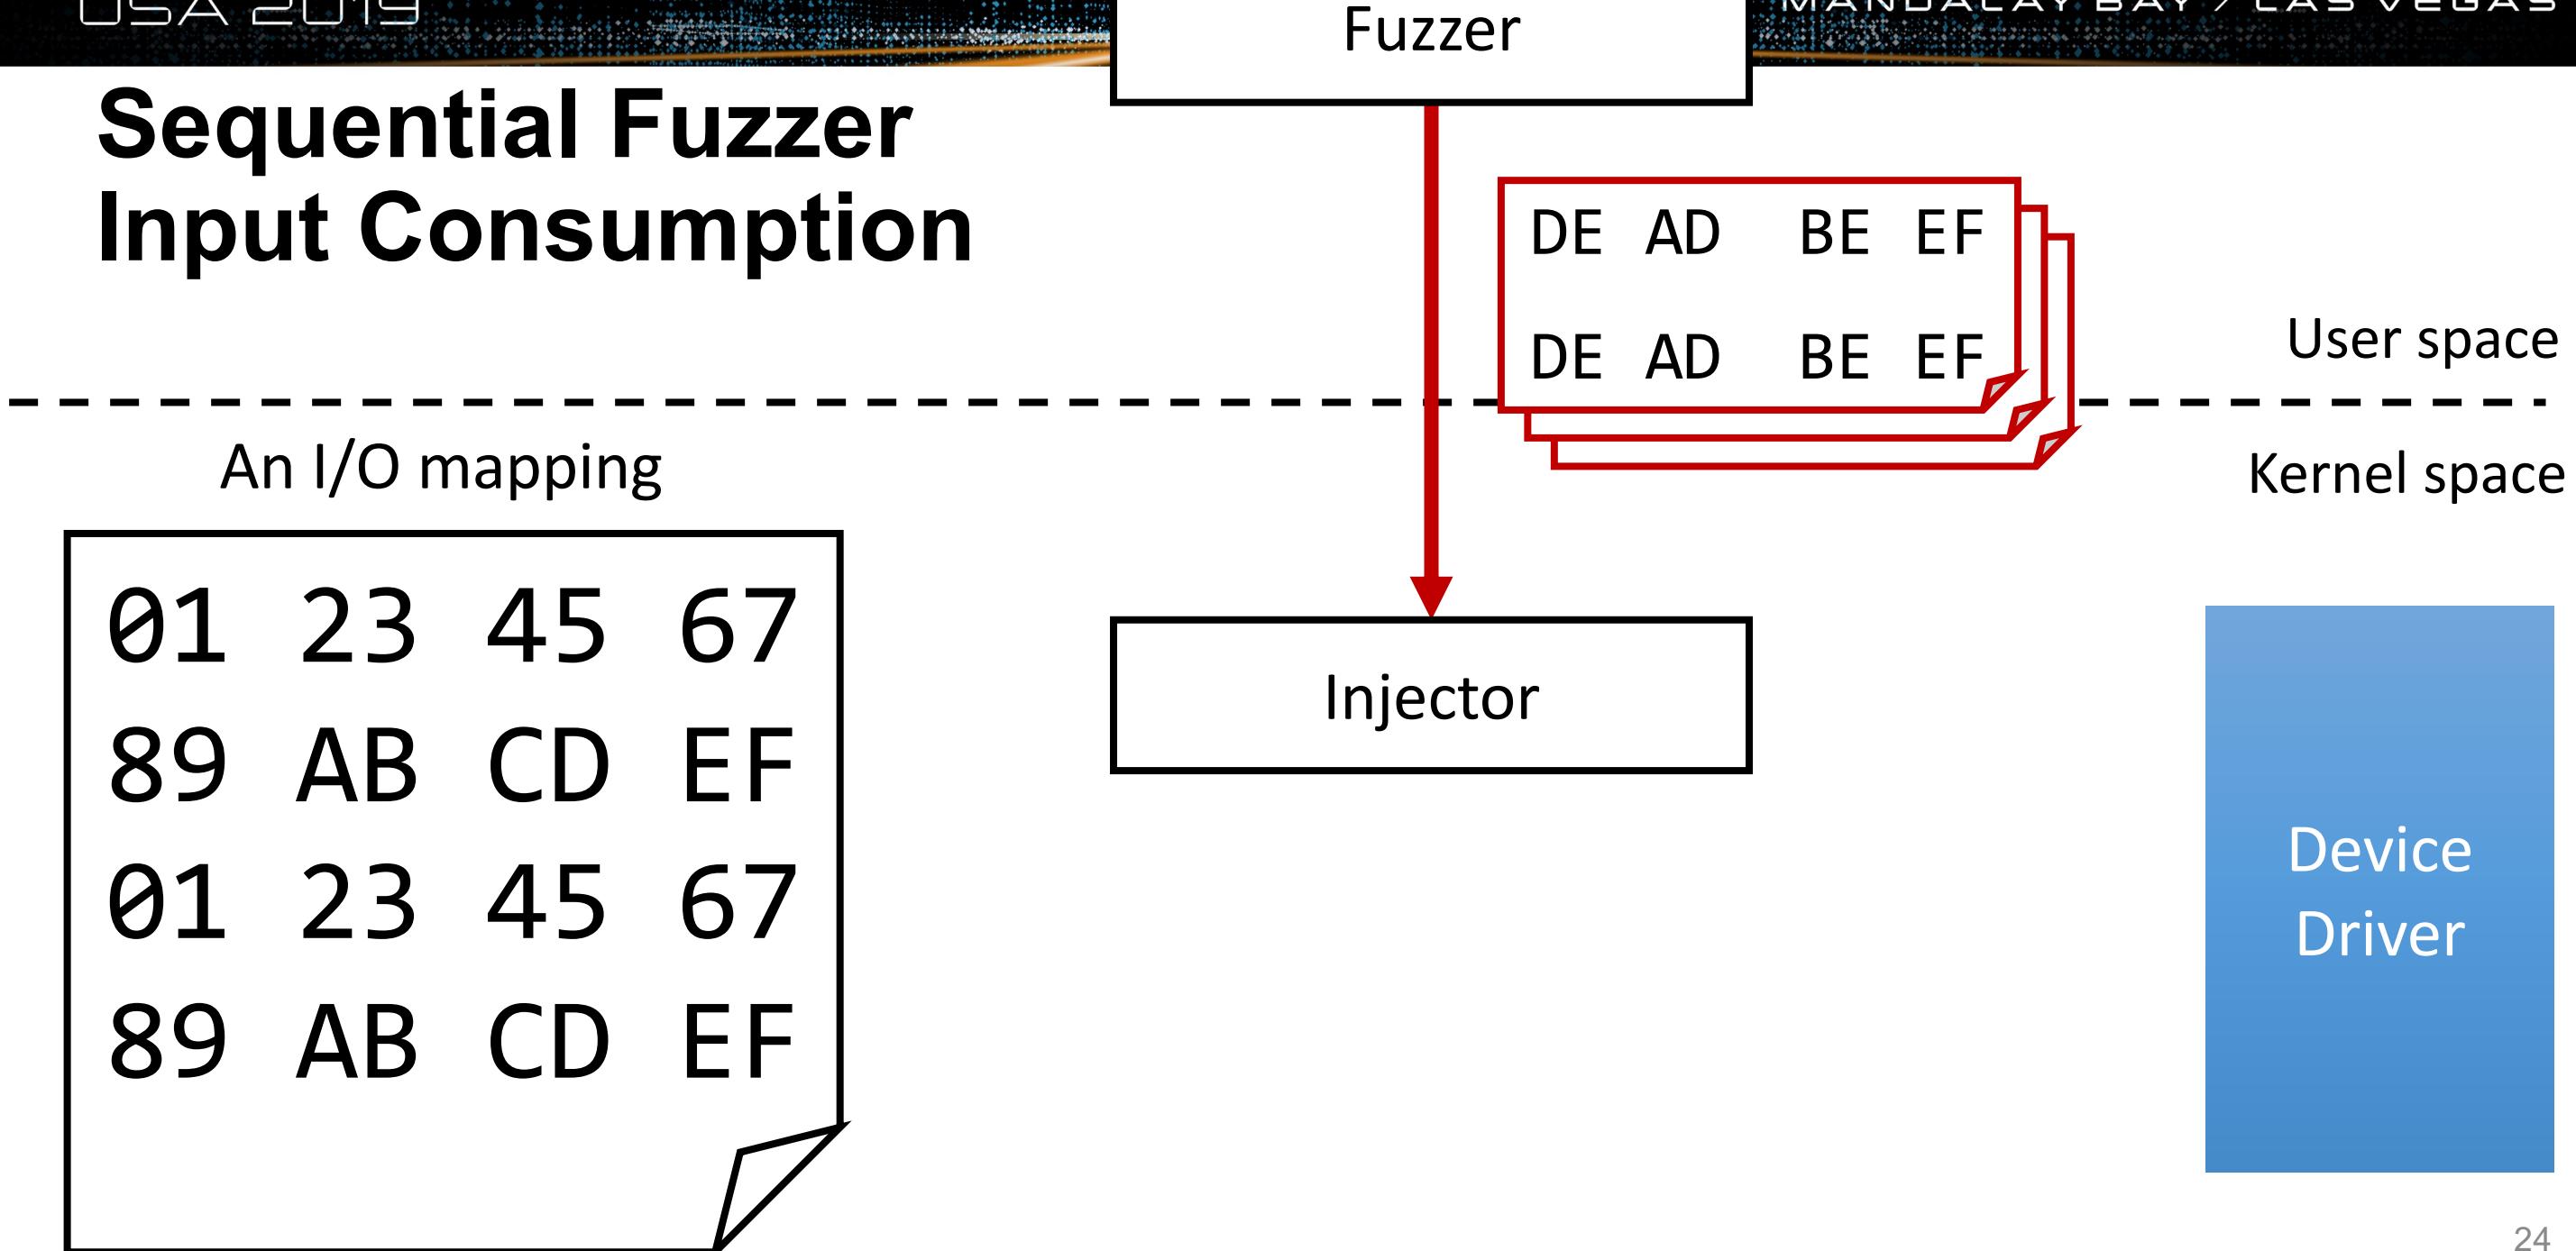
① First fetch
& check passes

③ Overlapping fetch
(without rechecking)

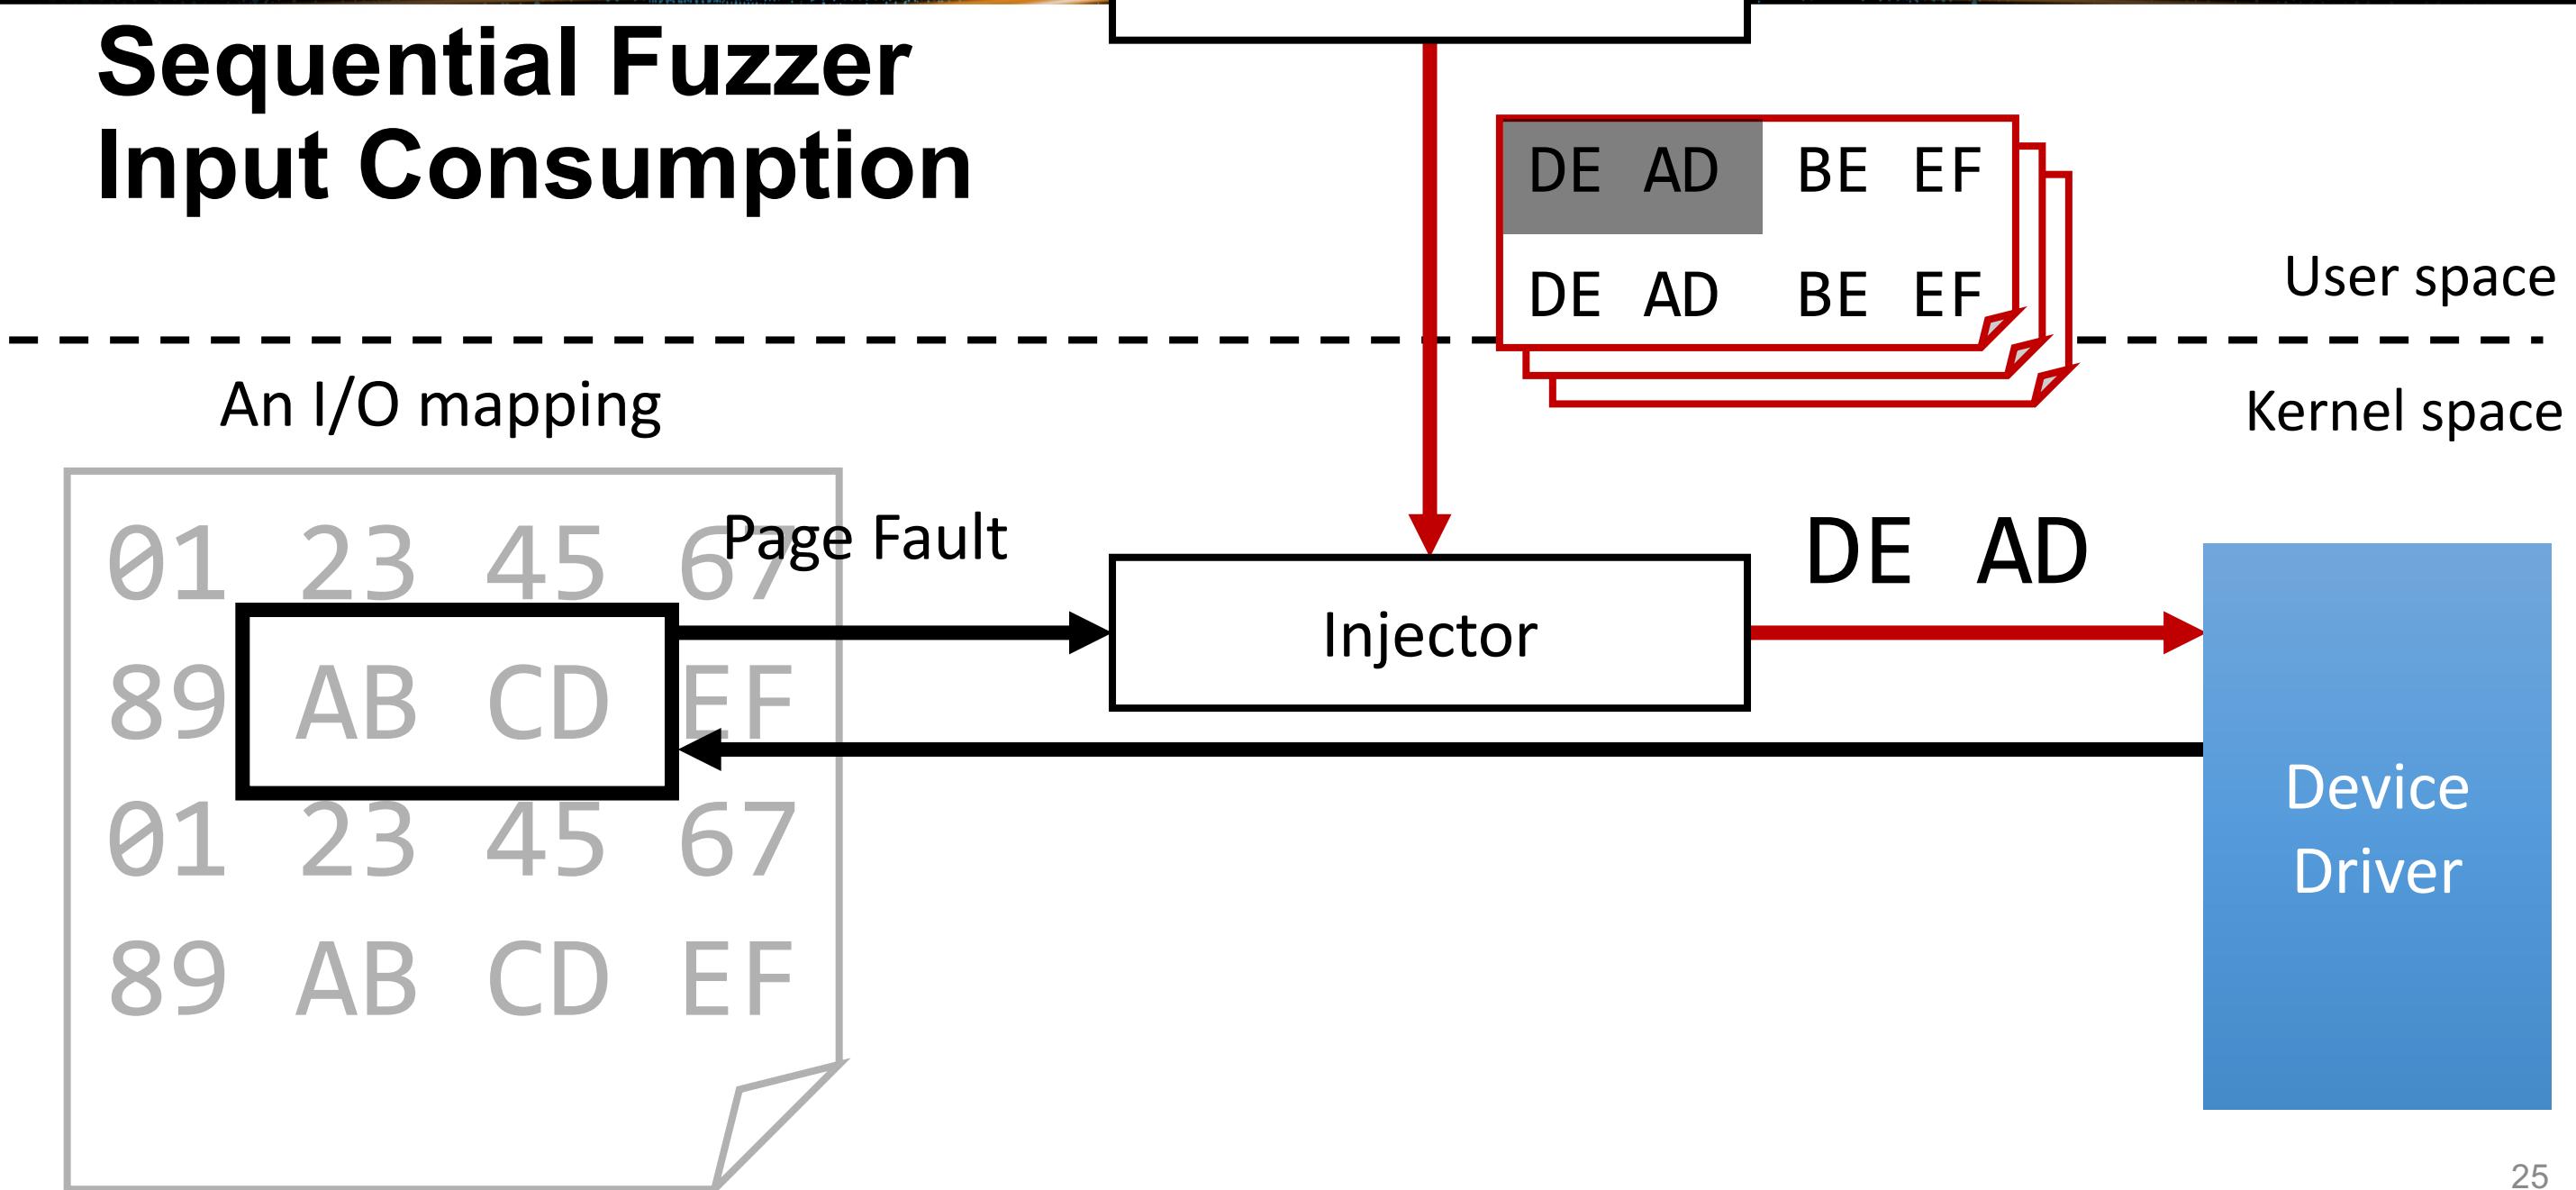
```
if (*map_ptr <= 0x00FF) {  
    ...  
    array[*map_ptr] = ...;
```

23

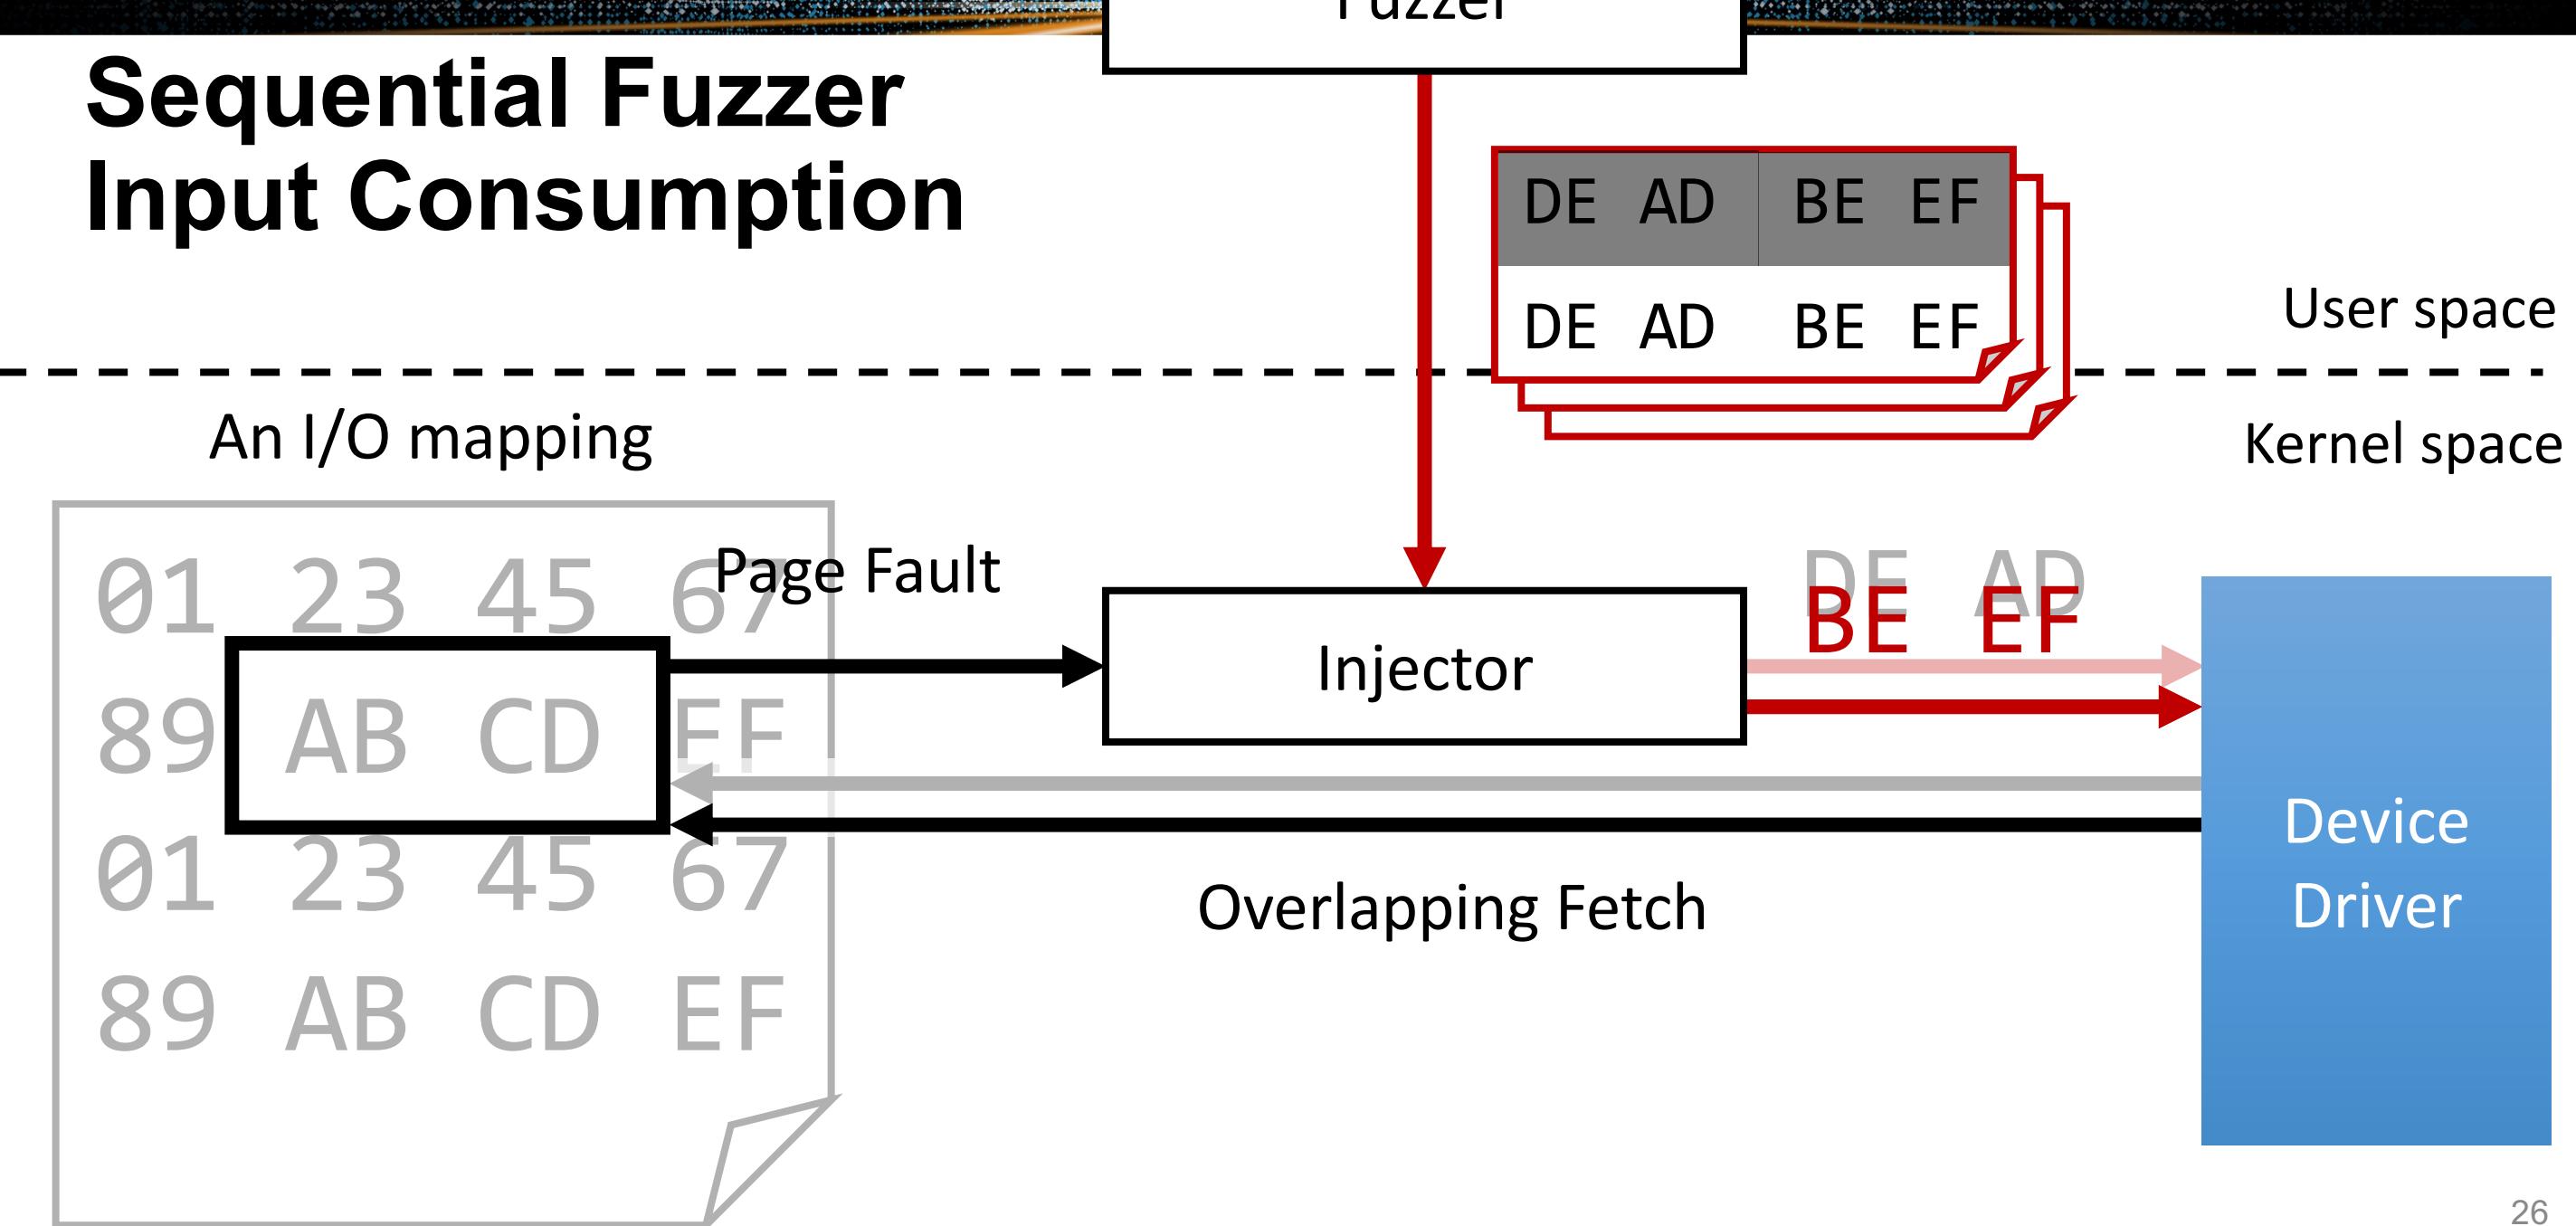
Sequential Fuzzer Input Consumption



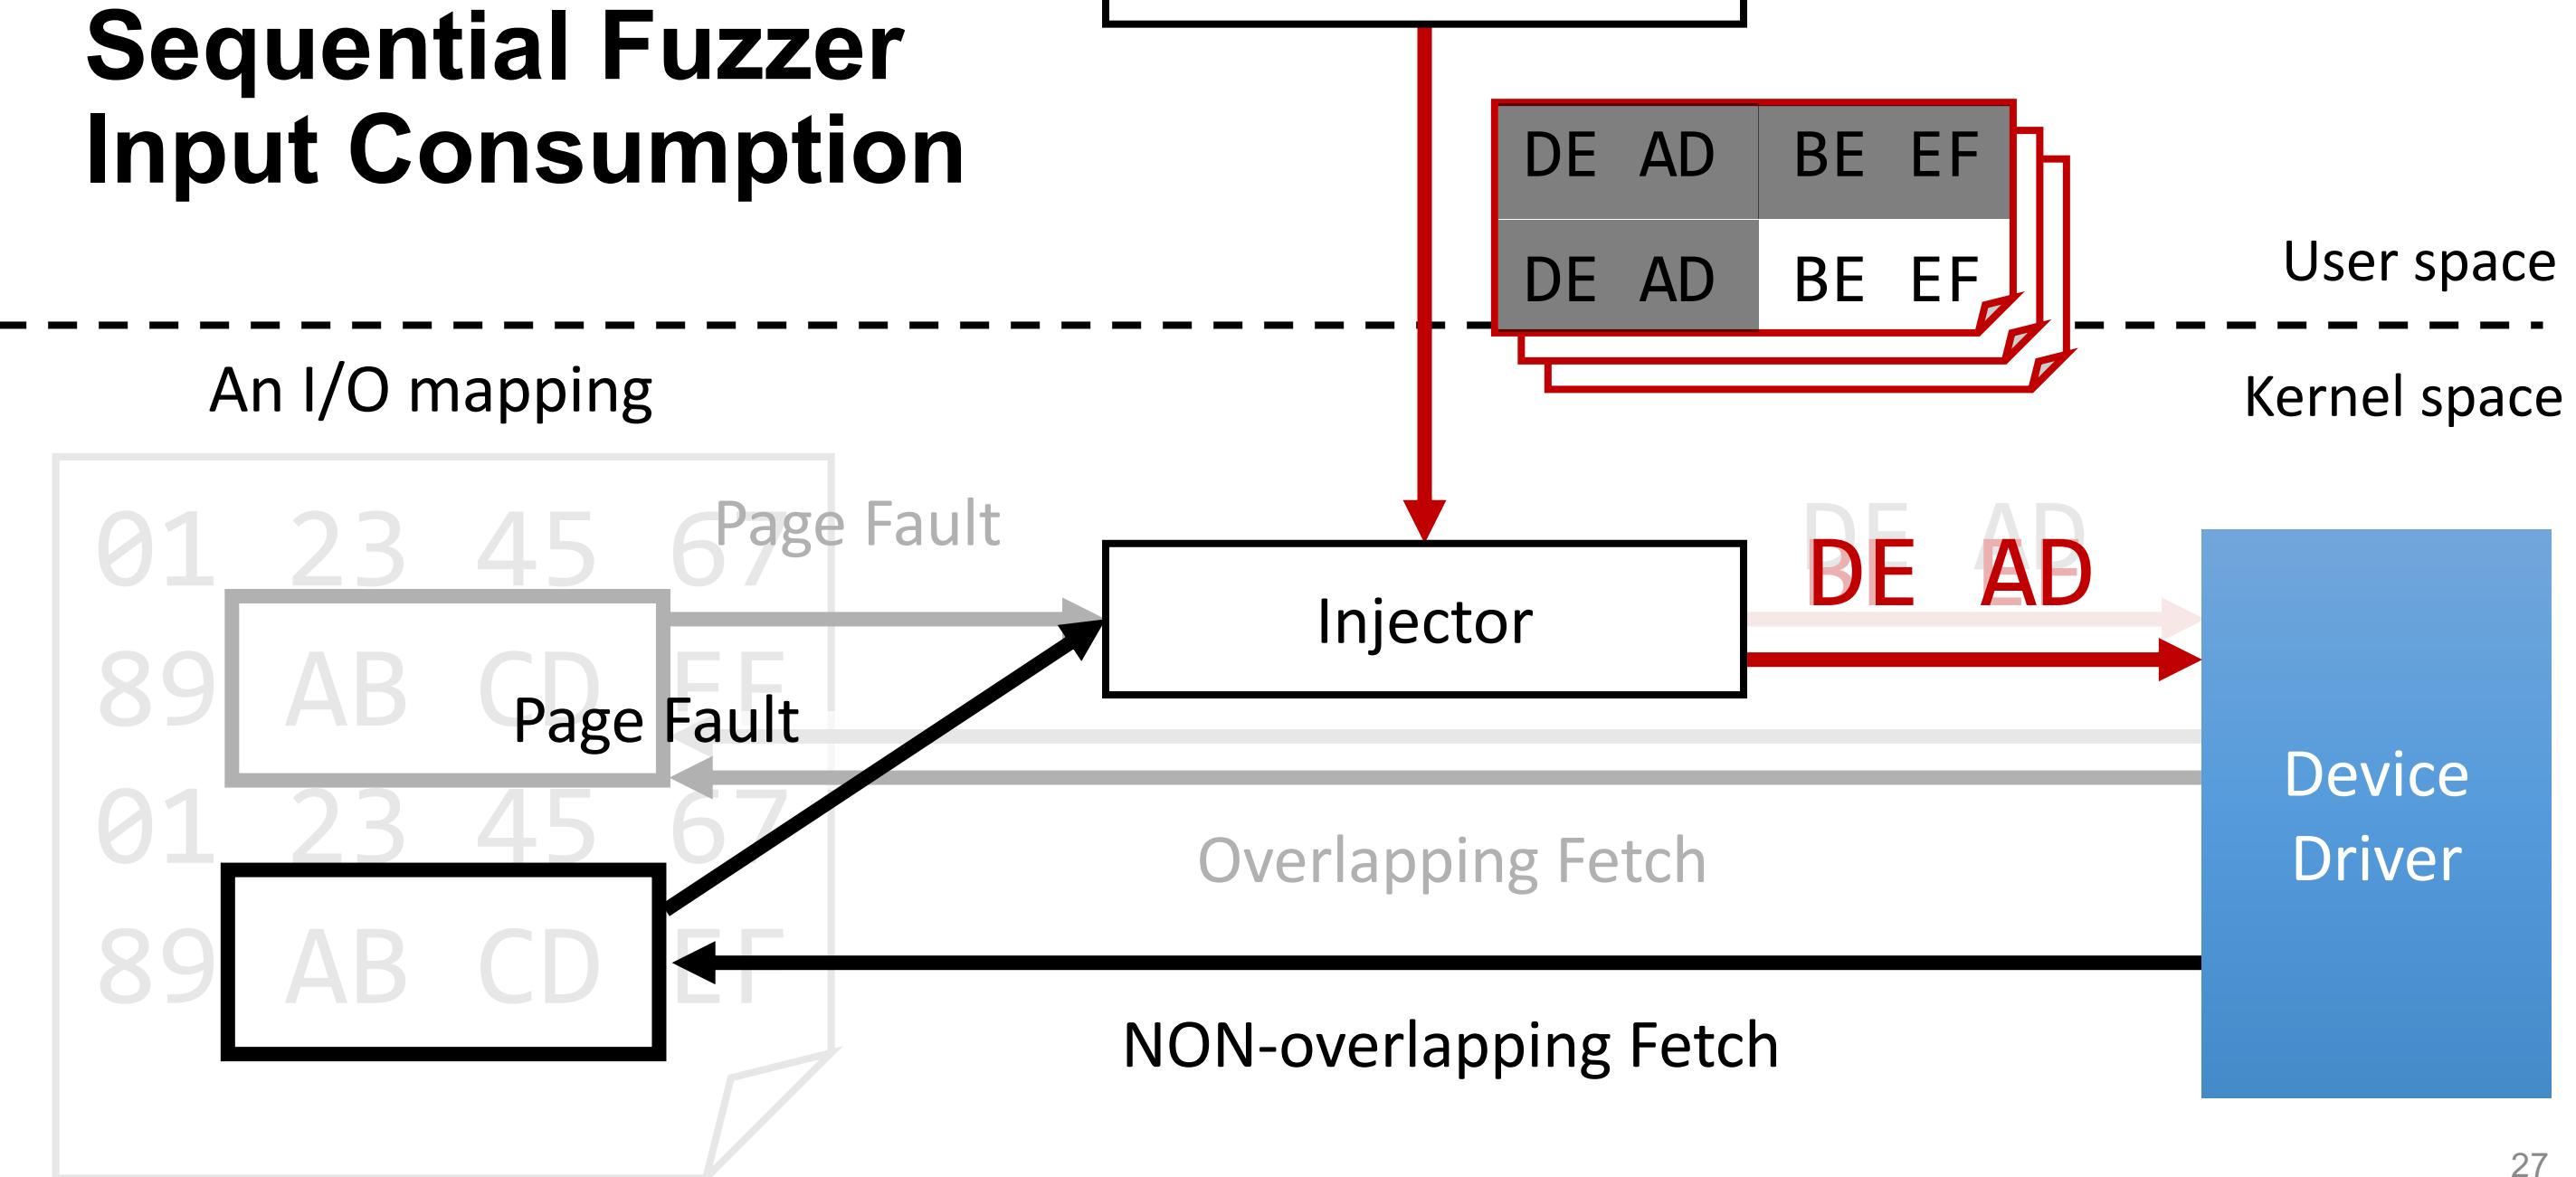
Sequential Fuzzer Input Consumption



Sequential Fuzzer Input Consumption

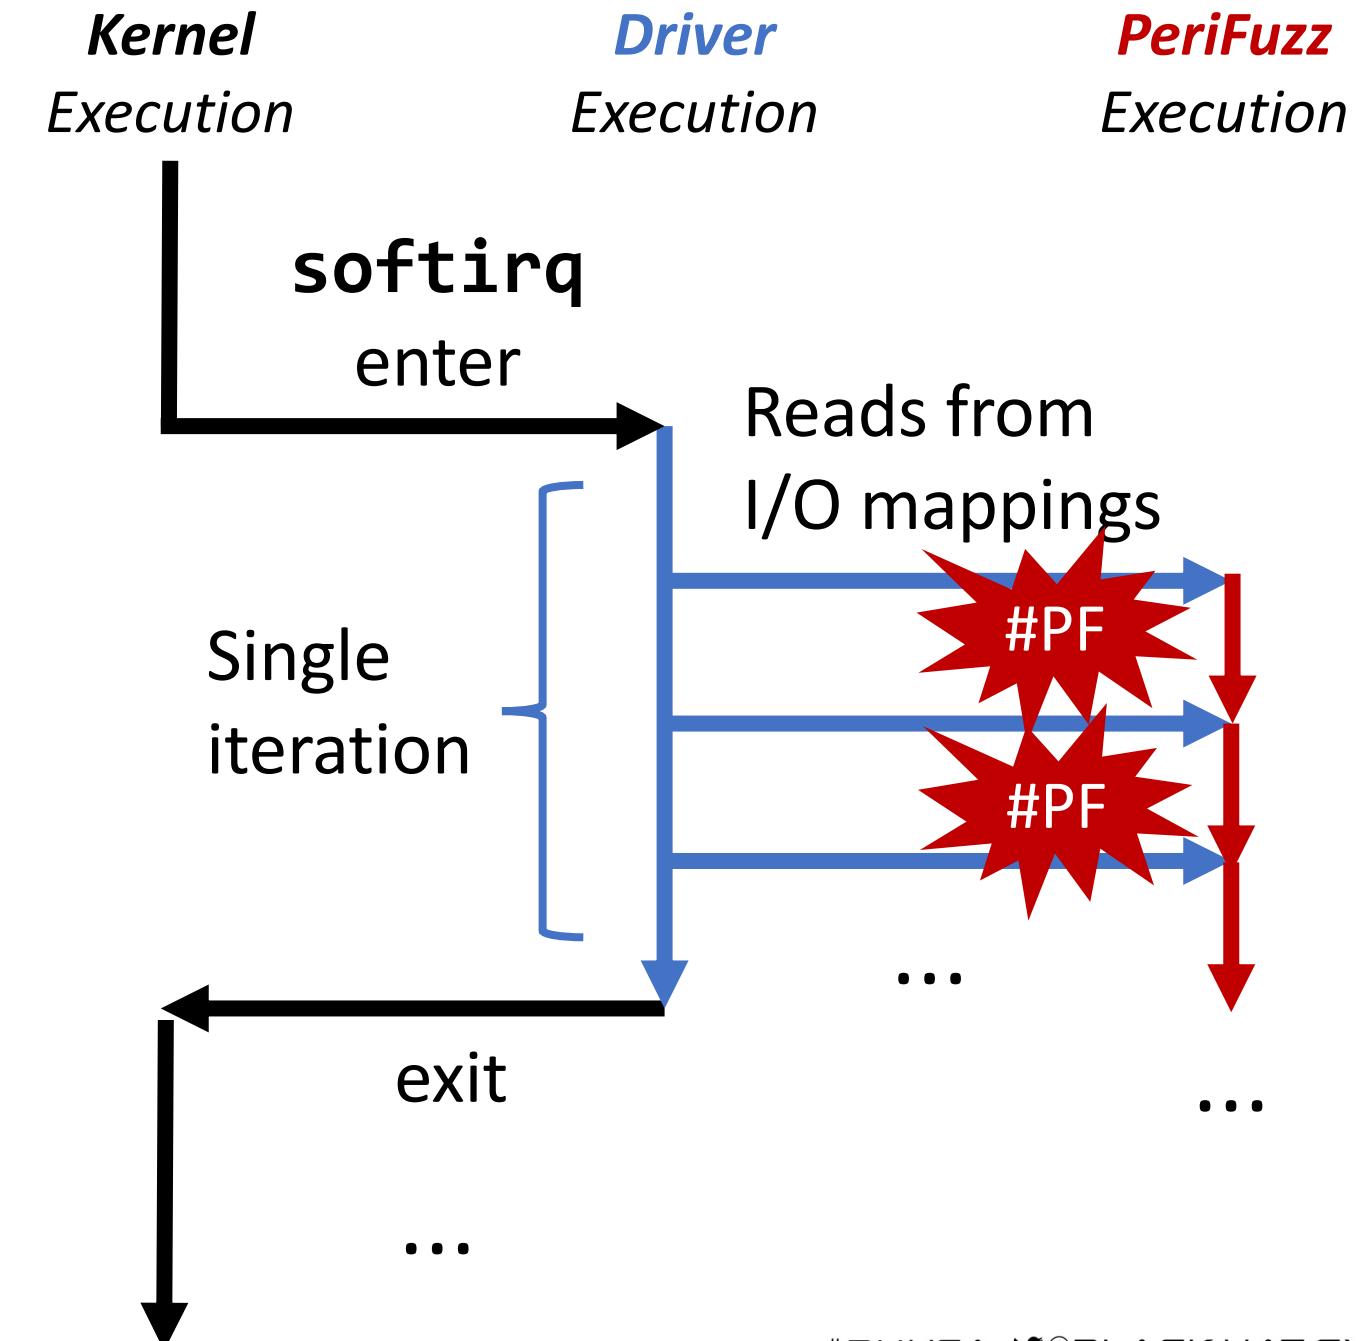


Sequential Fuzzer Input Consumption



Fuzzing Loop

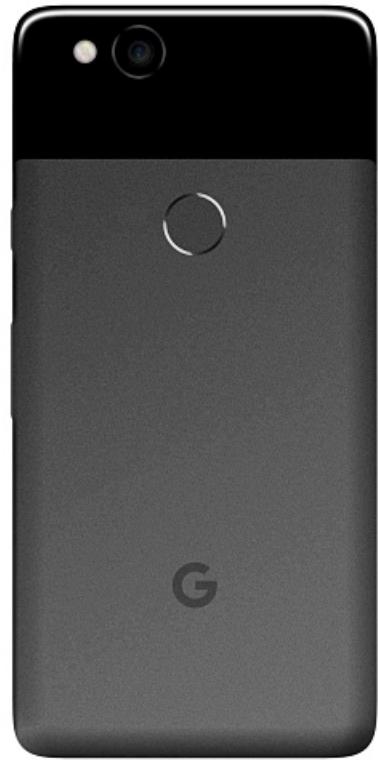
- Each iteration of the fuzzing loop consumes a single fuzzer-generated input
- aligned to the execution of software interrupt (softirq) handler's enter & exit
- can have **one or more reads** from I/O mappings.



Prototype Implementation

- Based on Linux kernel 4.4 for AArch64 (Google Pixel 2)
- Ported to 3.10 (Samsung Galaxy S6)
- AFL 2.42b as *PeriFuzz* front-end

Fuzzing Target: Wi-Fi Drivers



Qualcomm's Wi-Fi driver
in Google Pixel 2



Broadcom's Wi-Fi driver
in Samsung Galaxy S6

Fuzzing Target: Wi-Fi Drivers

1. Large codebase
 - Qualcomm's: 443,222 SLOC and Broadcom's: 122,194 SLOC
2. Highly concurrent
 - heavy use of bottom-half handlers, kernel threads, etc.
3. Lots of code runs in interrupt & kernel thread contexts
 - rather than system call contexts
4. No virtual device implementation available
5. No hypervisor support
 - EL2 not available in production smartphones

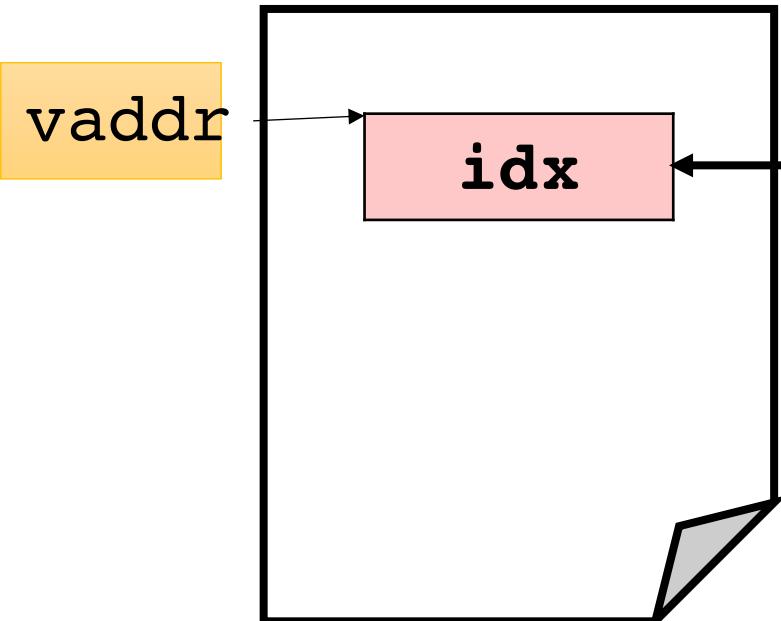
Bugs Found

- Different classes of bugs
 - 9 buffer overreads or overwrites
 - 4 double-fetch issues
 - 1 kernel address leak
 - 3 reachable assertions
 - 2 null pointer dereferences
- In total, 15 vulnerabilities discovered
 - 9 previously unknown
 - 8 new CVEs assigned

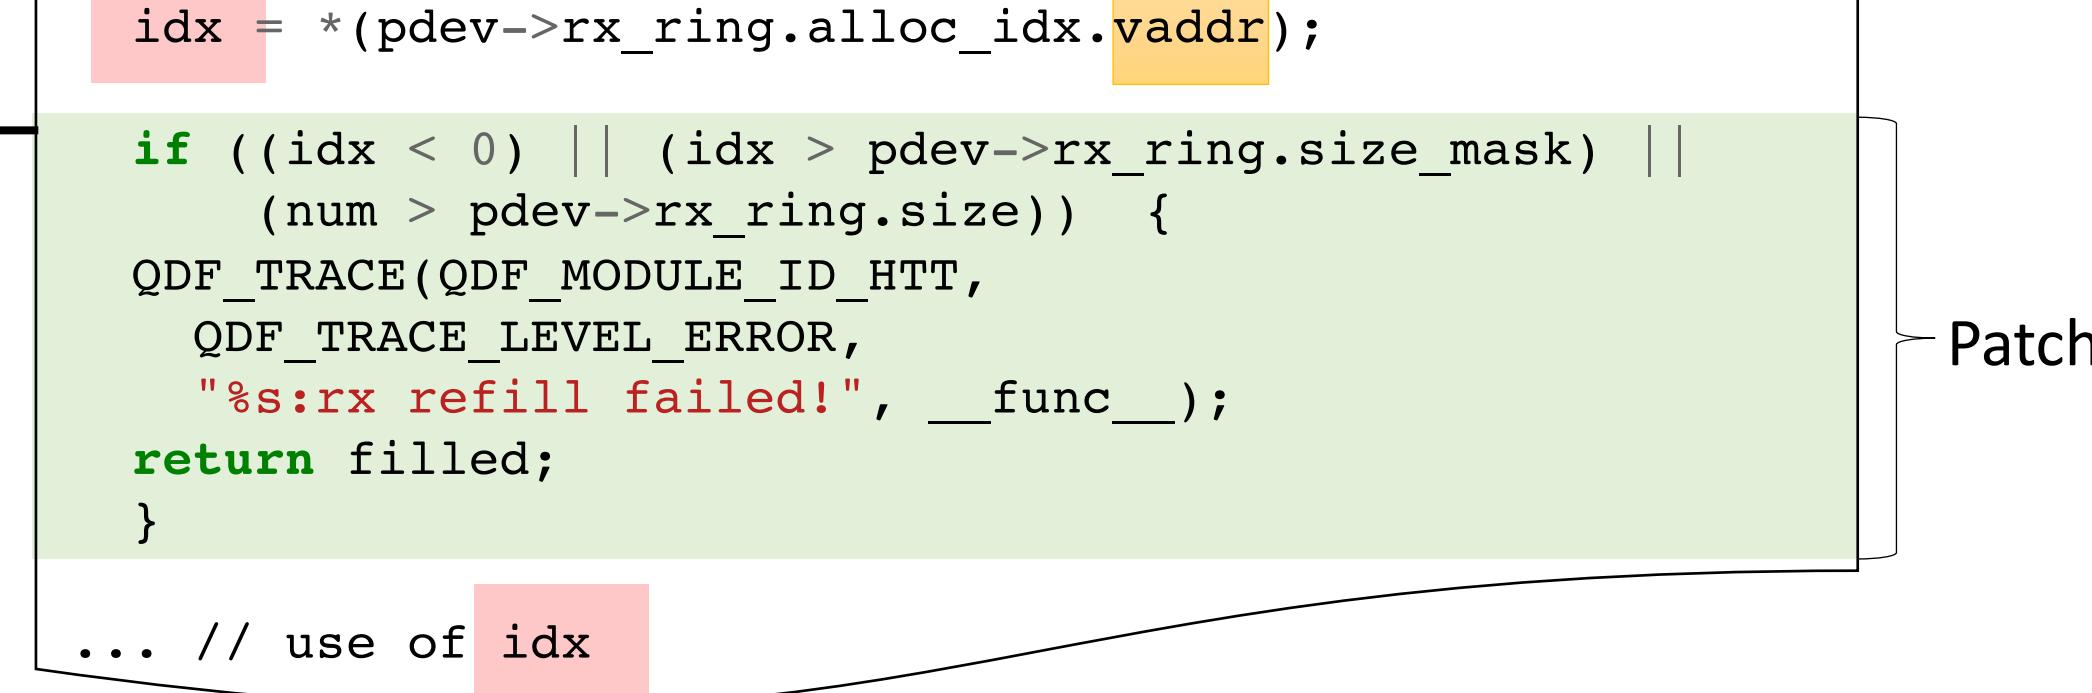
Buffer Overflow (CVE-2018-11902)

Driver used a value read from a DMA mapping as an index into an array without validation (now patched!)

DMA I/O mapping



Driver Source Code



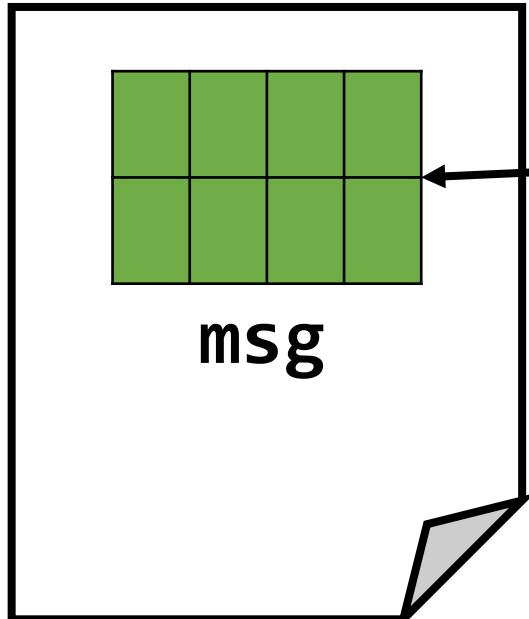
```
idx = *(pdev->rx_ring.alloc_idx.vaddr);  
  
if ((idx < 0) || (idx > pdev->rx_ring.size_mask) ||  
    (num > pdev->rx_ring.size)) {  
    QDF_TRACE(QDF_MODULE_ID_HTT,  
              QDF_TRACE_LEVEL_ERROR,  
              "%s:rx refill failed!", __func__);  
    return filled;  
}  
  
... // use of idx
```

Patch

Double-fetch Bug – Initial Fetch & Check

- 1 The driver computes and verifies the checksum of a message

DMA I/O mapping



Driver Source Code

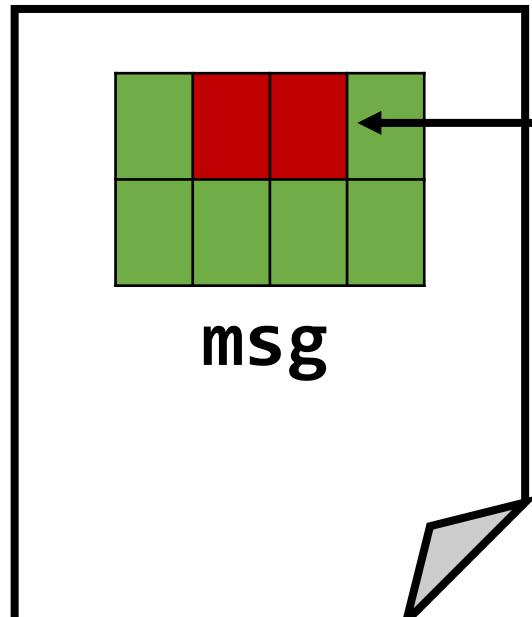
```
static uint8 dhd_prot_d2h_sync_xorcsu(m...)
```

...
prot_checksum = bcm_compute_xor32((volatile uint32 *)msg, num_words);
if (prot_checksum == 0U) { /* checksum is OK */
 if (msg->epoch == ring_seqnum) {
 ring->seqnum++; /* next expected sequence number */
 goto dma_completed;
 ...

Double-fetch Bug – Overlapping Fetch & OOB

- ② The driver fetches the same bytes again from msg

DMA I/O mapping



Driver Source Code

```
ifidx = msg->cmn_hdr.if_id;  
...  
ifp = dhd->iflist[ifidx];
```

Overlapping fetch (fuzzed)

Out-of-bounds access

Unable to handle kernel paging request at virtual address **2f6d657473797337**

Kernel panic - not syncing: Fatal exception in interrupt

Kernel Address Leak (CVE-2018-11947)

Unable to handle kernel paging request at virtual address
17000000d7ff0008

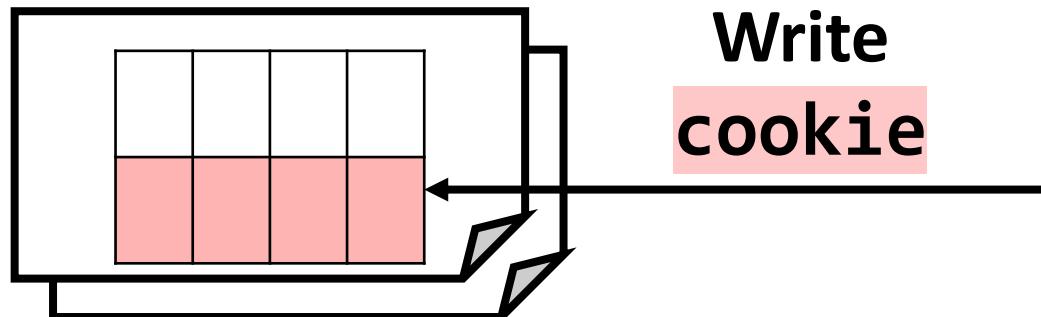
Kernel panic - not syncing: Fatal exception in interrupt

Symptom: A fuzzed value provided by *PeriFuzz* was *directly* being dereferenced.

Kernel Address Leak (CVE-2018-11947)

① Driver sends a kernel pointer to the device

DMA I/O mappings

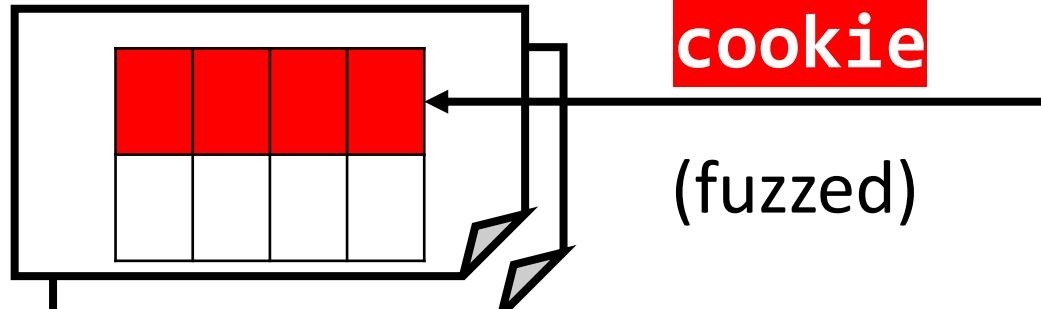


Driver Source Code

```
non_volatile_req = qdf_mem_malloc(sizeof(*non_volatile_req));  
...  
// use pointer as cookie (which is later sent to the device)  
cookie = ol_txrx_stats_ptr_to_u64(non_volatile_req);  
...
```

② Device sends the cookie back, which is then dereferenced by the driver

Read cookie



```
req = ol_txrx_u64_to_stats_ptr(cookie);  
...  
req->... // A value read from I/O mapping is dereferenced
```

Fuzzing Throughput

- Fuzzing throughput is about 7~24 inputs/sec depending on the nature of the I/O mapping being fuzzed.
- The number of page faults is the main contributor. (e.g., 50 page faults per iteration gives around 20 inputs/sec)
- Rooms for improvement. (Details in the paper)

Phone/Driver	I/O Mapping	Peak Throughput (# of test inputs/sec)
Pixel 2 - QCACLD-3.0	QC1	23.67
	QC2	15.64
	QC3	18.77
	QC4	7.63
Galaxy S6 - BCMDHD4358	BC1	9.90
	BC2	14.28
	BC3	10.49
	BC4	15.92

cf) On Pixel 2, Syzkaller achieves on average 24 program executions per second (max: ~60).
(1 proc ADB-based configuration measured for a 15-min period)

Future Work

- Minimizing the impact of shallow bugs
 - All bugs found in less than 10,000 inputs
 - Shallow bugs frequently hit, which causes system restarts (reboot takes 1 min)
 - We had to manually disable subpaths rooted at bugs already found
- Improving throughput
 - Slower than, for example, typical user-space fuzzing
 - Possible optimizations and trade-offs outlined in the paper

Conclusion

- Remote peripheral compromise poses a serious threat to OS kernel security.
- PeriScope and PeriFuzz are practical dynamic analysis tools that can analyze large, complex drivers along the hardware-OS boundary.
- PeriScope and PeriFuzz are effective at finding vulnerabilities along the HW-OS boundary.
 - Memory overreads/overwrites, address leak, null pointer dereferences, reachable assertions, and double-fetch bugs

Q & A

Thank you!

Contact

Dokyung Song

Ph.D. Student at UC Irvine

dokyungs@uci.edu